

Docker Fundamentals





WORKFORCE DEVELOPMENT



PARTICIPANT GUIDE

logistics



- **Class Hours:**
- Instructor will provide class start and end times.
- Breaks throughout class.

- **Lunch:**
- Yes, 1 hour and 15 minutes
- Extra time for email, phone calls, or simply a walk.



- **Telecommunication:**
- Turn off or set electronic devices to vibrate
- Reading or attending to devices can be distracting to other students

- **Miscellaneous**
- Courseware
- Bathroom

Course Objectives

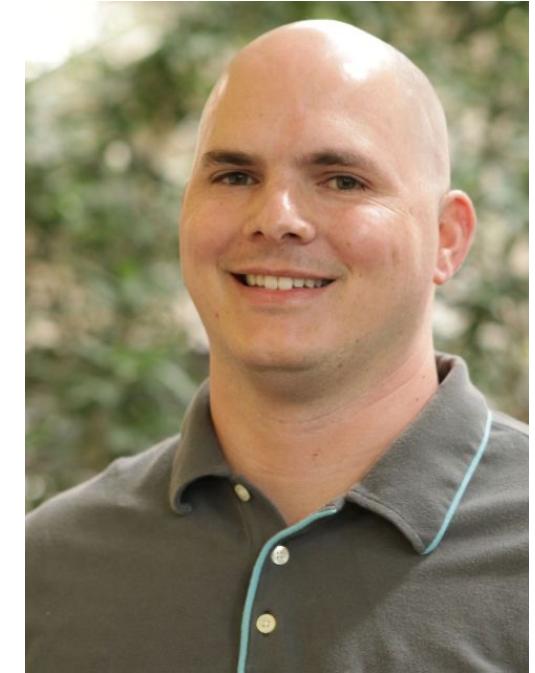
By the end of the course you will be able to:

- Describe the Docker platform and architecture
- Build, run, and manage Docker containers & images
- Configure container networks
- Maintain data in containers
- Configure a Docker Swarm cluster
- Install Docker on laptop for developer use

Hi!

Jason Smith

Cloud Consultant with a Linux sysadmin background.
Focused on cloud-native technologies: automation,
containers & orchestration



github

<https://github.com/jruels>

Expertise

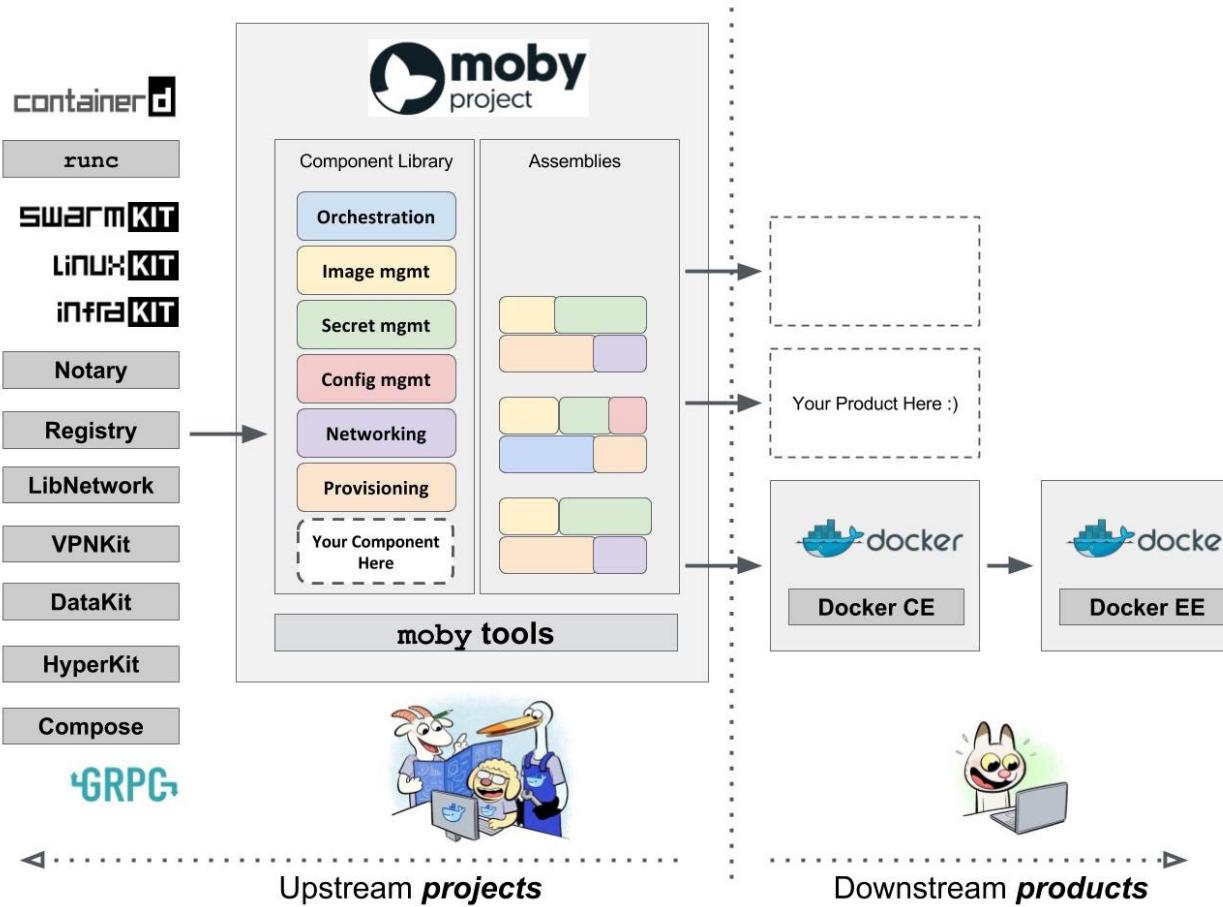
- Cloud
- Automation
- CICD
- Docker
- Kubernetes

Introductions

- Name
- Job Role
- Your experience with Docker (scale 1 - 5)
- Experience using a command-line text editor
- Expectations for course (please be specific)

Docker Overview

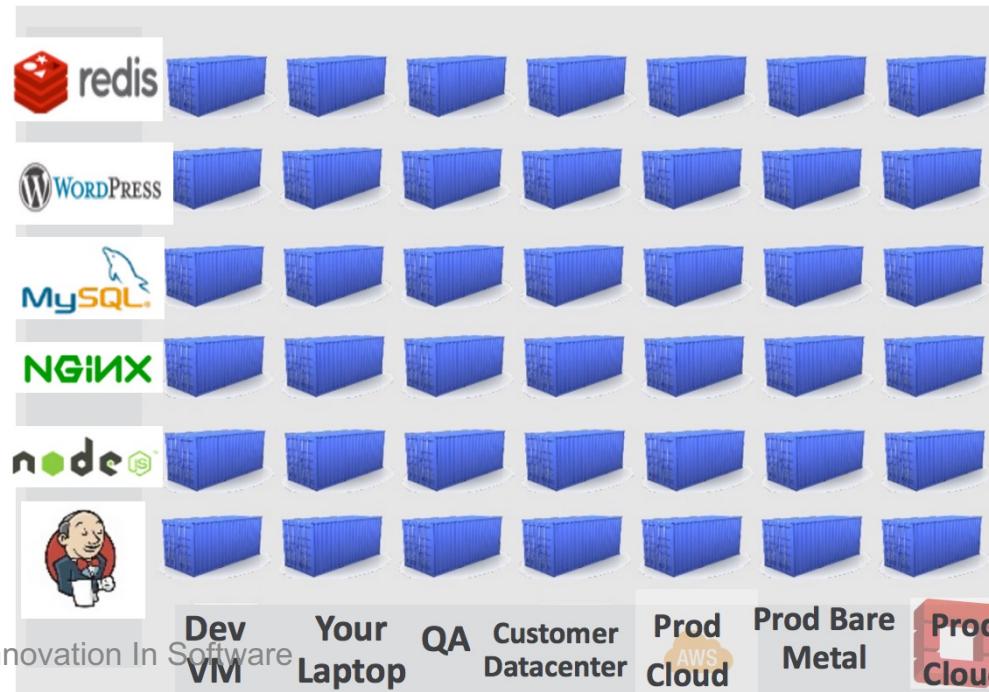
What is Docker?



What is Docker?

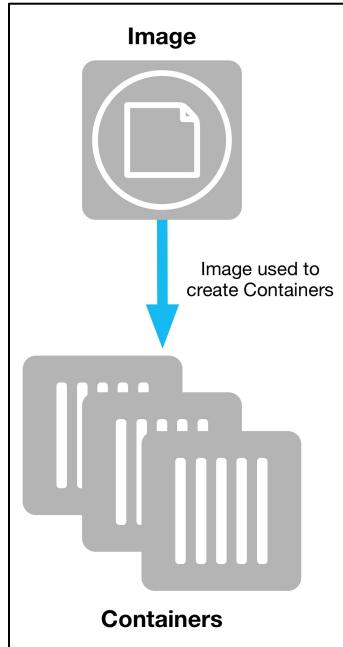


Docker allows you to package an application with all of its dependencies into a standardized unit for software development.



Terminology

Image



Read only template used to create containers

Built by you or other Docker users
Stored in Docker Hub, Docker Trusted Registry or your own Registry

Terminology

Image

Read only template used to create containers

Built by you or other Docker users

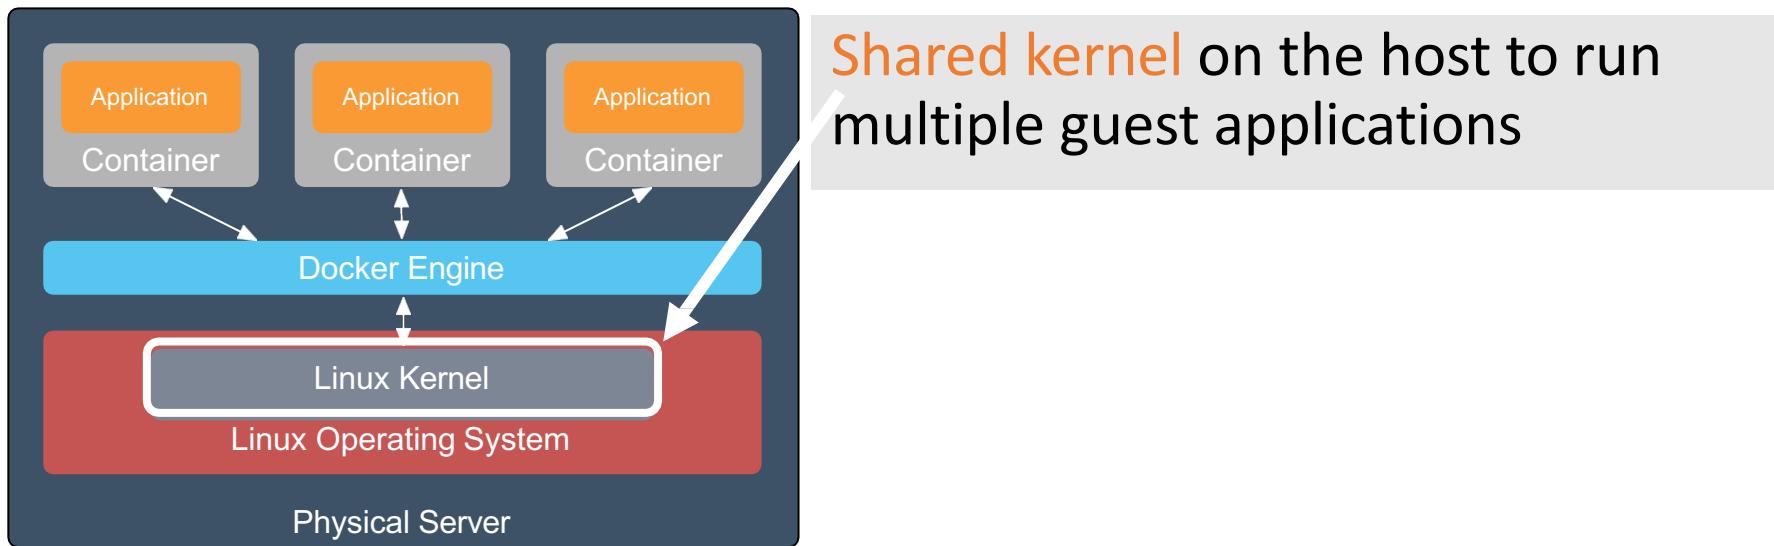
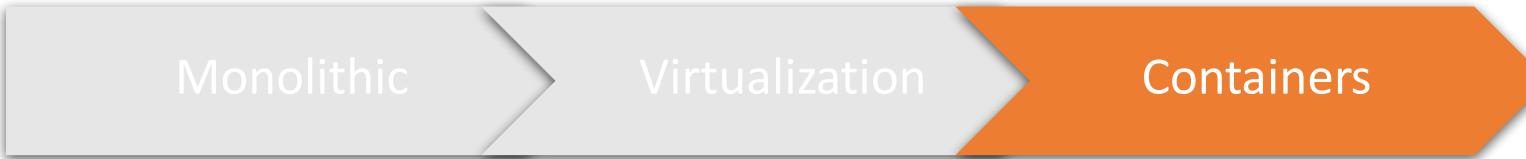
Stored in Docker Hub, Docker Trusted Registry or your own Registry

Container

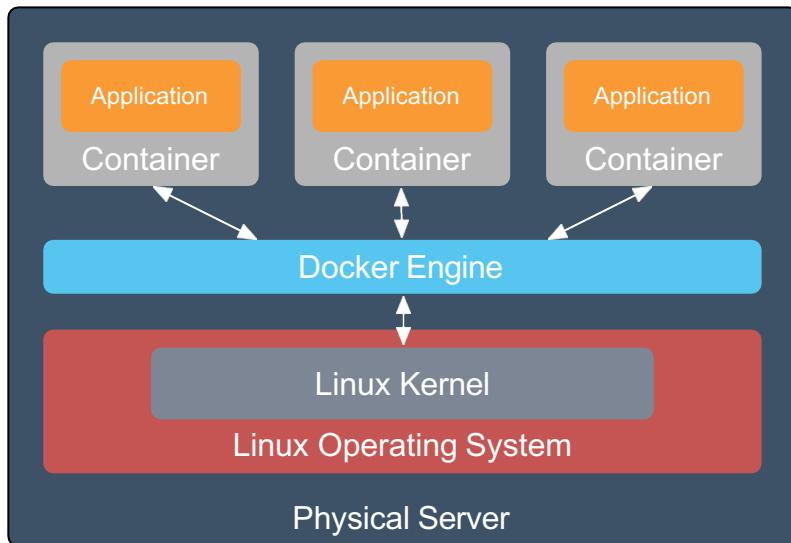
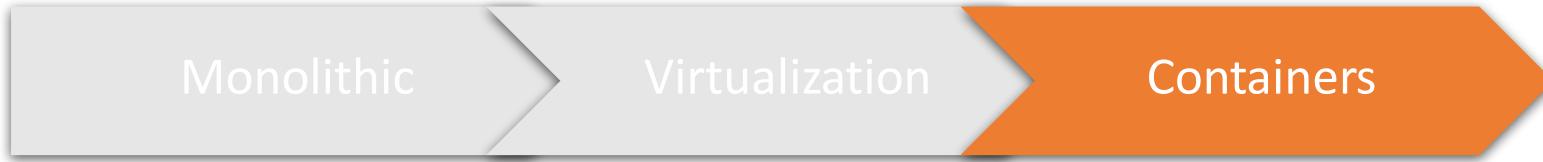
- Isolated application platform
- Contains everything needed to run your application
- Based on one or more images

Containers

Containers



Containers - Advantages

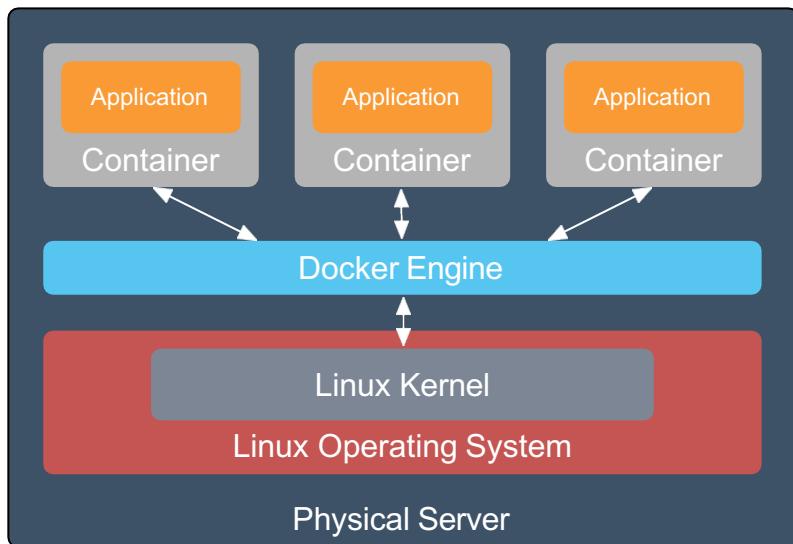
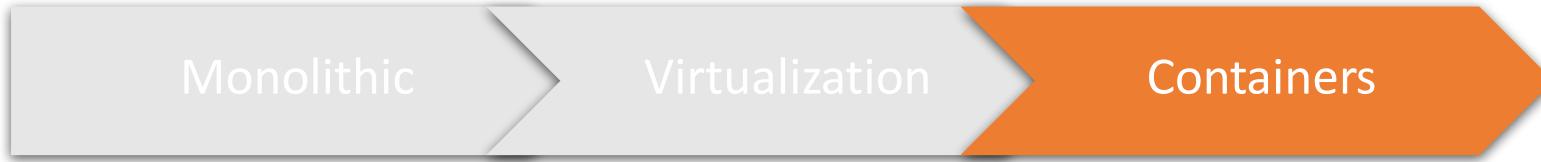


Shared kernel on the host to run multiple guest applications

Advantages over VMs

- Containers are more lightweight
- No need to install a guest Operating System
- Less CPU, RAM, storage overhead
- More containers per machine
- Greater portability

Containers - Challenges



Shared kernel on the host to run multiple guest applications

Container Challenges

- Early Docker focused on single-node operations
- Up to user to cluster Docker hosts and manage deployment of containers on cluster
- User solves for automatic scale out of applications
- User solves for service discovery between application components (microservices)

Container – Concept of Operations

Container based virtualization

Uses the kernel on the host operating system to run multiple guest instances

- Each guest instance is a container
- Each container has its own

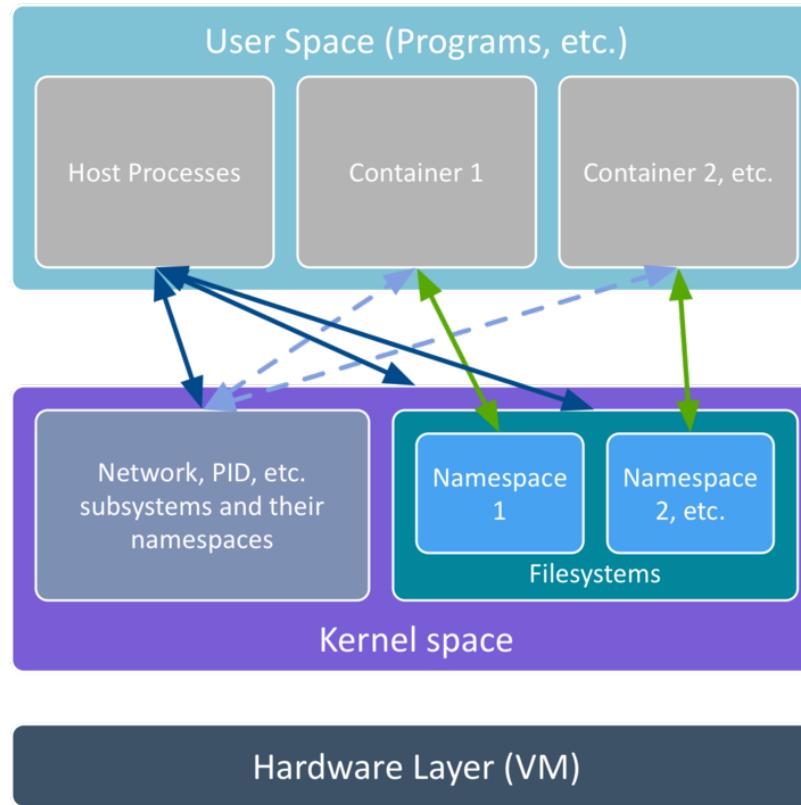
- Root filesystem
- Processes
- Memory
- Devices
- Network Ports



Isolation with Namespaces

Namespaces - Limits what a container can see (and therefore use)

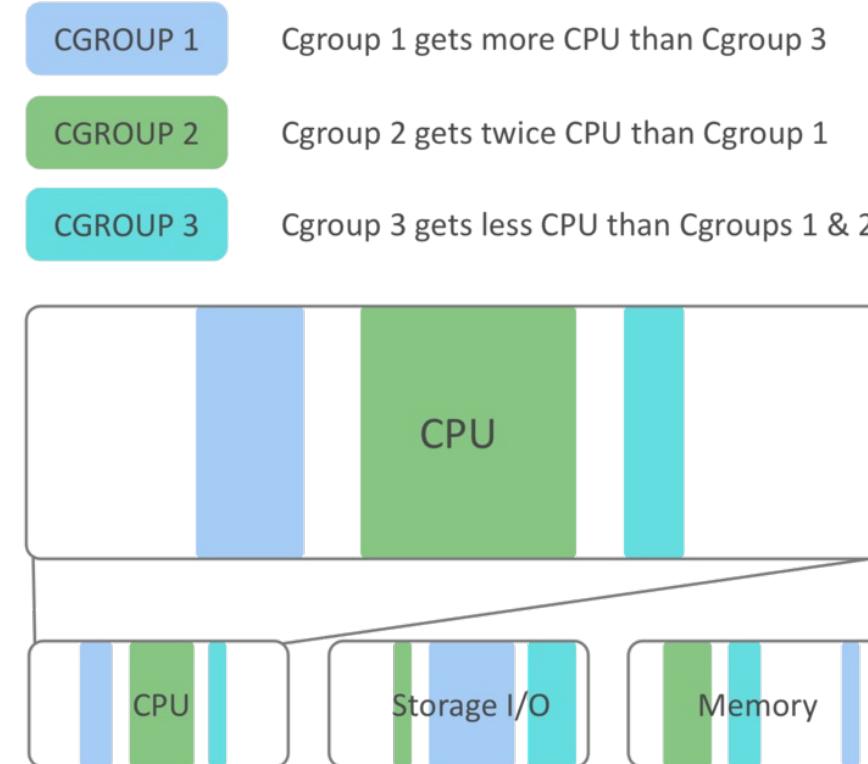
- Namespace wrap a global system resource in an abstraction layer
- Processes running in that namespace think they have their own, isolated resource
- Isolation includes:
 - Network stack
 - Process space
 - Filesystem mount points
 - etc.



Isolation with Control group (Cgroups)

Cgroups - Limits what a container can use

- Resource metering and limiting
 - CPU
 - MEM
 - Block/I/O
 - Network
- Device node (`/dev/*`) access control



Container Use Cases

DevOps



Developers

Focus on applications inside the container



Operations

Focus on orchestrating and maintaining
containers in production

Container Use Cases

Development

- Allows the ability to define the entire project configuration and tear-down/recreate it easily
- Supports multiple versions of application simultaneously

Container Use Cases

Development

- Allows the ability to define the entire project configuration and tear-down/recreate it easily
- Supports multiple versions of application simultaneously

Test Environments

- Same container that developers run is the container that runs in test lab and production – includes all dependencies
- Well formed API allows for automated building and testing of new containers

Container Use Cases

Development

- Allows the ability to define the entire project configuration and tear-down/recreate it easily
- Supports multiple versions of application simultaneously

Test Environments

- Same container that developers run is the container that runs in test lab and production – includes all dependencies
- Well formed API allows for automated building and testing of new containers

Micro-Services

- Design applications as suites of services, each written in the best language for the task
- Better resource allocation
- One container per microservice vs. one VM per microservice
- Can define all interdependencies of services with templates



Questions

The Docker Platform Components

Docker Engine

Docker Engine

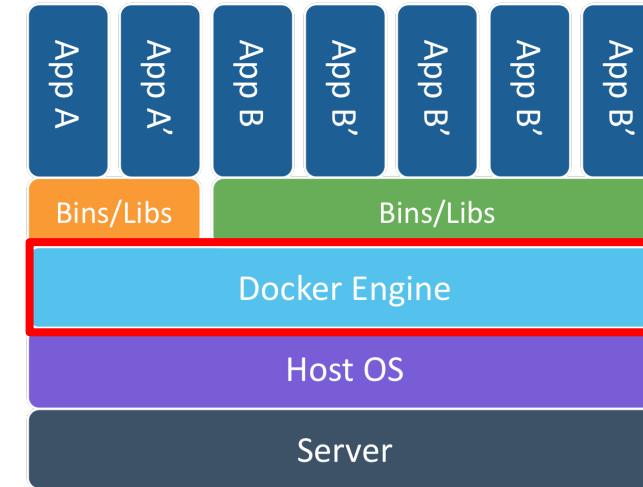
Docker Registry

Docker Compose

Docker Swarm

*Lightweight runtime program to **build, ship, and run Docker containers***

- Also known as **Docker Daemon**
- Uses Linux Kernel namespaces and control groups
 - **Linux Kernel (>= 3.10)**
- Namespaces provide an isolated workspace



Docker Engine

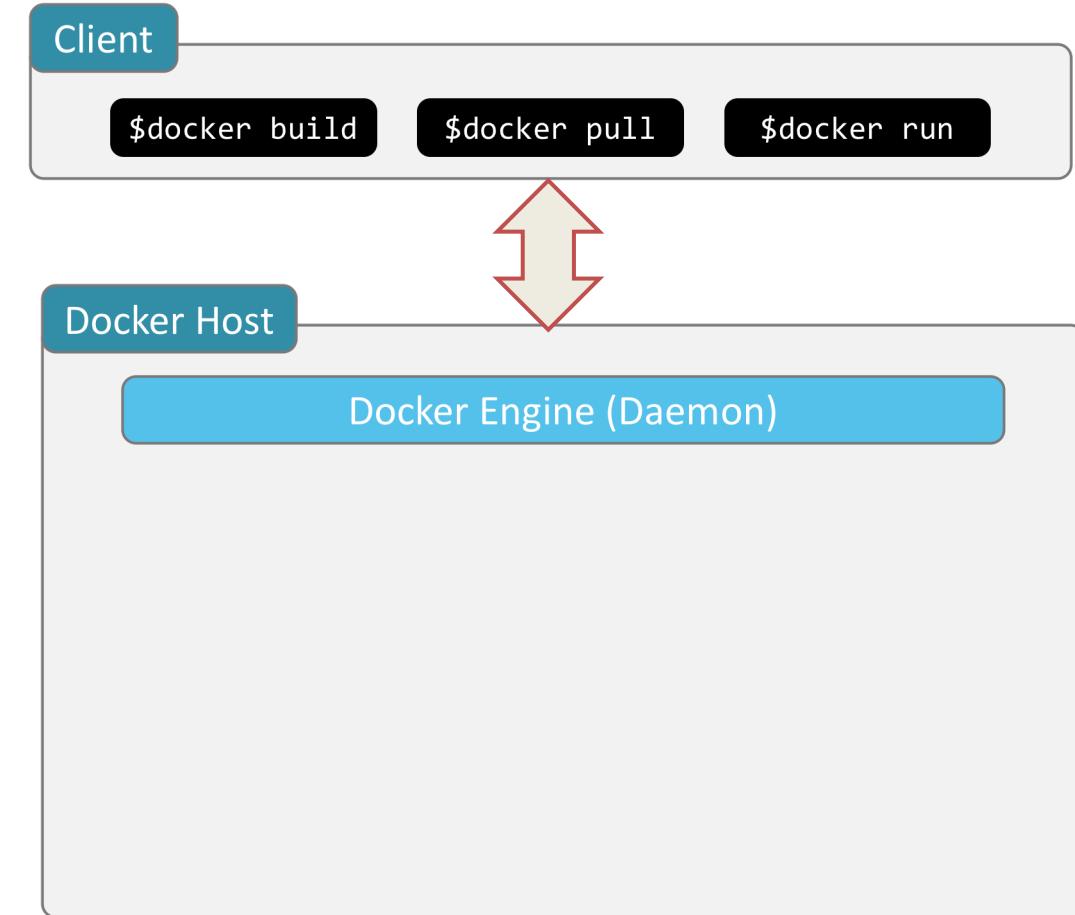
Docker Engine

Docker Registry

Docker Compose

Docker Swarm

- The Docker Client is the docker binary
 - Primary interface to the Docker Host
 - Accepts commands and communicates with the Docker Engine (Daemon)



Docker Engine

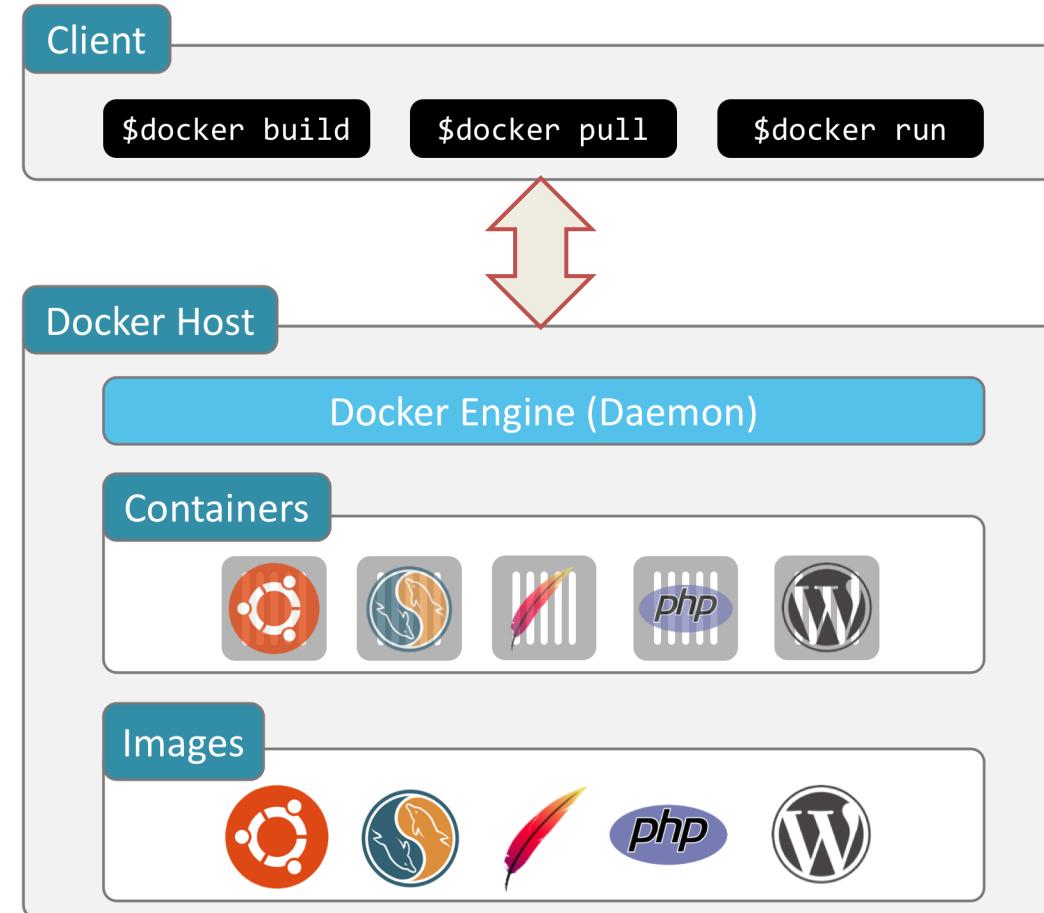
Docker Engine

Docker Registry

Docker Compose

Docker Swarm

- Lives on a Docker host
- Creates and manages containers on the host



Docker Registry

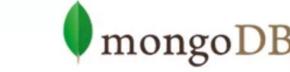
Docker Engine

Docker Registry

Docker Compose

Docker Swarm

Image Storage & Retrieval System



- Docker Engine **Pushes** Images to a Registry
- Version Control
- Docker Engine **Pulls** Images to Run

Docker Registry

Docker Engine

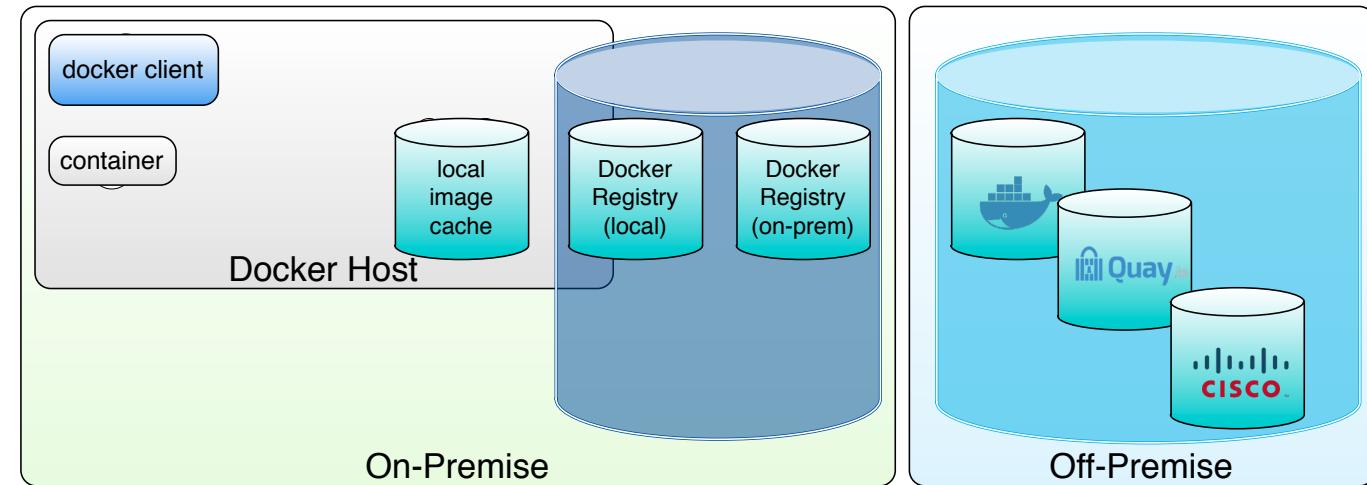
Docker Registry

Docker Compose

Docker Swarm

Types of Docker Registries

- Local Docker Registry (On Docker Host)
- Remote Docker Registry (On-Premise/Off-Premise)
- Docker Hub (Off-Premise)



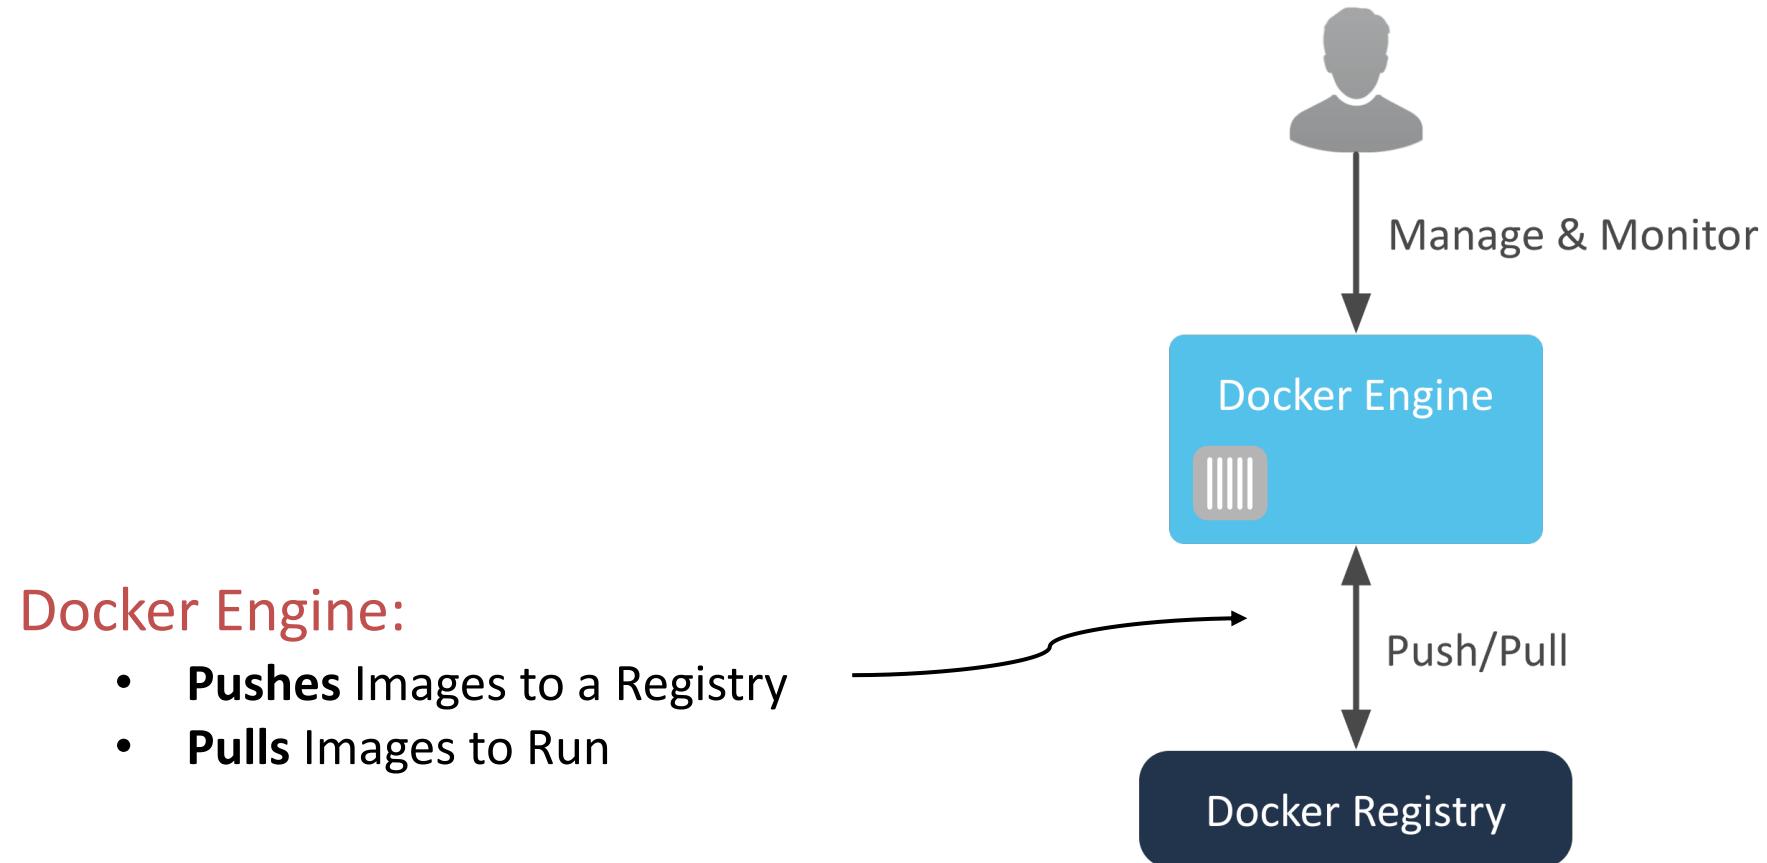
Docker Engine and Registry

Docker Engine

Docker Registry

Docker Compose

Docker Swarm



Docker Registry

Docker Engine

Docker Registry

Docker Compose

Docker Swarm

The registry and engine both present APIs

- All of Docker's functionality will utilize these APIs
- RESTFUL API
- Commands presented with Docker's CLI tools can also be used with curl and other tools

Docker Compose

Docker Engine

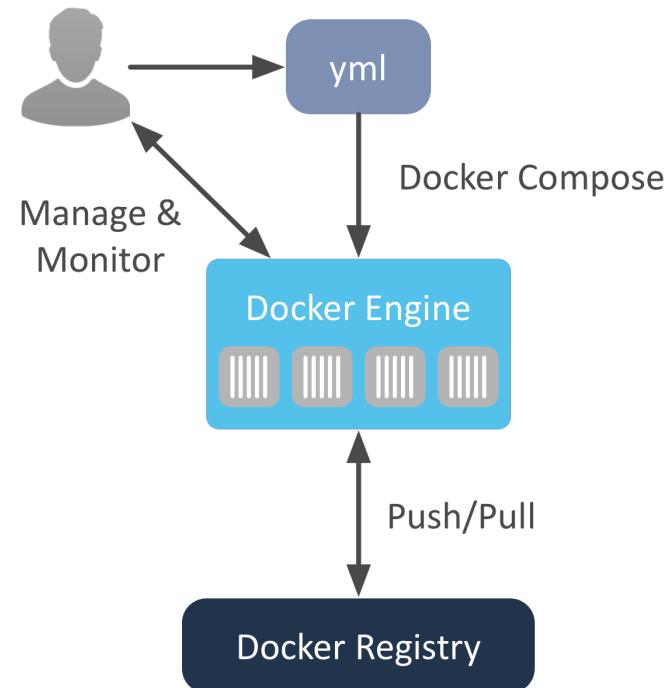
Docker Registry

Docker Compose

Docker Swarm

Tool to create and manage multi-container applications

- Applications defined in a single file:
docker-compose.yml
- Transforms applications into individual containers that are linked together
- Compose will start all containers in a single command



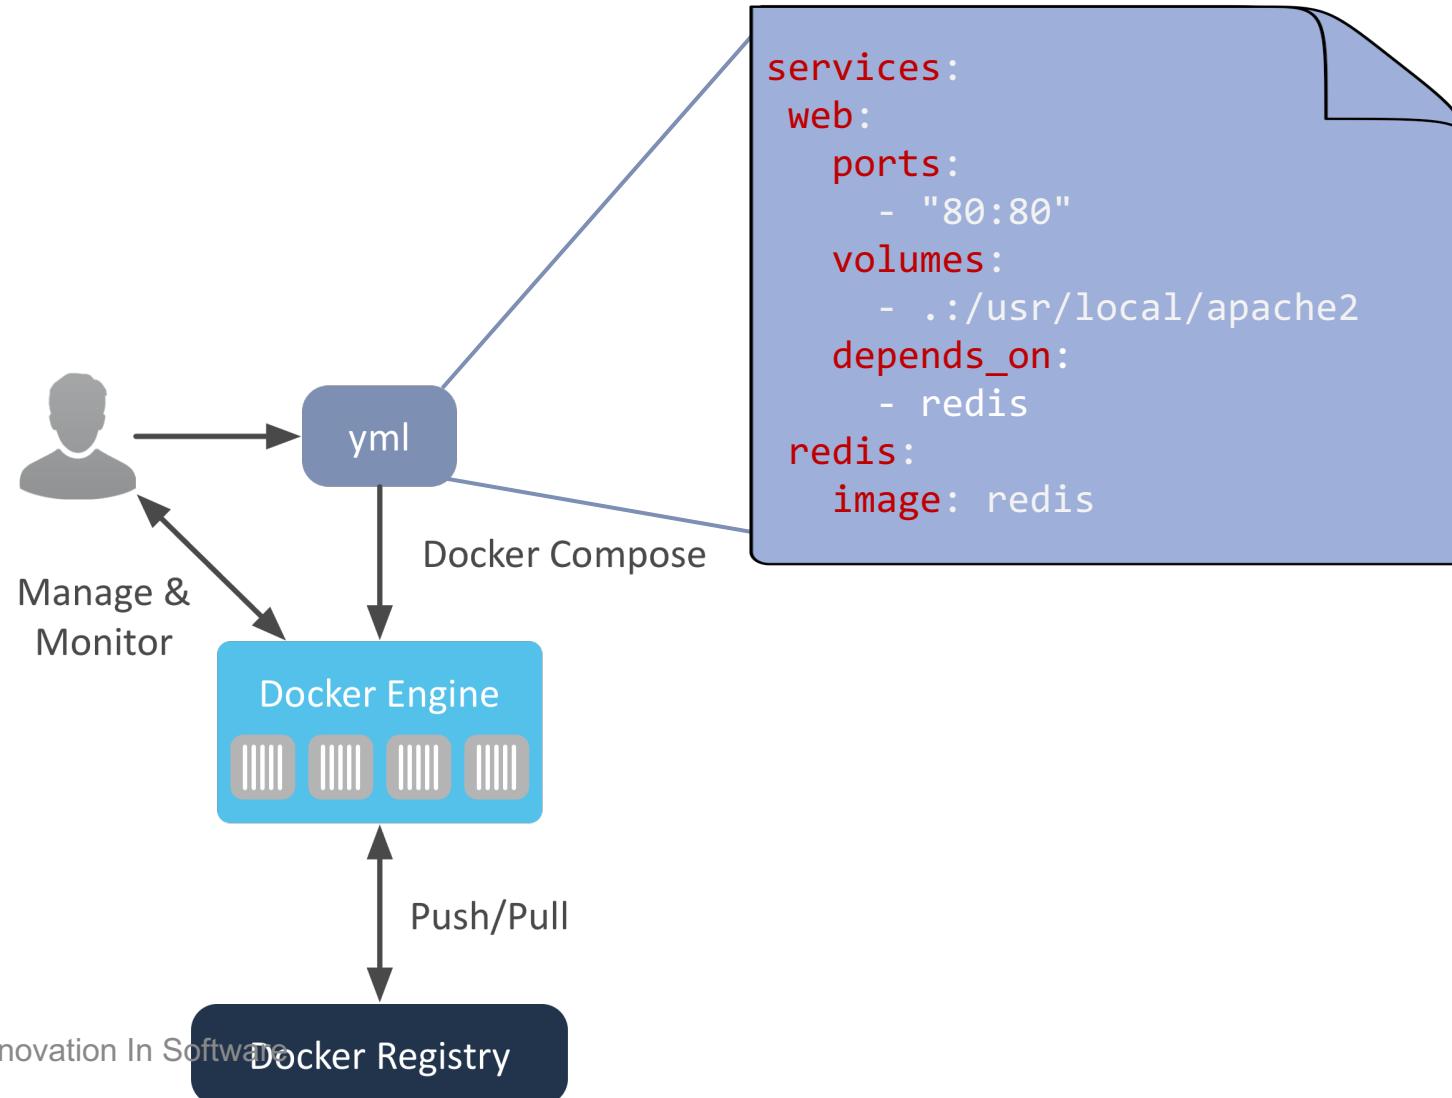
Docker Compose

Docker Engine

Docker Registry

Docker Compose

Docker Swarm



Docker Swarm

Docker Engine

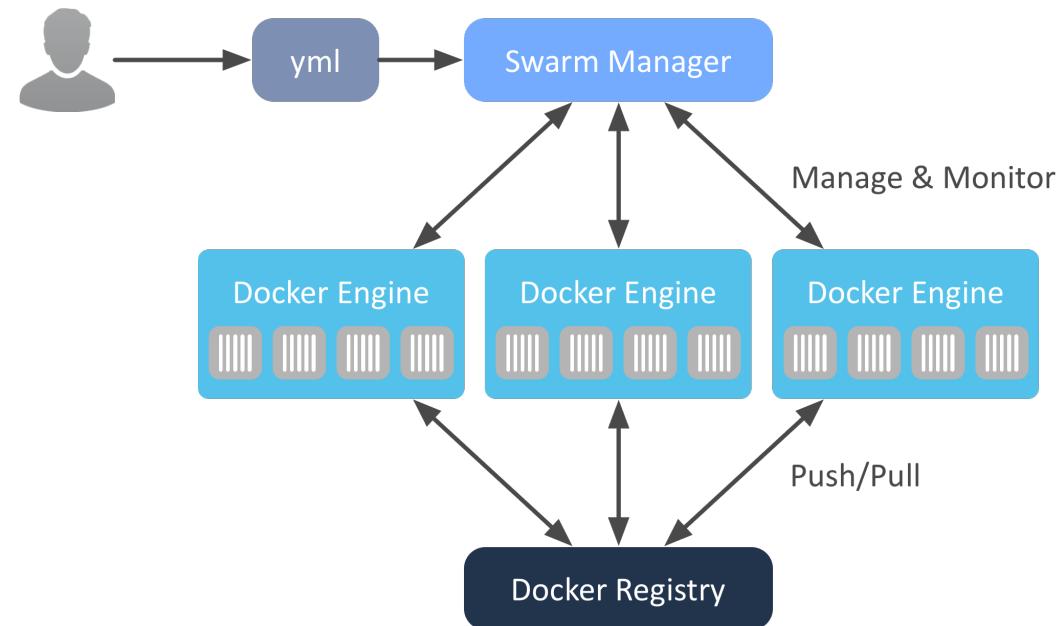
Docker Registry

Docker Compose

Docker Swarm

Clusters Docker hosts and schedules containers

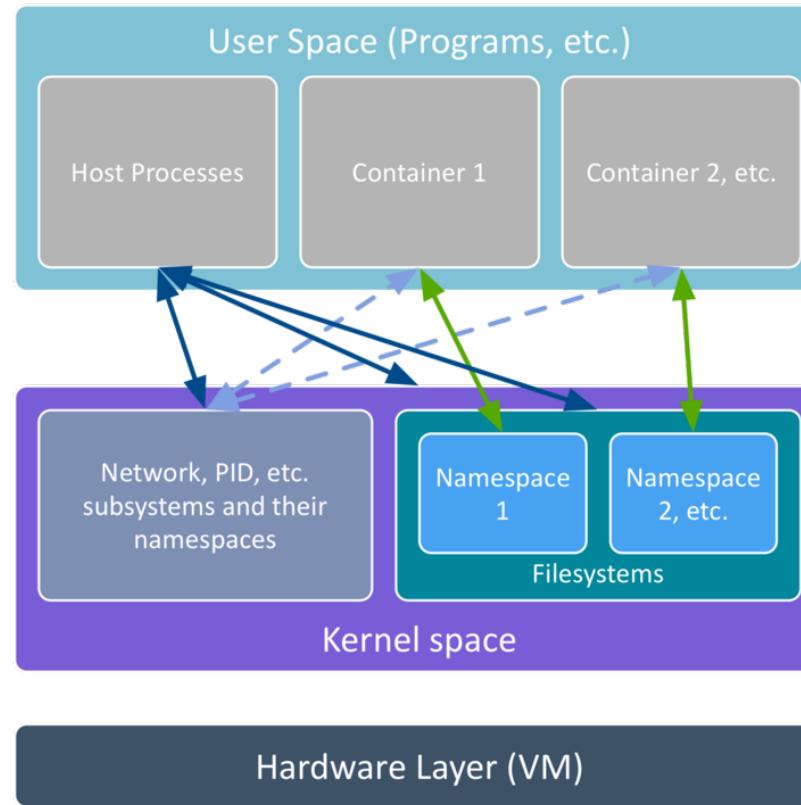
- Native Clustering for Docker
 - Turn a pool of Docker hosts into a single, virtual host
 - Serves the standard Docker API



Isolation with Namespaces

Namespaces - Limits what a container can see (and therefore use)

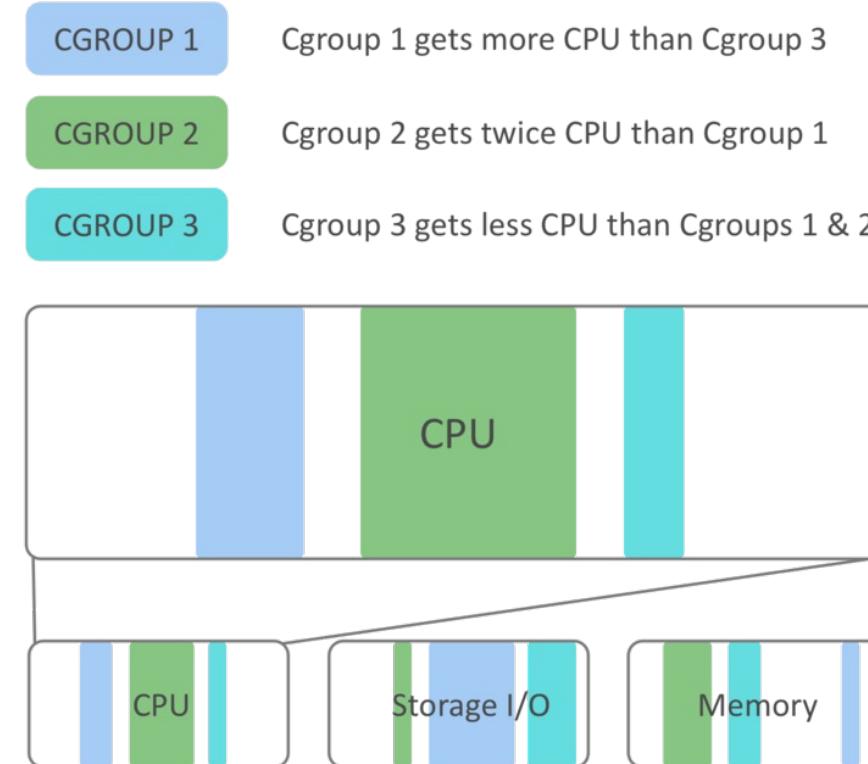
- Namespace wrap a global system resource in an abstraction layer
- Processes running in that namespace think they have their own, isolated resource
- Isolation includes:
 - Network stack
 - Process space
 - Filesystem mount points
 - etc.



Isolation with Control group (Cgroups)

Cgroups - Limits what a container can use

- Resource metering and limiting
 - CPU
 - MEM
 - Block/I/O
 - Network
- Device node (`/dev/*`) access control

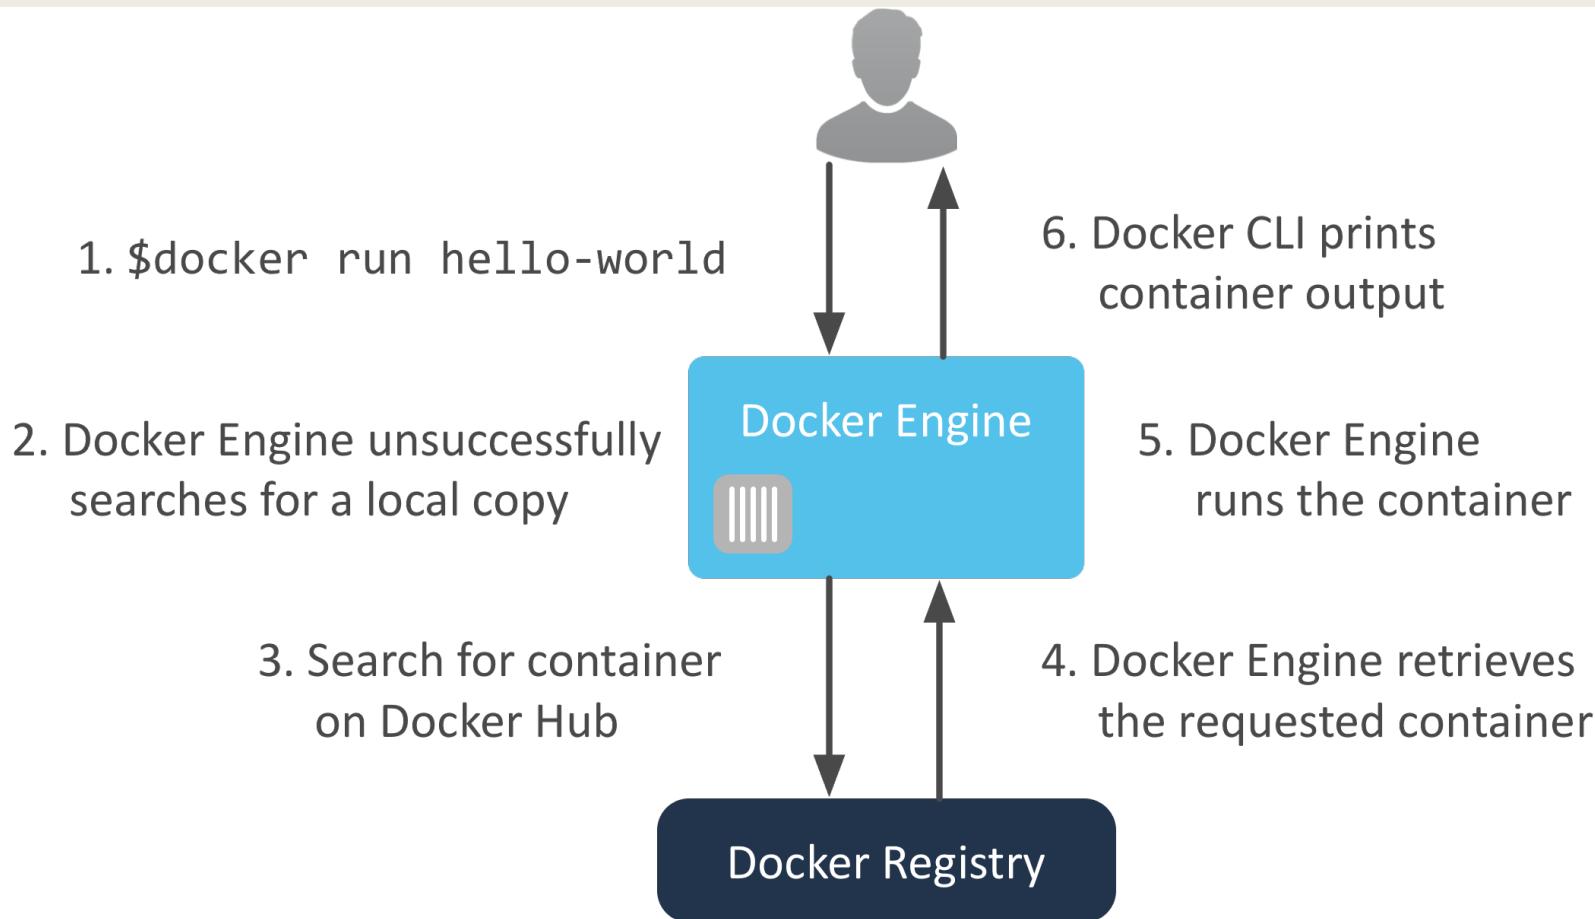


Docker Lab Setup

Run a Container

Container Review

What does running a container require?



Run a Container

In your Lab Environment, execute the following command:

```
$ docker run hello-world
```

The container will run and display the following:

```
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
ad01a741289f: Pull complete
Digest:
sha256:af451bc5c4d6e87b64303f707f7bd98debd7fd3175a676c08f4bf46cb48813d4
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker!

This message shows that your installation appears to be working correctly.

Run a Container

Containers can accept additional arguments

In your Lab Environment, execute the following command:

```
$ docker run docker/whalesay \
    cowsay 'Go Docker'
```

The container will run and display the following:

```
Unable to find image 'docker/whalesay:latest' locally
latest: Pulling from docker/whalesay

e190868d63f8: Pull complete
909cd34c6fd7: Pull complete
0b9bfabab7c1: Pull complete
a3ed95caeb02: Pull complete
00bf65475aba: Pull complete
c57b6bcc83e3: Pull complete
8978f6879e2f: Pull complete
8eed3712d2cf: Pull complete
Digest:
sha256:178598e51a26abbc958b8a2e48825c90bc22e641de3d31e18aaf55f3258ba93b
Status: Downloaded newer image for docker/whalesay:latest
```

Run a Container

docker run also accepts additional arguments

In your Lab Environment, execute the following command:

```
$ docker run -d docker/whalesay \
    cowsay 'Detached'
```

The container will run and display the following:

```
3fa9d2d4f5ad3169d2e763f7ef48436c15ca4f8749a044e3017e6cb96cc60a2d
```

Run a Container

What just happened?

```
$ docker run -d docker/whalesay \
    cowsay 'Detached'
```

Run a Container

```
$ docker run -d docker/whalesay \  
    cowsay 'Detached'
```

The *-d [detach]* flag instructs the Docker Engine to run the container in the background.

```
docker run -help  
  
Usage: docker run [OPTIONS] IMAGE [COMMAND] [ARG...]  
Run a command in a new container  
  
Options:  
-d, --detach          Run container in background and print container ID
```

Run a Container

Why Run a Container in the Background?

Discussion

Why Run a Container in the Background?

- Containers that do not require interaction can be started and left to run in their own process

Monitor Docker – Native Tools

Native Monitoring Tools

These following tools provide real-time metrics on a container or a set of containers.

docker ps <switch>

- list containers

docker logs <container>

- displays console output of the container

docker top <container>

- lists all the processes running in a container

docker stats <container>

- streams real time stats of containers

docker inspect <container>

- displays detailed container configuration

Docker ps (list containers)

Show running Docker containers

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND
ece82d202fd6	grafana/grafana:2.0.2	"/usr/sbin/grafana-se"
2be9d338d8e4	google/cadvisor	"/usr/bin/cadvisor -s"
ffa57323dad8	tutum/influxdb:0.8.8	"/run.sh"

Show Docker containers id

```
$ docker ps -q
```

```
ece82d202fd6
2be9d338d8e4
ffa57323dad8
c5817eea595e
```

Show all Docker containers

```
$ docker ps -a
```

CONTAINER ID	IMAGE	COMMAND
ece82d202fd6	grafana/grafana:2.0.2	"/usr/sbin/grafana-se"
2be9d338d8e4	google/cadvisor	"/usr/bin/cadvisor -s"
ffa57323dad8	tutum/influxdb:0.8.8	"/run.sh"
c5817eea595e	nginx/nginx-php	"/usr/local/bin/run"

Docker ps (list containers)

In your Lab Environment, execute the following command:

```
$ docker ps
```

The command will display the output:

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
--------------	-------	---------	---------	--------	-------	-------

This is expected since we have no running containers...

Docker ps (list containers)

What about all containers?

In your Lab Environment, execute the following command:

```
$ docker ps -a
```

The command will display output similar to:

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
f774f8618226	hello-world	"/hello"	7 seconds ago	Exited (0) 6 seconds ago		timely_terapin
3fa9d2d4f5ad	docker/whalesay	"cowsay 'Go Docker'"	50 minutes ago	Exited (0) 50 minutes ago		ecstatic_feynman
8f5e737d15bf	docker/whalesay	"cowsay 'Detached'"	About an hour ago	Exited (0) 59 minutes ago		happy_leavitt

The process we executed in these containers finished, so the containers exited

Make note (copy) a CONTAINER ID

Docker logs <container>

Display the Console Output of a Container

```
$ docker logs f774f8618226
```

Hello from Docker!

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker Hub account:

<https://hub.docker.com>

For more examples and ideas, visit:

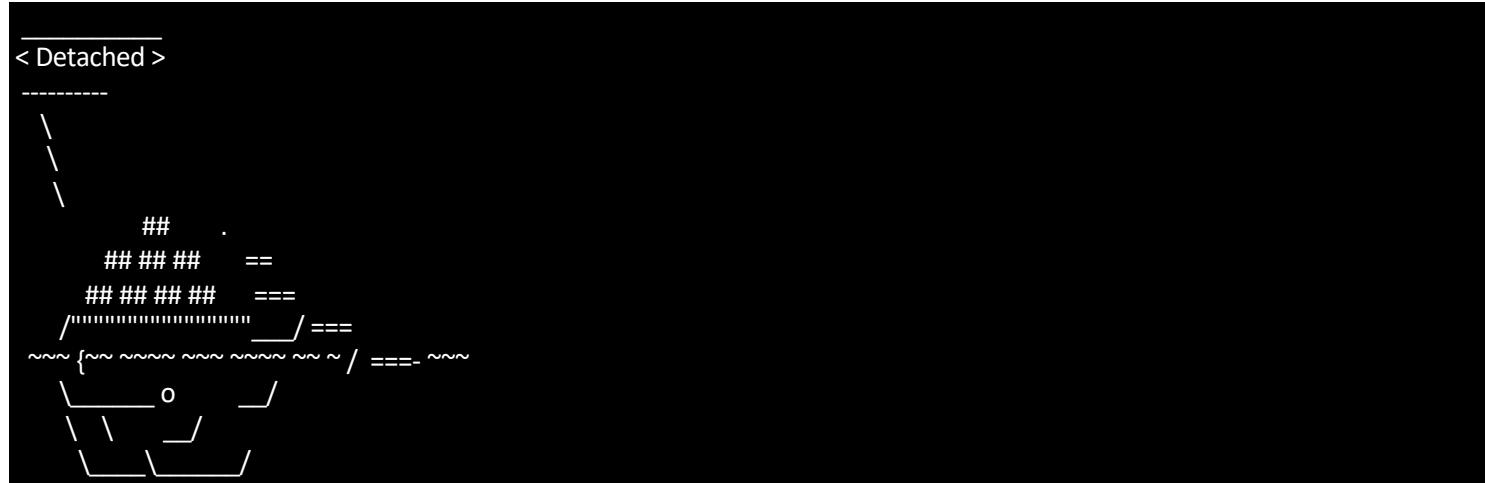
<https://docs.docker.com/engine/userguide/>

Docker logs <container>

Try it!

In your Lab Environment, execute the following command: `$ docker logs <container_id>`

The command will display the output:



Docker top <container>

Running the **docker top** command on all host container

- Flags used with regular top can be used **-aux -faux**

```
$ docker top -help
Usage: docker top [OPTIONS] CONTAINER [ps OPTIONS]
Display the running processes of a container
--help=false      Print usage
```

Docker top <container>

Try it!

In your Lab Environment, execute the following command:

```
$ docker run -d busybox top
```

This will pull and run the latest build of busybox and run the ‘top’ command.

Once the container has started, the console will display the Container ID:

```
7ed45bb572d0886686c5d9ac7484094716cf0c581087d6499193a799a3ed6345
```

Docker top <container>

Now that we have a running container:

In your Lab Environment, execute the following command:

```
$ docker top <container_id>
```

The console will display the running process in the container:

PID	USER	TIME	COMMAND
3372	root	0:03	top

Docker stats <container>

docker stats command

```
$ docker stats --help
Usage: docker stats [OPTIONS] CONTAINER [CONTAINER...]
Display a live stream of container(s) resource usage statistics
--help=false          Print usage
--no-stream=false     Disable streaming stats and only pull the first
result
```

docker stats command on a single container

```
$ docker stats <container_id>
CONTAINER           CPU %               MEM USAGE / LIMIT
cdfd04b5aeeef      0.15%              29.52 MB / 2.098 GB
MEM %               NET I/O            BLOCK I/O
1.41%              6.684 kB / 648 B    0 B / 49.15 kB
```

Docker stats <container>

Try it!

In your Lab Environment, execute the following command:

```
$ docker stats <container_id>
```

The console will display the real-time stats of the container:

CONTAINER	CPU %	MEM USAGE / LIMIT	MEM %	NET I/O
	BLOCK I/O	PIDS		
7ed45bb572d0886686c5d9ac7484094716cf0c581087d6499193a799a3ed6345	0.25%	147 MiB / 1.954 GiB	7.35%	648 B /
	648 B	0 B / 0 B	24	

Docker inspect <container>

docker inspect command

```
$ docker inspect --help
Usage: docker inspect [OPTIONS] CONTAINER|IMAGE|TASK
Return low-level information on a container, image or task
-f, --format      Format the output using the given go template
--help            Print usage
-s, --size        Display total file sizes if the type is container
--type           Return JSON for specified type, (e.g image, container or task)
```

docker inspect command on a single container

```
$ docker inspect 506fe8cf11bb
[
  {
    "Id": "506fe8cf11bbbb74bf9db26f6561f218ec437cc2b8267d214affb31493e4c251",
    "Created": "2016-07-11T15:51:34.148392848Z",
```

Docker inspect <container>

Try it!

In your Lab Environment, execute the following command:

```
$ docker inspect <container_id>
```

The console will display the configuration details of the container

```
[  
  {  
    "Id":  
      "8404dd7fcb18ba49e6d839bf87ce877297566cfac6ca9a6f0c93846e2bf5e366",  
    "Created": "2016-09-09T14:46:28.911819016Z",  
    "Path": "top",
```

Start, Stop & Remove Containers

Manage Containers

The following tools allow management of containers

docker stop <container>

- Stop a container

docker start <container>

- Start a container

docker rm <container>

- Remove a container

docker stop <container>

Stop a Docker container

```
$ docker stop --help
```

```
Usage: docker stop [OPTIONS] CONTAINER [CONTAINER...]
```

```
Stop one or more running containers
```

```
Options:
```

```
    --help      Print usage
    -t, --time int    Seconds to wait for stop before killing it (default
                      10)
```

docker stop <container>

Try it!

In your Lab Environment, execute the following command:

```
$ docker stop <container_id>
```

This will stop the container. The console will acknowledge with the Container ID

```
7ed45bb572d0
```

docker start <container>

Start a Docker container

```
$ docker start --help
```

```
Usage: docker start [OPTIONS] CONTAINER [CONTAINER...]
```

Start one or more stopped containers

Options:

-a, --attach	Attach STDOUT/STDERR and forward signals
--detach-keys string	Override the key sequence for detaching a container
--help	Print usage
-i, --interactive	Attach container's STDIN

docker start <container>

Try it!

In your Lab Environment, execute the following command:

```
$ docker start <container_id>
```

This will start the container.

The console will acknowledge with the Container ID.

```
7ed45bb572d0
```

docker rm <container>

Remove a Docker container

```
$ docker rm --help
```

```
Usage: docker rm [OPTIONS] CONTAINER [CONTAINER...]
```

```
Remove one or more containers
```

Options:

```
-f, --force      Force the removal of a running container (uses  
SIGKILL)
```

```
--help          Print usage
```

```
-l, --link       Remove the specified link
```

```
-v, --volumes    Remove the volumes associated with the container
```

docker rm <container>

Try it!

In your Lab Environment, execute the following command:

```
$ docker rm <container_id>
```

This will remove the container.

The console will acknowledge with the Container ID.

```
7ed45bb572d0
```

But what if I wanted to delete ALL containers?
In your Lab Environment, execute the following
command: `$ docker rm $(docker ps -a -q)`

This will remove the container. The console will
acknowledge with all the Container IDs.

```
1028e737f3ab
b72a77b7a22f
f774f8618226
3fa9d2d4f5ad
8f5e737d15bf
```

Use with Caution



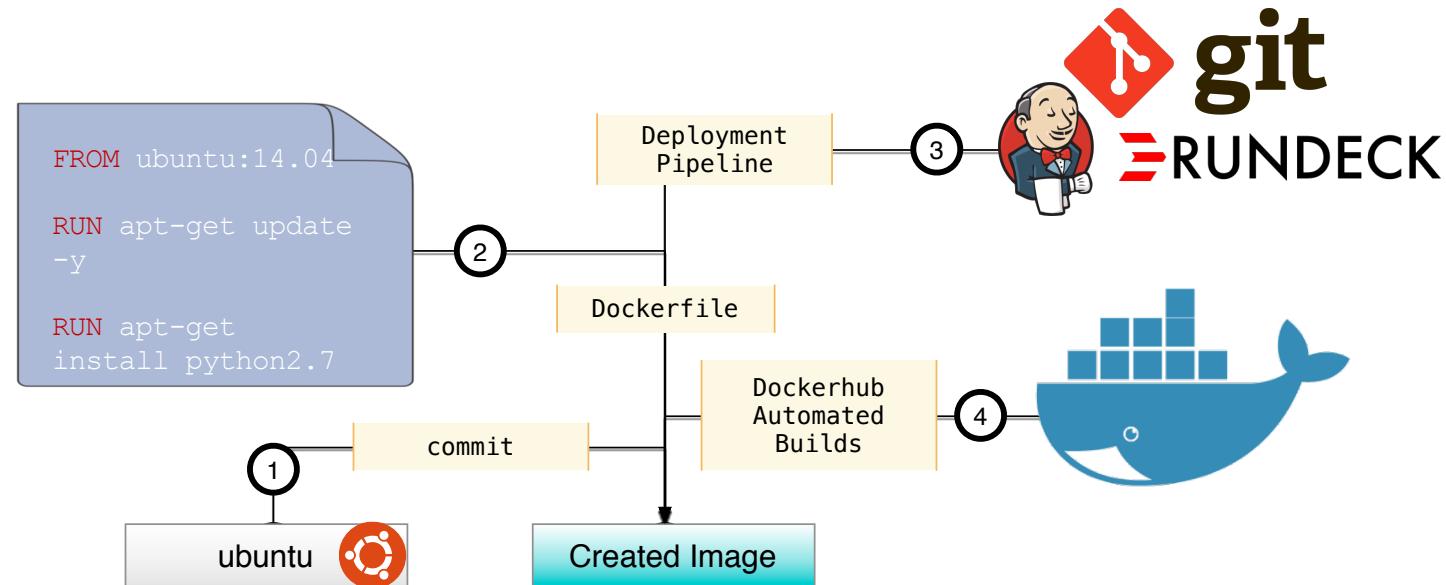
Questions

Docker Images

Methods to Create an Image

Images can be created from:

1. The command-line
2. An image from a Dockerfile
3. An image from a deployment pipeline
4. An image from Docker automated builds



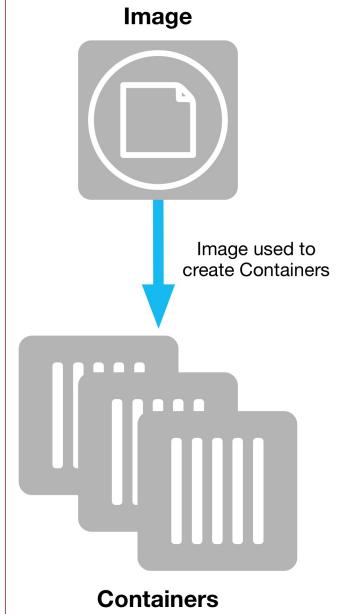
Anatomy of Docker Image

Docker Image Terminology



Hierarchy of files, with meta-data for how to create/run a container

- Read only template used to create containers
- Can be exported or modified to new images
- Created manually or through automated processes
- Stored in a Registry (Docker Hub, Docker Trusted Registry, etc.)



Docker Image Terminology

Image

Union
Filesystem

UnionFS – Used by Docker to layer images

- Not a distributed File System

Dockerfile

Container

Docker Image Terminology

Image

Union
Filesystem

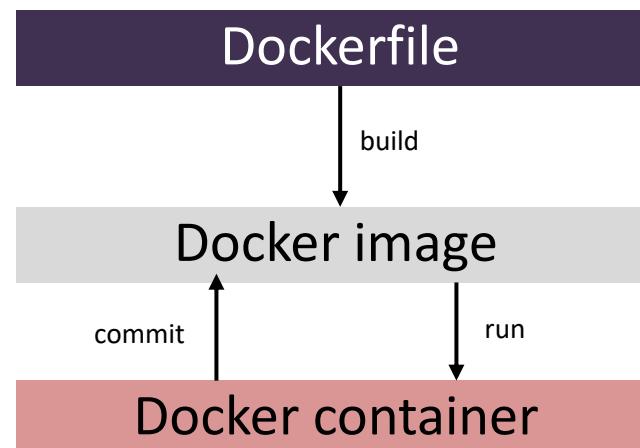
Dockerfile

Container

Configuration File (script) for creating images.

Defines:

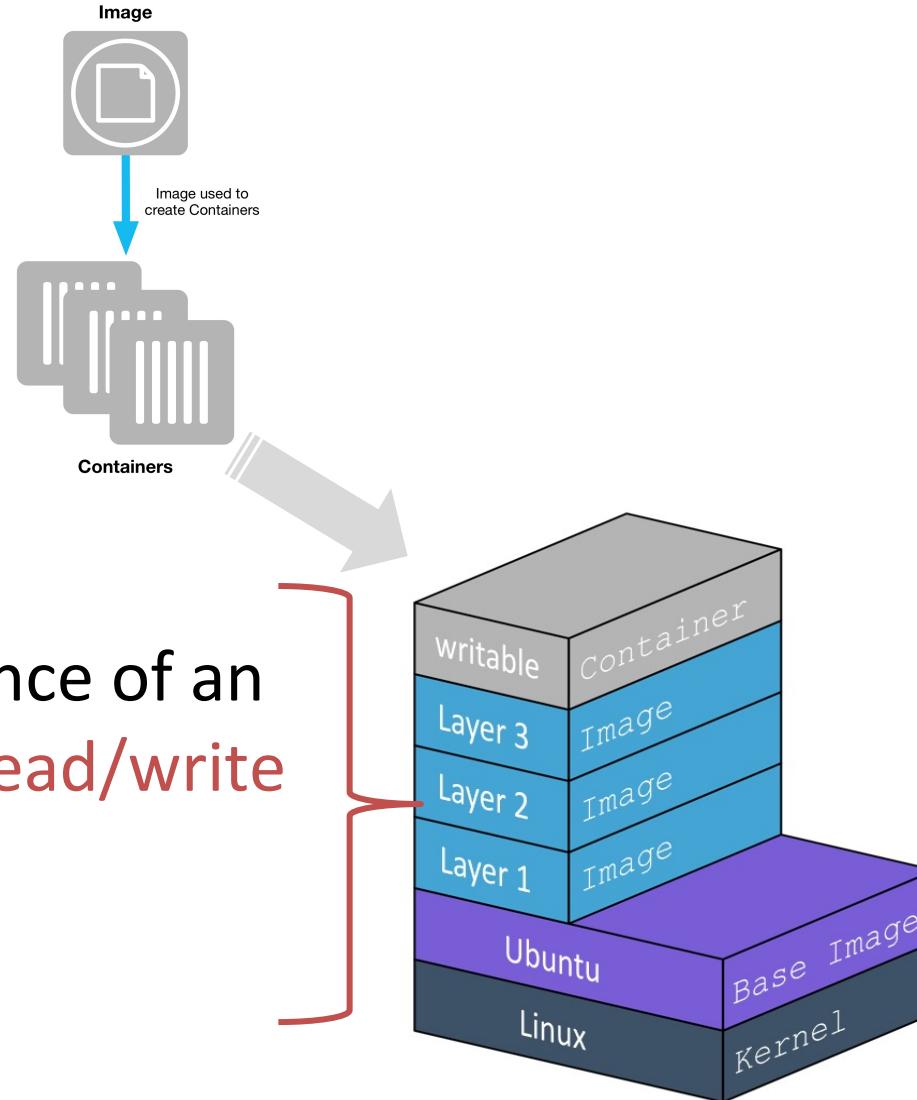
- Existing image to be the starting point
- Set of instructions to augment that image (each of which results in a new layer of the file system)
- Meta-data such as ports exposed
- The command to execute when the image is run



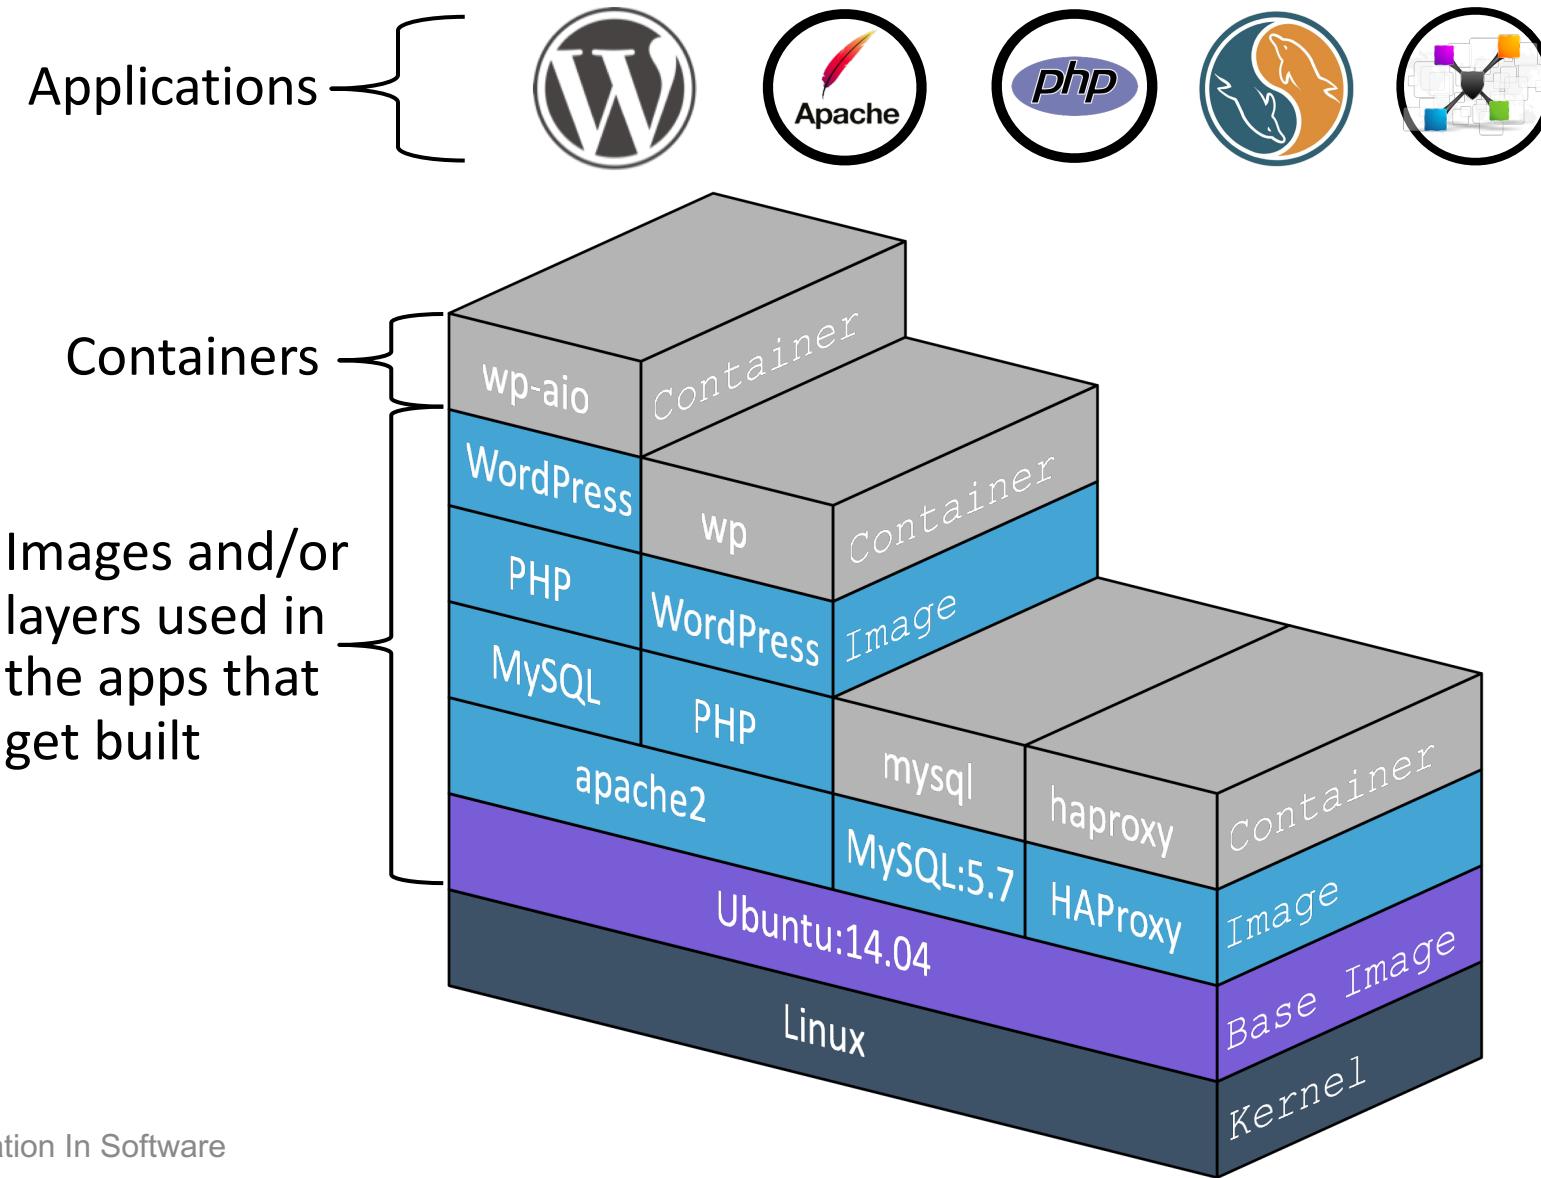
Docker Image Terminology



Runtime instance of an
image **plus a read/write
layer**



Applications, Containers, and Images



Union Files System

Image/Container Interactions

Pull an image from a Docker Registry

Run a container from an image

Add a file to a running container –

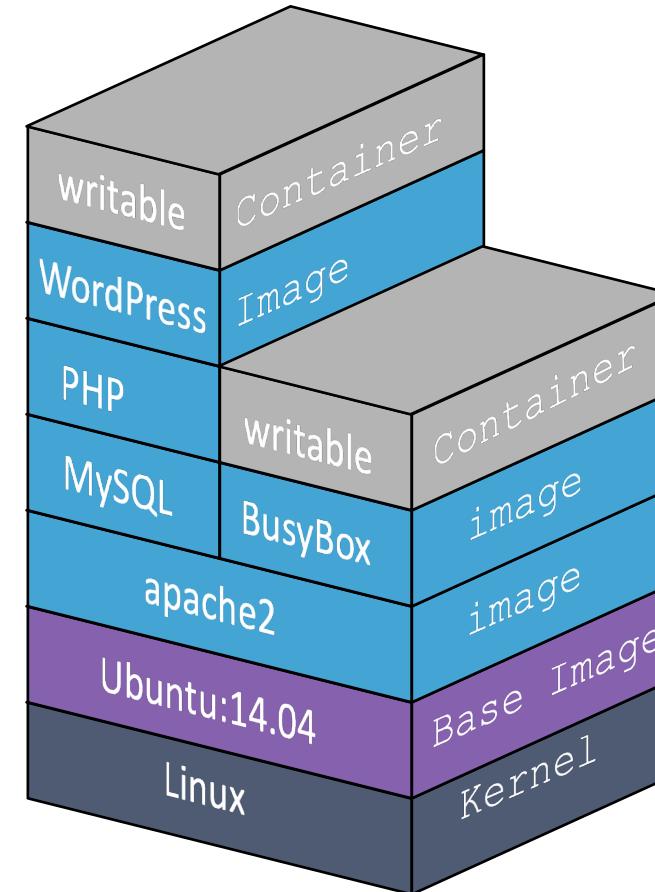
Example, the image requires an additional file called index.html

Change to a running container –

Example, the image requires an update of the index.php file

Delete files from a running container –

All containers share the same host kernel



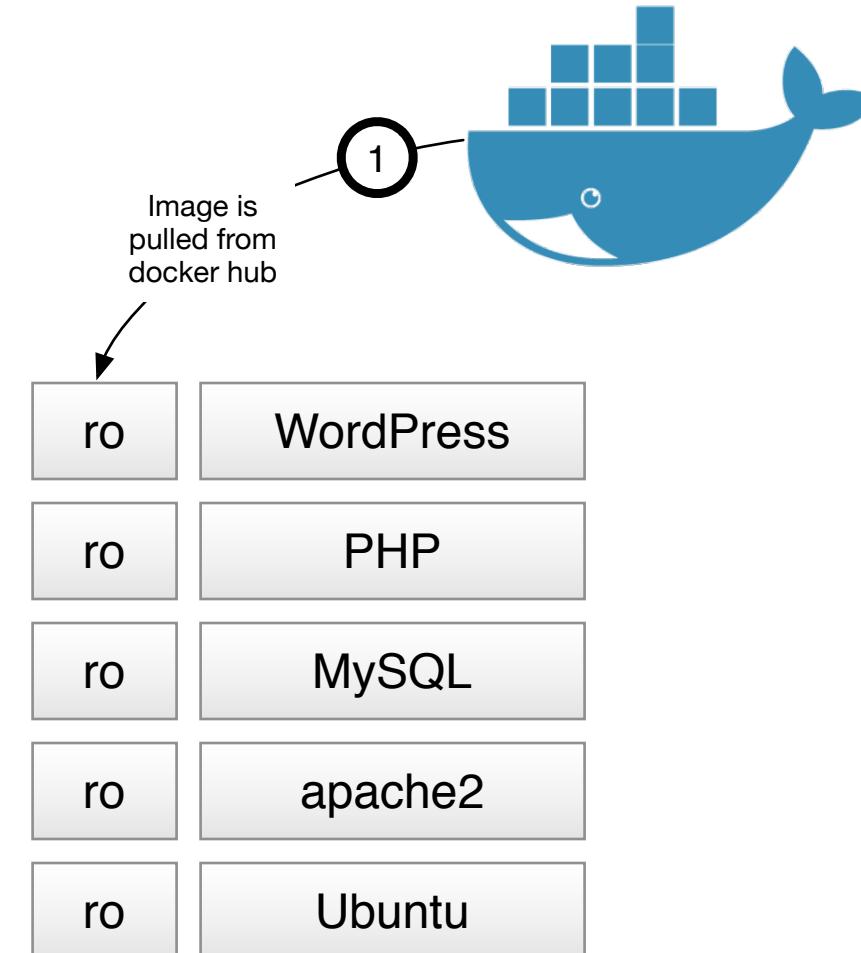
Pull the WordPress AIO

The WordPress All-In-One
image is pulled from
Docker Hub

```
$docker inspect wordpress-aio
```

Docker images are:

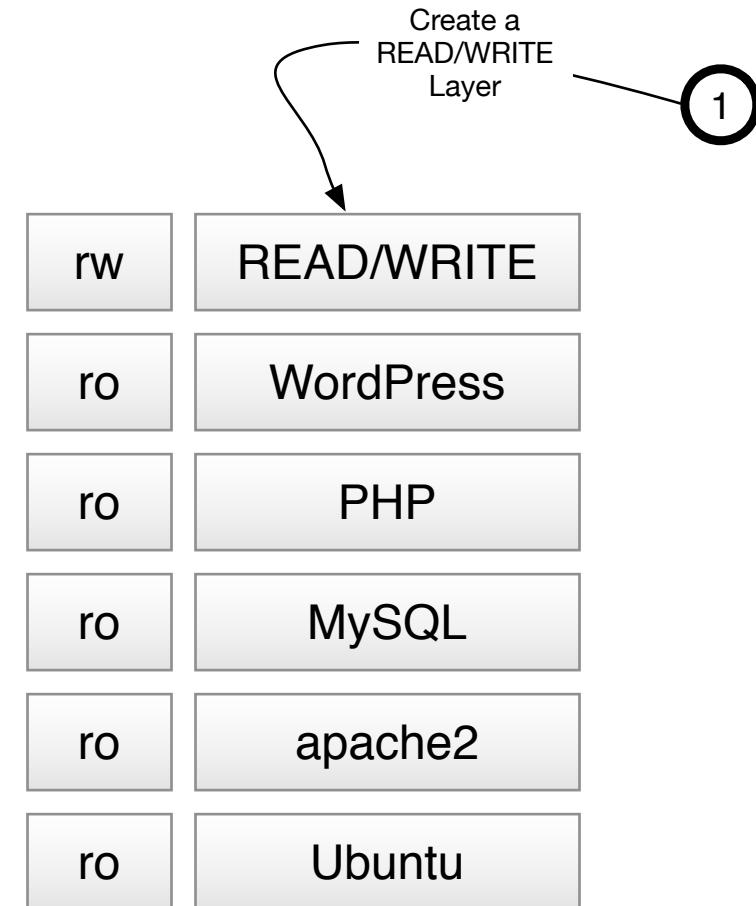
- Pulled or created
- Stored in a local image repo
- READ-ONLY
- The basis of all containers



Run the WordPress AIO

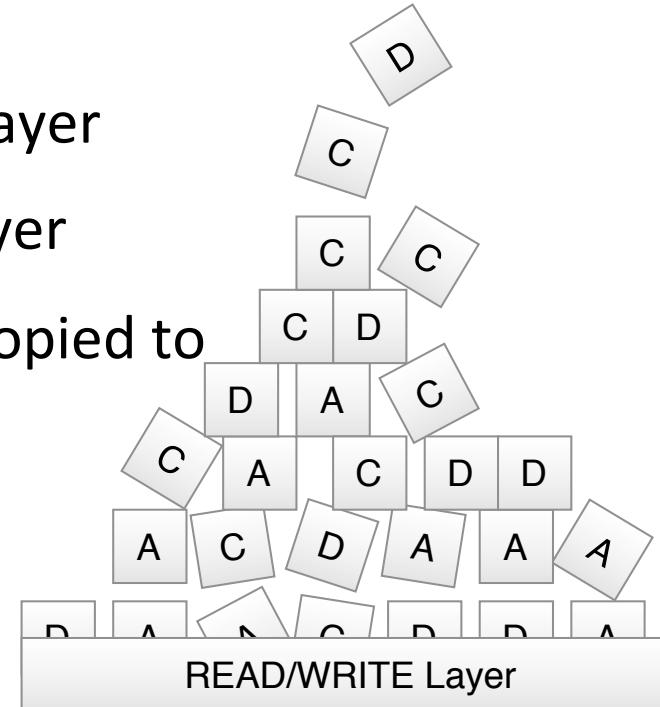
A container is made from the previously pulled WordPress AIO image

1. `$docker run wordpress-aio`
2. Container created
3. READ/WRITE layer created
4. All **ADD, CHANGE, DELETE** are committed to the READ/WRITE Layer



Key Takeaways

- **Images** are:
 - Highly portable and readily available
 - Reusable for many deployments
- **ADD** – ADD DATA to READ/WRITE Layer
- **CHANGE** – ADD DATA to READ/WRITE Layer
- **DELETE** – ADD DATA to READ/WRITE Layer
- **ALL ADDs, CHANGEs, and DELETEs** are copied to the READ/WRITE Layer

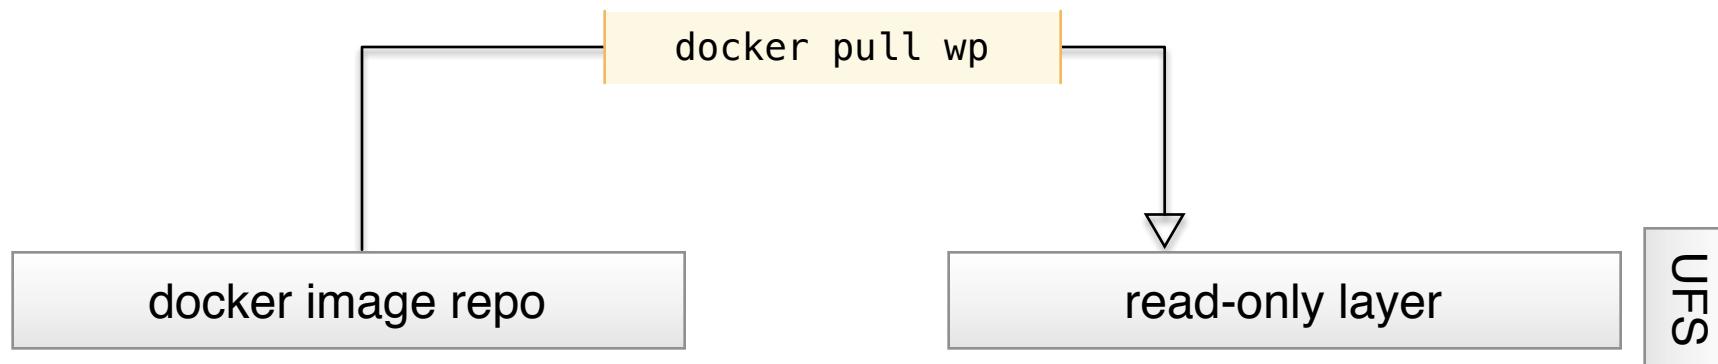


Commits on Containers

Union Filesystem

docker pull– The image is pulled from a Docker registry

```
$ docker image pull <image>
```

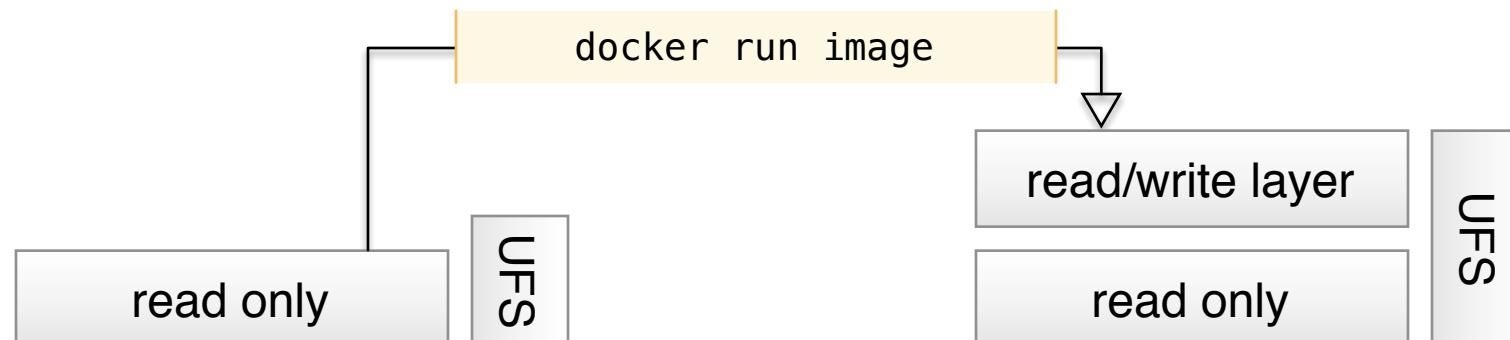


Union Filesystem

RUN – Creates a read/write layer for the container to use

ADDS, CHANGES, DELETES are written to the read/write layer,
the COPY ON WRITE layer.

```
$ docker container run --options <image>
```



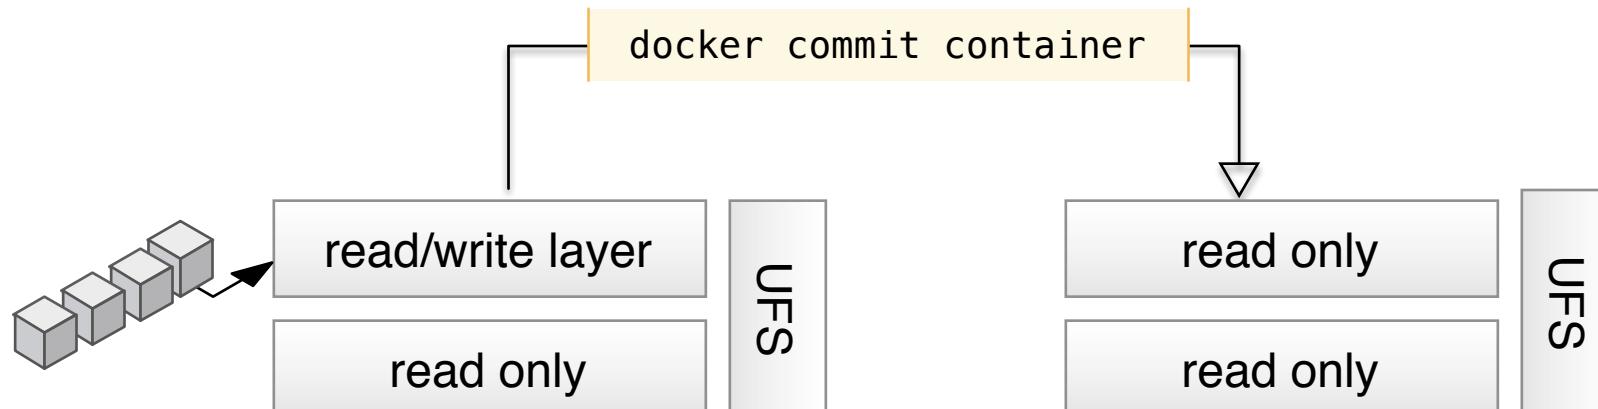
Union Filesystem

Container is committed

- When a container is committed, the READ/WRITE layer becomes a READ ONLY layer

```
docker commit -m "comment" -a "author" demo  
author/demo:new
```

```
$ docker commit -m "comment" -a "author" wp <repo>/wp:tag
```

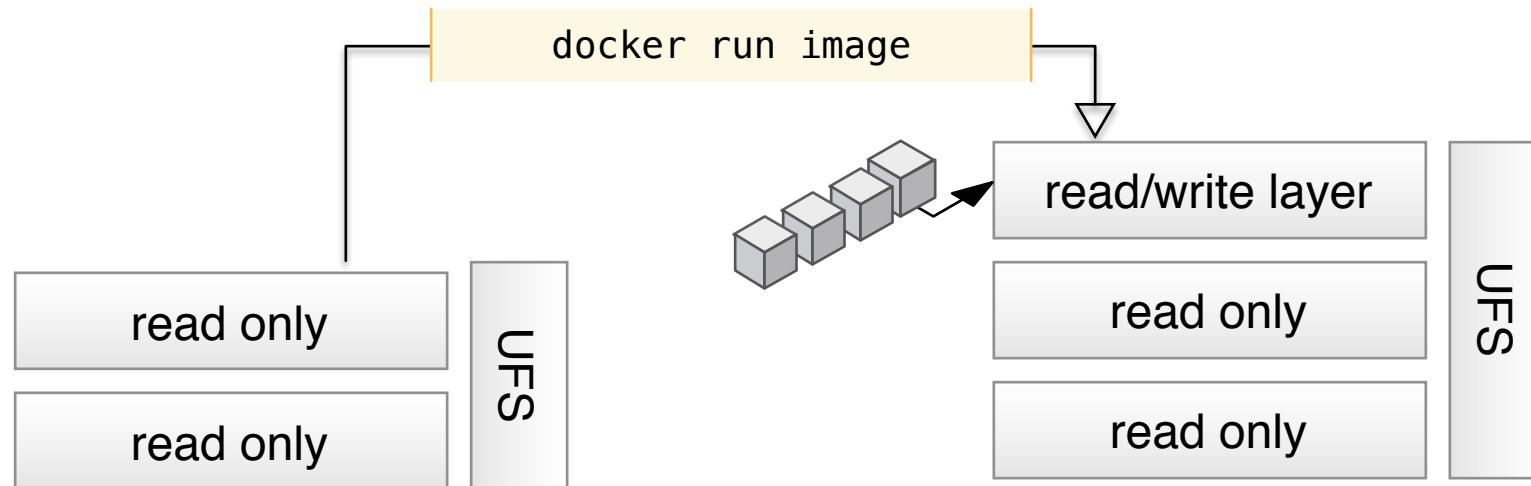


Union Filesystem

A Container run from the newly created image

- The UFS READ/WRITE READ-ONLY process starts again
- A new READ/WRITE layer is created
- Only when changes are made to to R/W UnionFS is data written

```
$ docker container run --options <repo>/wp:tag
```



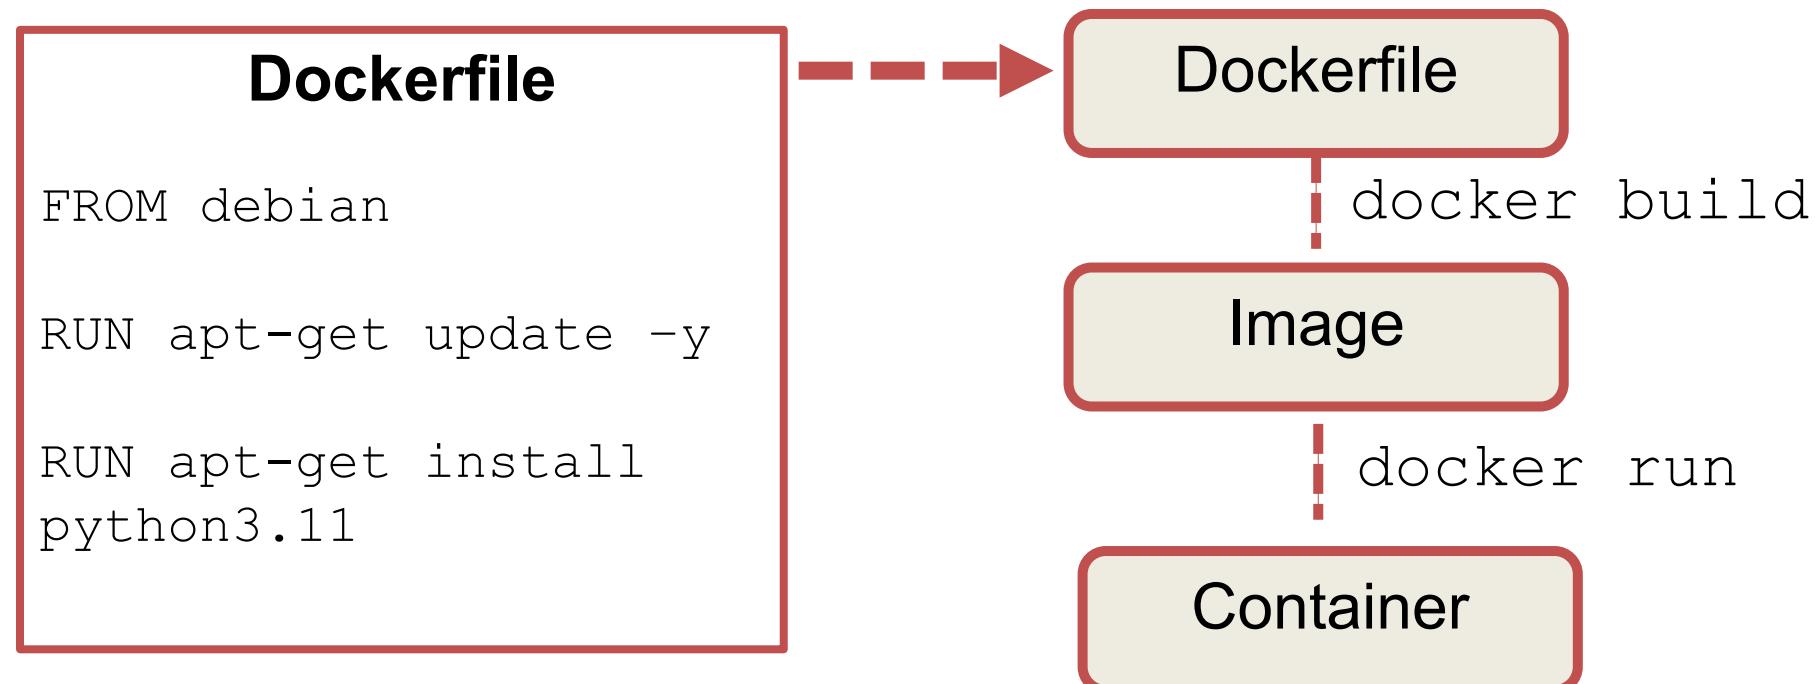
Lab: Build Docker Images (Manually)

Build Docker Images (Dockerfile)

Dockerfile

Docker can build images automatically by reading the instructions from a **Dockerfile**

E.g. dockerfile that installs python 3.11 on top of debian image



Dockerfile Syntax

Dockerfiles have an easily readable format

comments

INSTRUCTIONS

arguments



```
#My first Dockerfile
FROM ubuntu:20.04
RUN apt-get update
```

Lines starting with the “#” are comments

Comments can be placed anywhere

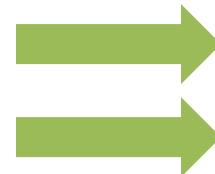
Comments help tell others what is intended

Dockerfile Syntax

comments

INSTRUCTIONS

arguments



```
#My first Dockerfile  
FROM ubuntu:20.04  
RUN apt-get update
```

INSTRUCTIONS should be CAPITALIZED

INSTRUCTIONS are run during image build

The first INSTRUCTION is always FROM

Dockerfile Syntax

comments
INSTRUCTIONS
arguments

arguments are what gets fed to the builder
arguments can be run sequentially in the same INSTRUCTION with an escape

```
#My first Dockerfile
FROM ubuntu:20.04
RUN apt-get update
```

Escape Parser Directive

- form # directive=value
- Windows uses “\” for directories

```
# escape=`
FROM microsoft/windowsservercore
SHELL ["powershell", "-command"]
RUN New-Item -ItemType Directory `

C:\\Example
ADD Execute-MyCmdlet.ps1 c:\\example
```

Dockerfile Example

- Dockerfile takes the latest mongo image
- Updates it
- Installs an entrypoint script to streamline execution

```
FROM mongo:latest
RUN apt-get update && apt-get install -y dos2unix
EXPOSE 27017
COPY docker-entrypoint-init.sh /entrypoint-init.sh
RUN dos2unix /entrypoint-init.sh && \
    apt-get --purge remove -y dos2unix && \
    rm -rf /var/lib/apt/lists/*
RUN chmod ugo+x /entrypoint-init.sh
ENTRYPOINT ["/entrypoint-init.sh"]
CMD ["mongod"]
```

Dockerfile Instructions

Command	Description
#	Comment line
MAINTAINER	Provides name and contact info of image creator
FROM	Tells Docker which base image to build on top of (e.g. centos7)
COPY	Copies a file or directory from the build host into the build container
RUN	Runs a shell command inside the build container
CMD	Provides a default command for the container to run. May be overridden or changed
ADD	Copies new files, directories or remote file URLs
LABEL	Adds metadata to an image

source: <https://docs.docker.com/v1.8/reference/builder/>

Dockerfile Instructions

Command	Description
VOLUME	Exposes any database storage area, configuration storage, or files/folders created by your Docker container
USER	Change to a non-root user
WORKDIR	Sets the working directory for any RUN, CMD, ENTRYPOINT, COPY or ADD instruction
ONBUILD	Executes after the current Dockerfile build completes. ONBUILD executes in any child image derived FROM the current image
EXPOSE	Informs Docker that the container will listen on the specified network ports at runtime
ENTRYPOINT	Allows you to configure a container that will run as an executable
ENV	Sets the environment variable <key> to the value

source: <https://docs.docker.com/v1.8/reference/builder/>

ADD – Not Recommended

```
ADD http://example.com/big.tar.xz /usr/src/things/  
RUN tar -xJf /usr/src/things/big.tar.xz -C  
/usr/src/things  
RUN make -C /usr/src/things all
```

- Image size matters
- Using ADD to fetch packages from remote URLs
INCREASES images sizes
 - Not recommended
- Use curl or wget
 - Gain ability to delete the files you no longer need after they've been extracted
 - Reduces unnecessary layers in your image

Best Practice: Use Curl

```
RUN mkdir -p /usr/src/things \
&& curl -SL http://example.com/big.tar.xz \ | tar
-xJC /usr/src/things \
&& make -C /usr/src/things all
```

- **Use curl instead of ADD**
 - Allows for cleaning up the tar file after it's been extracted and application installed

Build Docker Images (Docker build)

Build a Docker Image (Dockerfile)

1. The `docker build` command is used to "bake" an image
 - Uses Dockerfile
 - Most common flag is `-t` which names the image and (optionally) tags it

build flags	Description
<code>-- pull</code>	Pull new version of the image
<code>-m, --memory</code>	Memory limit
<code>--no-cache</code>	Do not use cache when building the image
<code>-q, --quiet</code>	Suppress the verbose output generated by the containers

Build a Docker Image (Dockerfile)

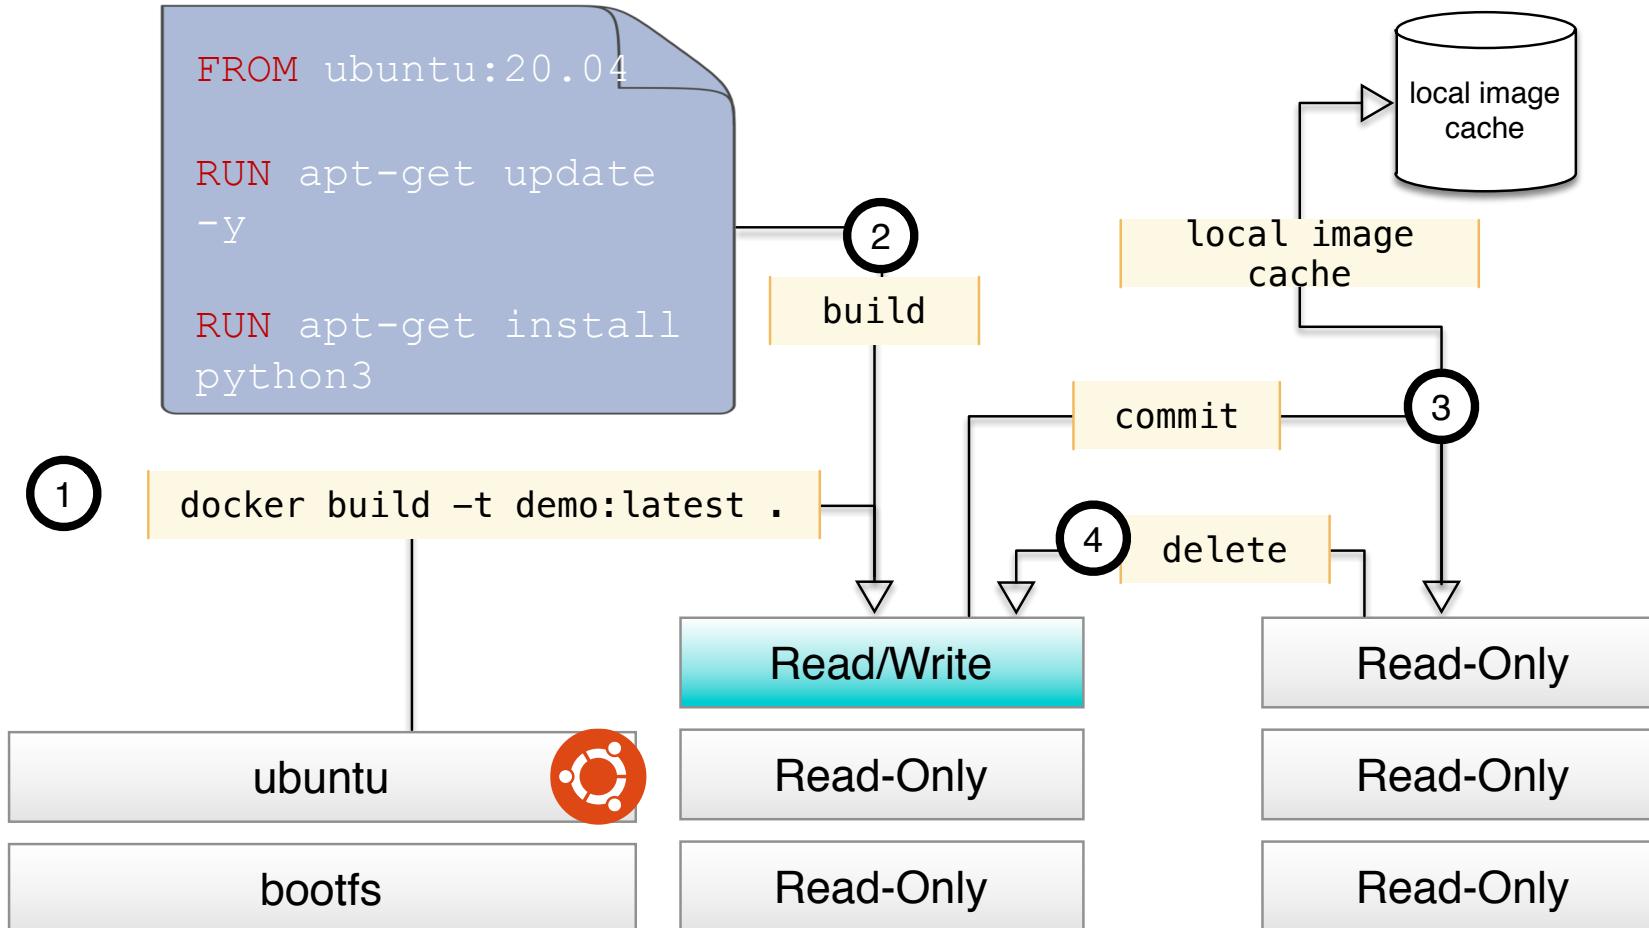


Image Tags

Version tags simplify image identification and tracking

- Common - especially with base operating systems (i.e. `ubuntu:20.04`)
- Allows iteration on known good versions of images, as well as enable us to know exactly what version of our image we are running
- Use multiple tags on an image for easy handling of version for dev, test, etc.
 - Ex. `myimage:dev`, `myimage:1.1`, `myimage:latest`

Docker stop & start signals

- ENTRYPPOINT/CMD starts as PID1
- SIGTERM sent for graceful shutdown
 - --time=30 (30 seconds until SIGKILL sent)
- SIGKILL sent for immediate shutdown
- docker kill --signal=SIGQUIT nginx
- docker kill --signal=SIGWINCH apache

Dockerfile kill signals

- "lazy" syntax
 - "python /myapp/server.py"
 - Shell initialized →
 - /bin/bash = PID1 →
 - server.py
 - Exec syntax
 - ["python", "/myapp/server.py"]
 - server.py = PID1
 - Script used to call app MUST use 'exec python /myapp/server.py'

Lab: Build Docker Images (Dockerfile)

Advanced Docker builds

- A great docker image is small e.g. < 20MiB
- The ideal image contains only the application*
- A great image has very few layers*
- The ideal image has 1 layer*
- The ideal build is 100% reproducible
- The ideal build is as fast as possible



*except when it doesn't

Q: When does performance matter?

ALWAYS

Q: When does performance not matter?

NEVER

Typical order of operations

1. Correctness and Integrity
2. Readability and Usability
3. Performance

That said, let's take performance where it's free

Why care?

Smaller Images Mean:

- Faster builds
- Faster pulls
- Faster pushes
- Faster development
- Faster testing
- Faster deployment
- Faster feedback

Minimal base images

- scratch
- alpine (<5MiB)



Application Base Images

- redis:4-alpine
- mysql:5.7



redis



Stack Base Images

- gcr.io/distroless (python, node, java, dotnet, cc')
- python:3.6-alpine
- php:7.2-alpine3.7
- golang



Layer review

- Images are stacks of layers
- Each Dockerfile line is a layer

```
#Layer n~  
FROM nixos/nix~  
~  
#Layer n+1~  
ENV build=dev~  
~  
#Layer n+2~  
CMD /app/bin/start~  
~  
#Layer n+3~  
COPY app /app~  
~  
#Layer n+4~  
RUN /app/bin/build.sh~  
~
```

Add small layers

Might be small, might be HUGE

```
FROM alpine
|
|   COPY . .
```

Probably smaller

```
FROM alpine
|
|   COPY app /app
```

Independent RUNs at top

RUNs are cacheable

```
FROM alpine
-
RUN dd if=/dev/random of=/tmp/seed bs=1 count=1
-
# More stuff ...
```

Dependent RUNs immediately after the file they depend on

```
#Layer n
FROM nixos/nix

#
#Layer n+1
ENV build=dev

#
#Layer n+2
CMD /app/bin/start

#
#Layer n+3
COPY app /app

#
#Layer n+4
RUN /app/bin/build.sh
```

Chain your RUNs

Chained commands happen in the same layer

```
FROM node:6-alpine
-
CMD ["node", "/www/src/index.js"]
-
COPY package.json /www/
-
RUN apk add --update git && \
... npm --prefix=/www install && \
... apk del git && \
... rm -rf /var/cache/apk/*
-
COPY src/ /www/src
```

Cleanup your RUNs

This build requires git, but the runtime doesn't. Let's cleanup after the build.

```
FROM node:6-alpine
-
CMD ["node", "/www/src/index.js"]
-
COPY package.json /www/
-
RUN apk add --update git && \
    npm --prefix=/www install && \
    apk del git && \
    rm -rf /var/cache/apk/*
-
COPY src/ /www/src
```

Add layers sparingly

Each line should have a good reason to exist

Common extraneous layers:

RUN followed by RUN

MAINTAINER <in+version+control@example.com>

EXPOSE 80

WORKDIR

Sometimes extraneous:

VOLUME /app/data

Order matters

Anti-Example:

```
COPY . /tmp/-
RUN pip install --requirement /tmp/requirements.txt-
-
```

Not bad?:

```
COPY requirements.txt /tmp/-
RUN pip install --requirement /tmp/requirements.txt-
COPY . /tmp/-
-
```

Order matters

Good:

```
COPY LICENSE /tmp/-
COPY requirements.txt /tmp/-
RUN pip install --requirement /tmp/requirements.txt-
COPY *.py /tmp/-
[...]
```

Best:

```
COPY LICENSE /tmp/-
COPY requirements.txt /tmp/-
RUN pip install --requirement /tmp/requirements.txt-
COPY constants.py /tmp/-
COPY main.py /tmp/-
[...]
```

Layers are immutable, nth reminder

Bad*:

```
COPY ..
```

Probably Worse:

```
COPY ..
```

```
RUN rm -rf secrets/
```



.dockerignore - reduce build context

you can use `.dockerignore` in root of build context to exclude files from the build

`.dockerignore` makes these better:

```
COPY ..
```

```
COPY src/ /app/src
```

When simple, explicit is still preferable:

```
COPY foo /app/foo
```

```
COPY bar /app/bar
```

Cacheability vs Layers and Readability

We see a tradeoff between values:
minimizing number of layers and readability
maximizing cacheability

```
COPY *.py /app/-
└
# VS
└
COPY constants.py /app/-
COPY main.py /app/-
```

```
COPY *.py /app/-
└
# VS
└
COPY a.py /app/-
COPY b.py /app/-
COPY c.py /app/-
COPY d.py /app/-
# ...
```

When to not minimize an image

When it impacts usability or readability:

- run a debugger inside the container (copy in the debugger dependencies)
- run a profiler inside the container (copy in the profiler dependencies)
- run a supporting application inside the container (e.g. redis-cli)
- open a shell inside a container

Multi-stage Docker builds

- Introduced in Docker 17.05
- Multiple FROM statements can be used to decompose a docker build into “stages”
- Conditionally build a single stage, or all stages

Common multi-stage patterns

- Create a “build” layer that performs the actual compilation of your application
- Create a “test” layer that contains canned dummy data for your application to consume or expose for testing
- Create a “production” layer that is stripped down to only a lightweight base image (like alpine) and your binary

Benefits of multi-stage builds

Portability

- Utilizing a "build" layer in your Dockerfile allows a docker build to be ported to any build system (Jenkins, TravisCI, CircleCI, etc) that can build a container. No consideration has to be made as to whether the build system "supports" builds with your preferred runtime.

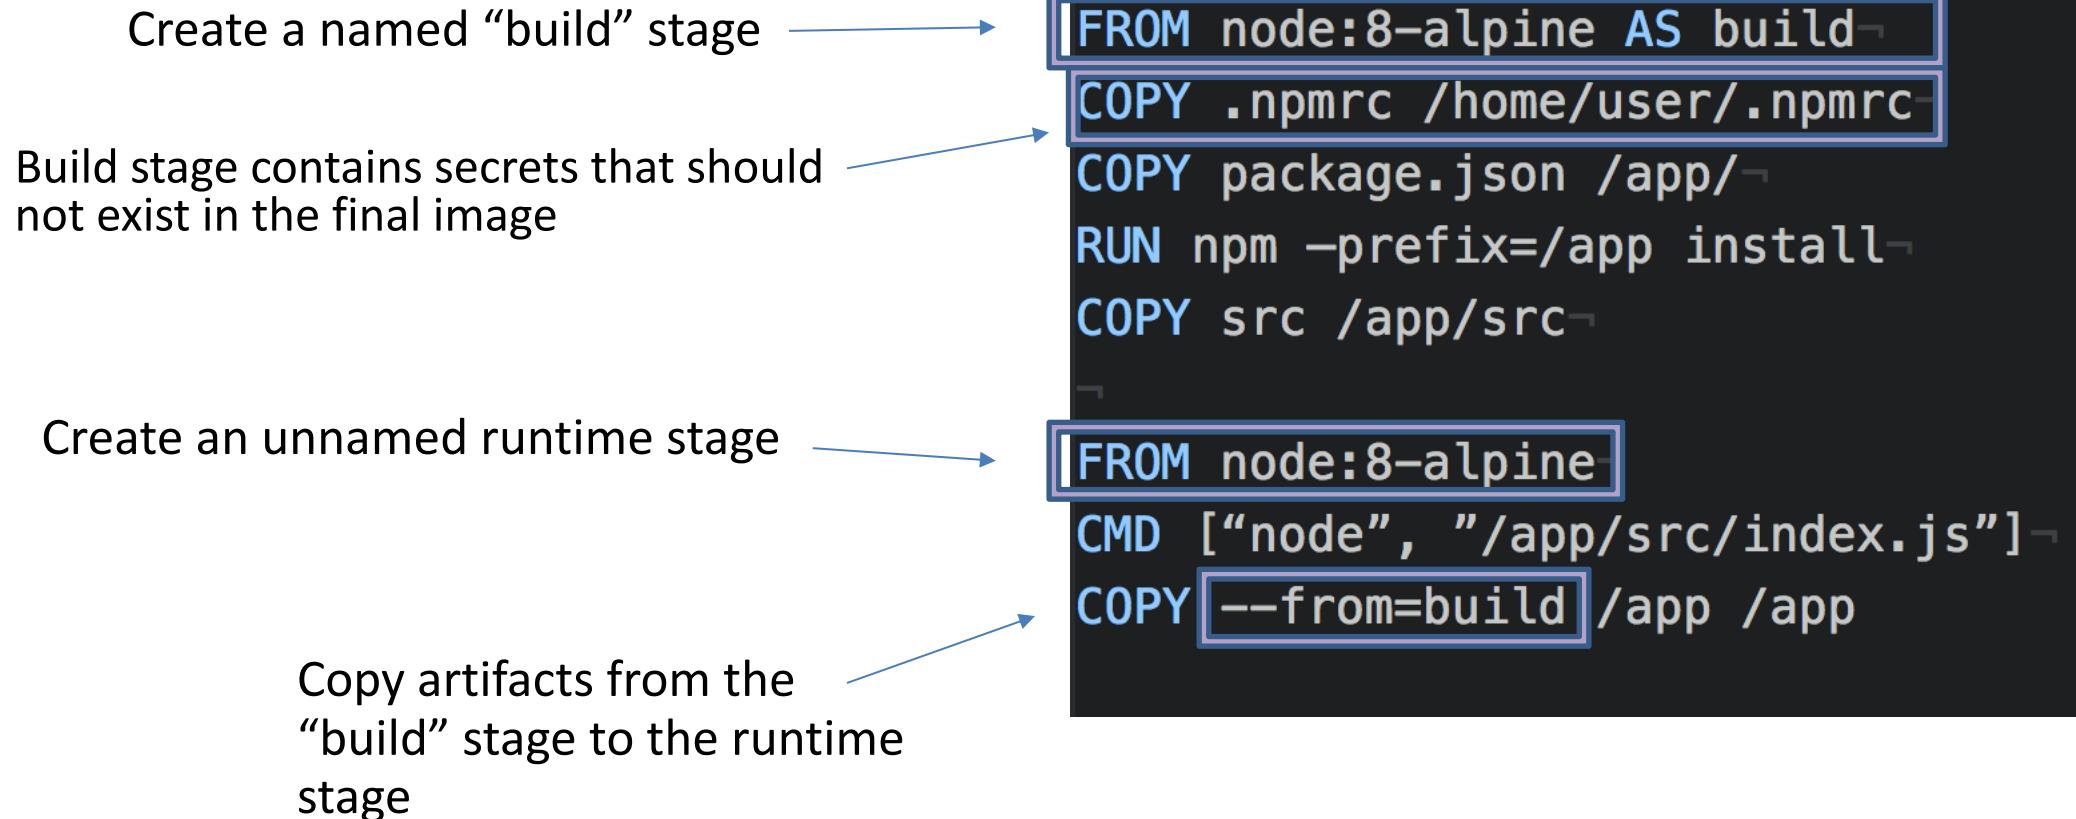
Security

- By only including explicitly what you want in the final image, you decrease your container attack surface

Simplicity

- Removes any requirement to have "test" and "production" Dockerfiles to facilitate testing.
- The number and size of layers in intermediate stages does not factor into the final deployable image.

Multi-stage build example





Questions

Trusting Upstream

FROM nginx:1.13.9

Does project use semver?

If yes are they consistent?

What version of semver?

What happens if upstream breaks my app?

Will tests catch it?



Reliable builds

```
#nginx:1.13.9~  
FROM nginx@sha256:4771d09578c7c6a65299e110b3ee1c0a2592f5ea2618d23e4ffe7a4cab1ce5de~  
COPY index.html /www/~  
~
```

- + Reliable base image for builds
- No automatic upstream updates
- Worse readability



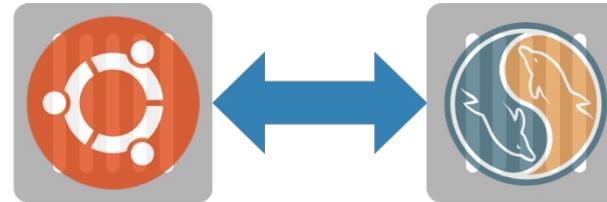
Discussion: Optimize Dockerfiles

Docker Networking

Docker Networking Overview

Docker Engine provides network access for containers that need:

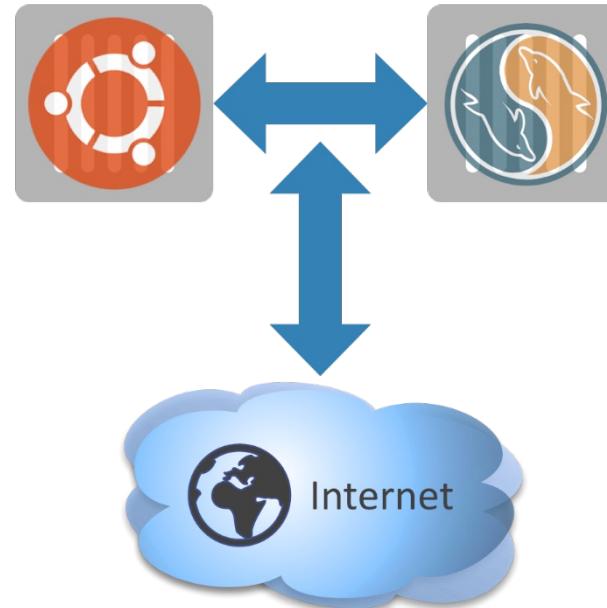
- Other Containers



Docker Networking Overview

Docker Engine provides network access for containers that need:

- Other Containers
- Networks



Docker Networking Overview

Docker Engine provides network access for containers that need:

- Other Containers
- Networks

Accomplished through:

1. IP Address Management (IPAM)
2. Service Port exposure
3. Manual Port Mapping
4. Dynamic Port Mapping

Docker Networking Architecture

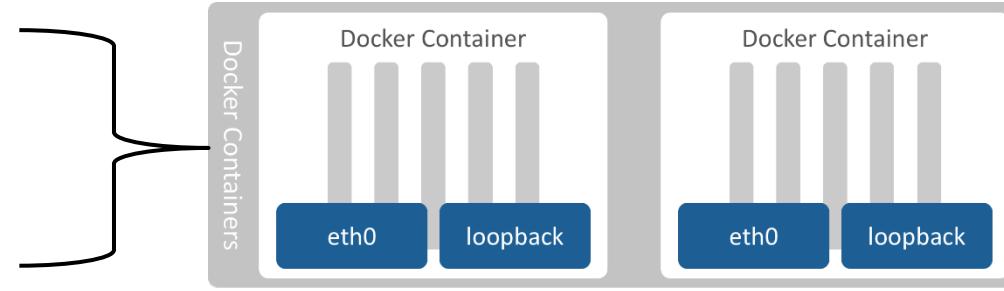
© 2025 by Innovation In Software



Network Interfaces

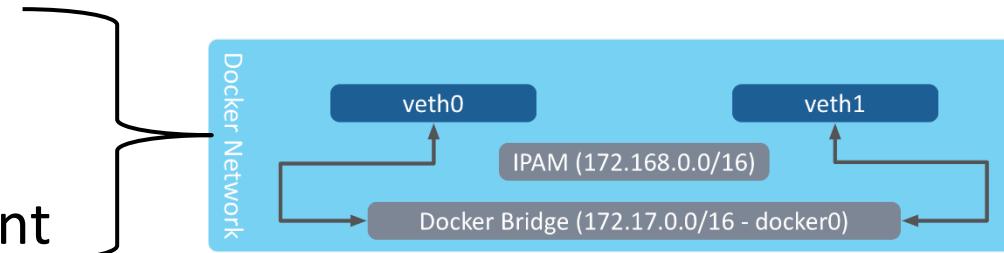
Docker Container

- eth0 Interface
- Loopback Interface



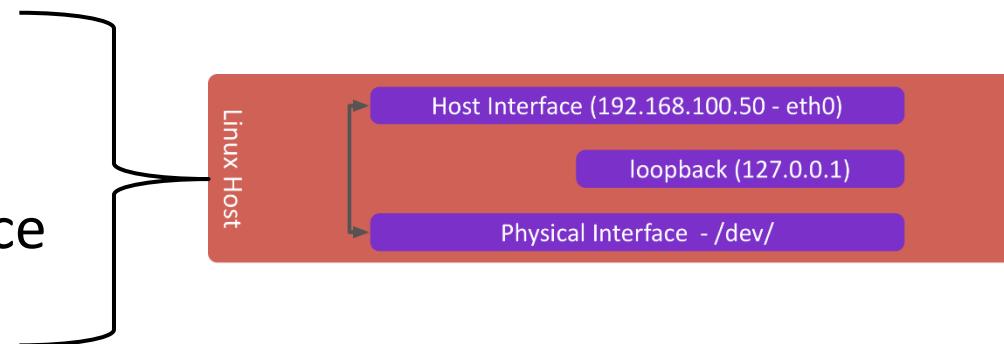
Docker Network

- Docker Bridge
- IP Address Management

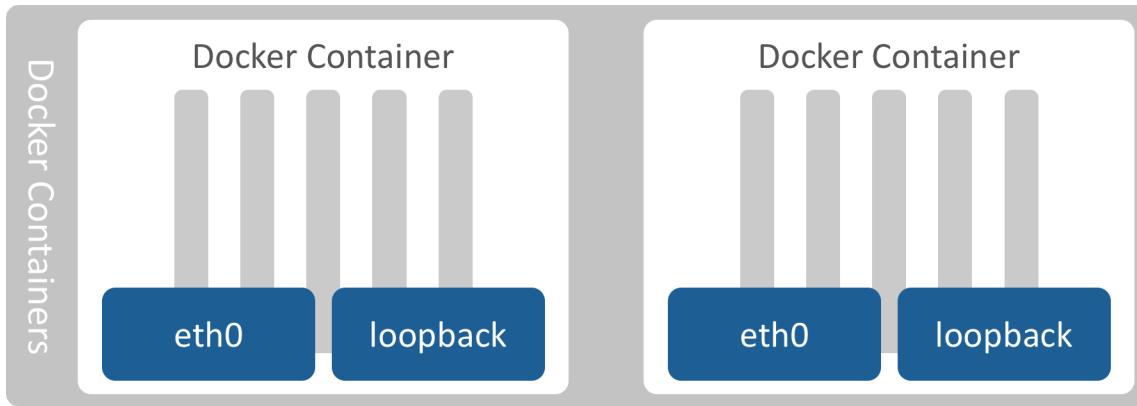


Linux Host Network

- Host Interfaces
- Host Loopback Interface
- Physical Interfaces



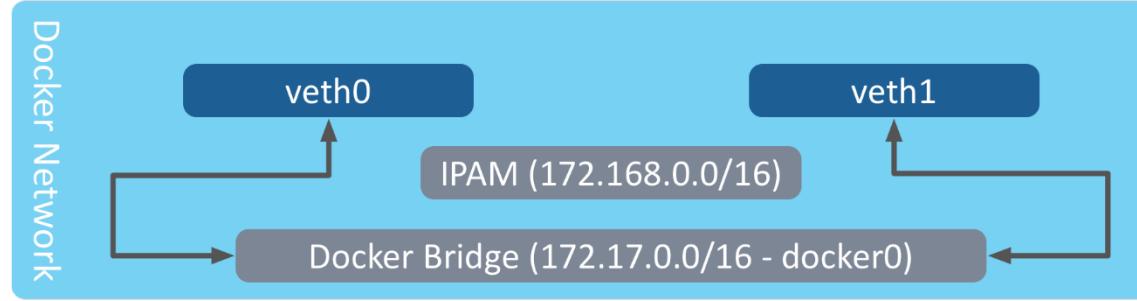
Docker Container



Docker Container

- **Loopback Interface** - used for internal communication
- **eth0 Interface** - used for external communication
- Container TCP Port Exposure

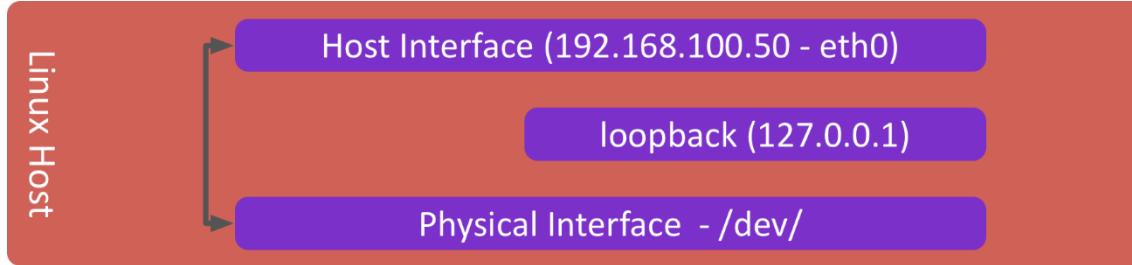
Docker Container



Docker Network

- **Docker Bridge** (e.g. docker0) - the bridge which manages traffic between Docker Engine and the Linux Host
- **veth Interfaces** - interfaces on the bridge to the containers
- **IP Address Management** - assigns IP addresses to container that require them
- Container to Host TCP Port Mapping

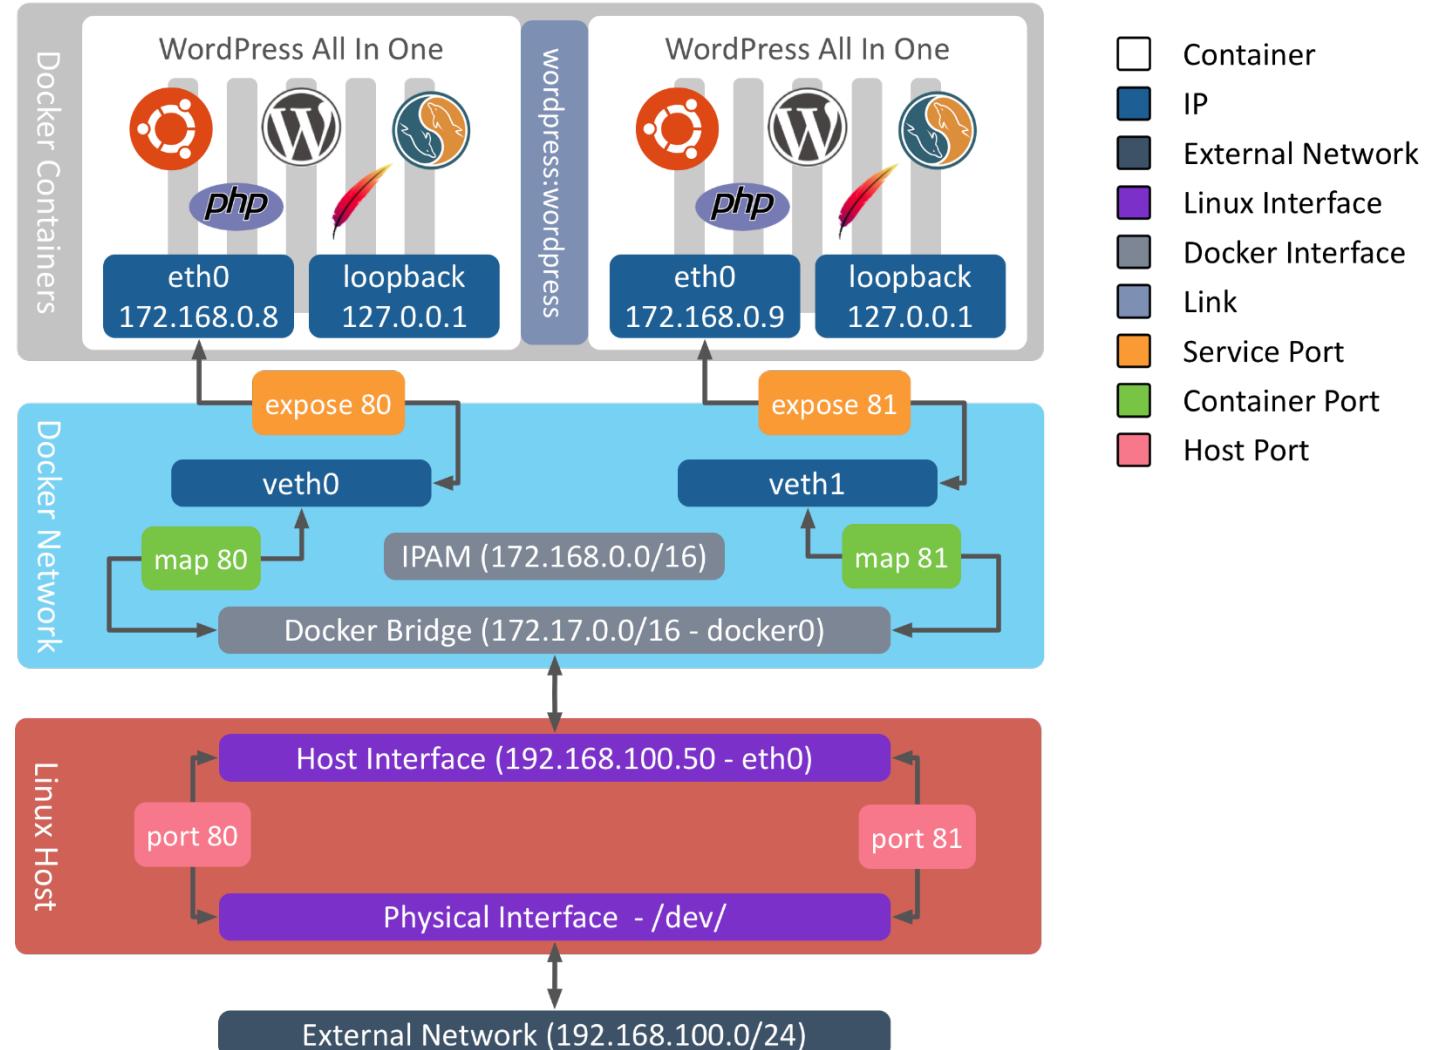
Docker Container



Linux Host Network

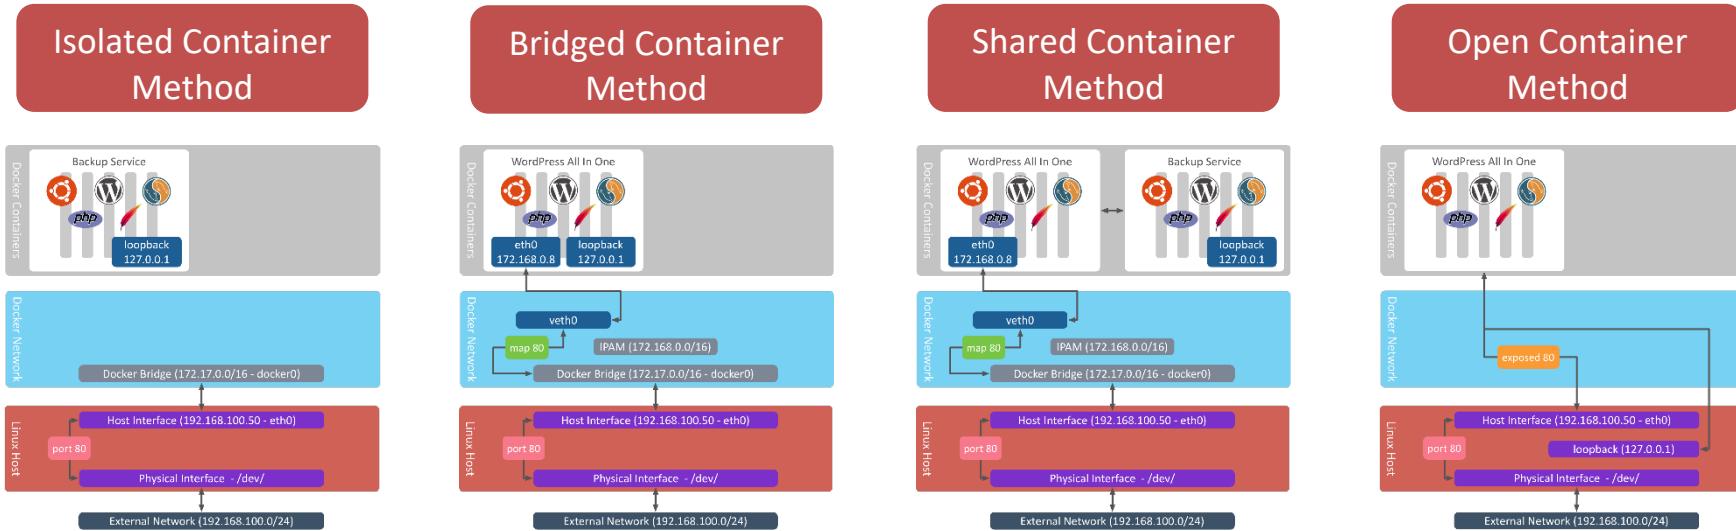
- **Host Loopback Interface** - host internal communication
- **Host Interfaces** - internal and external communication
- **Physical Interfaces** - sends/receives external communication to/from the outside world
- Linux Host TCP Ports

Docker Networking Architecture



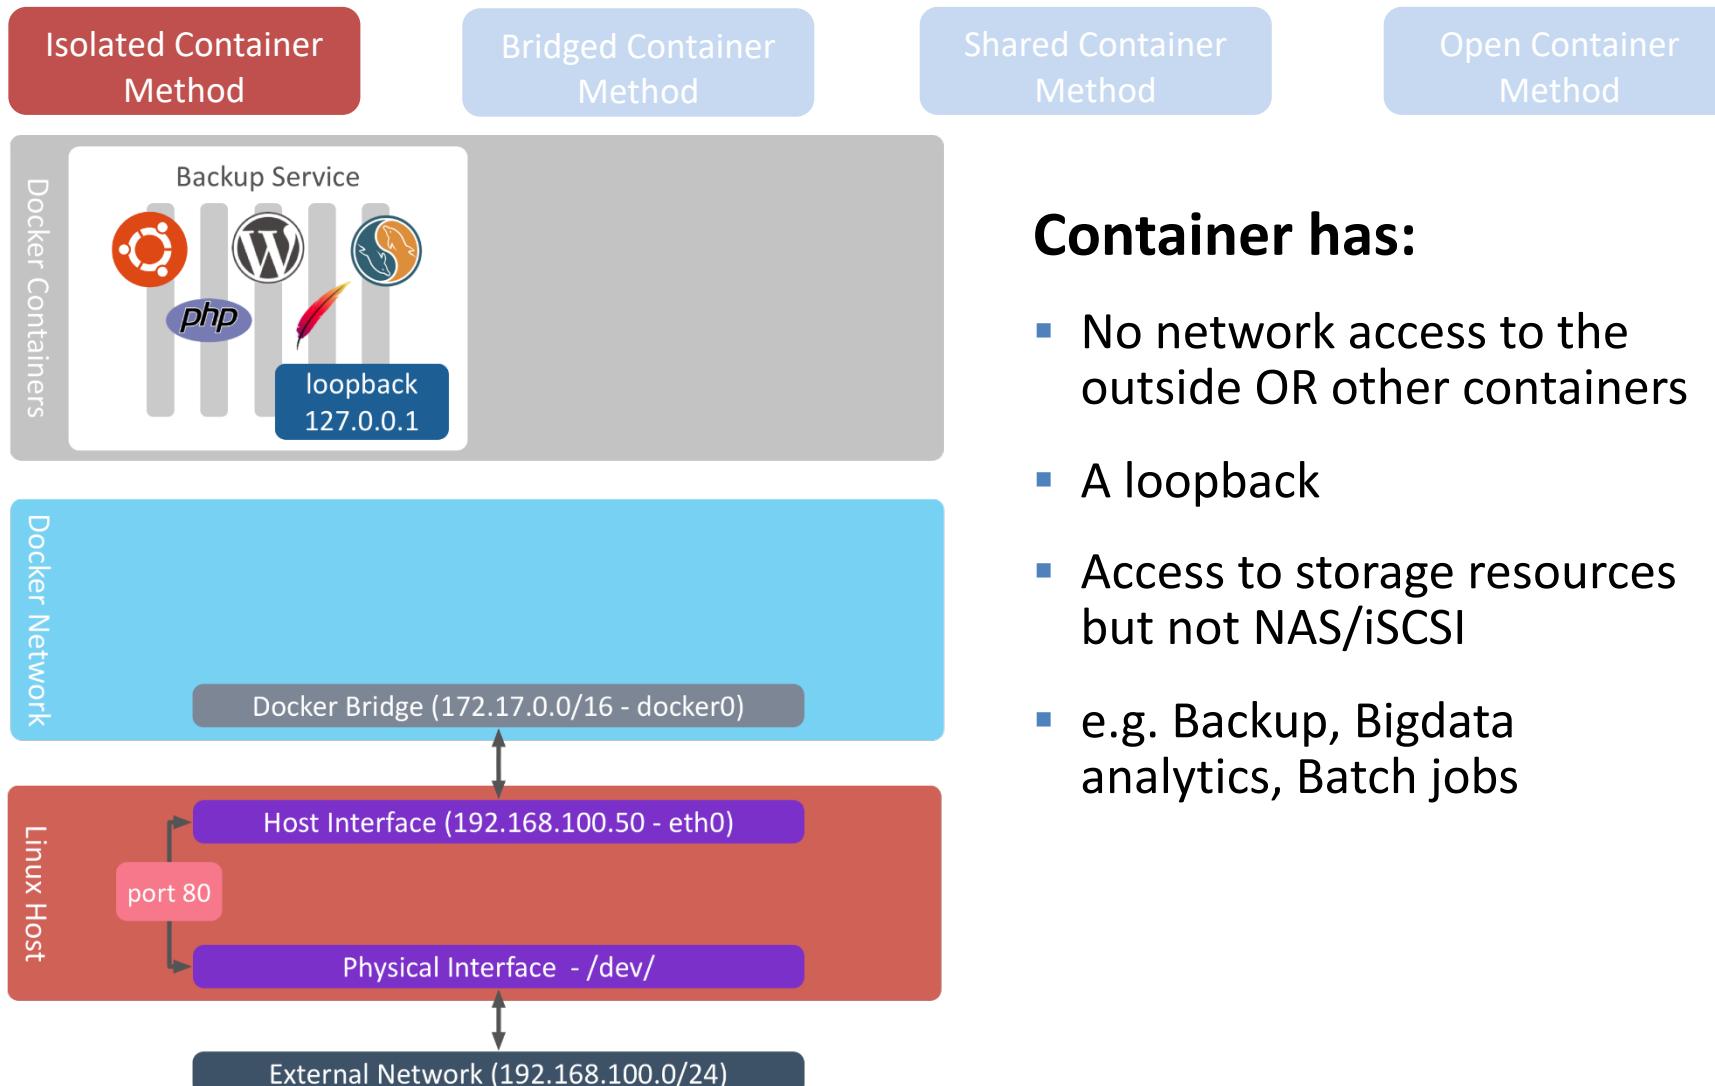
Container Network Deployment Methods

Container Network Models

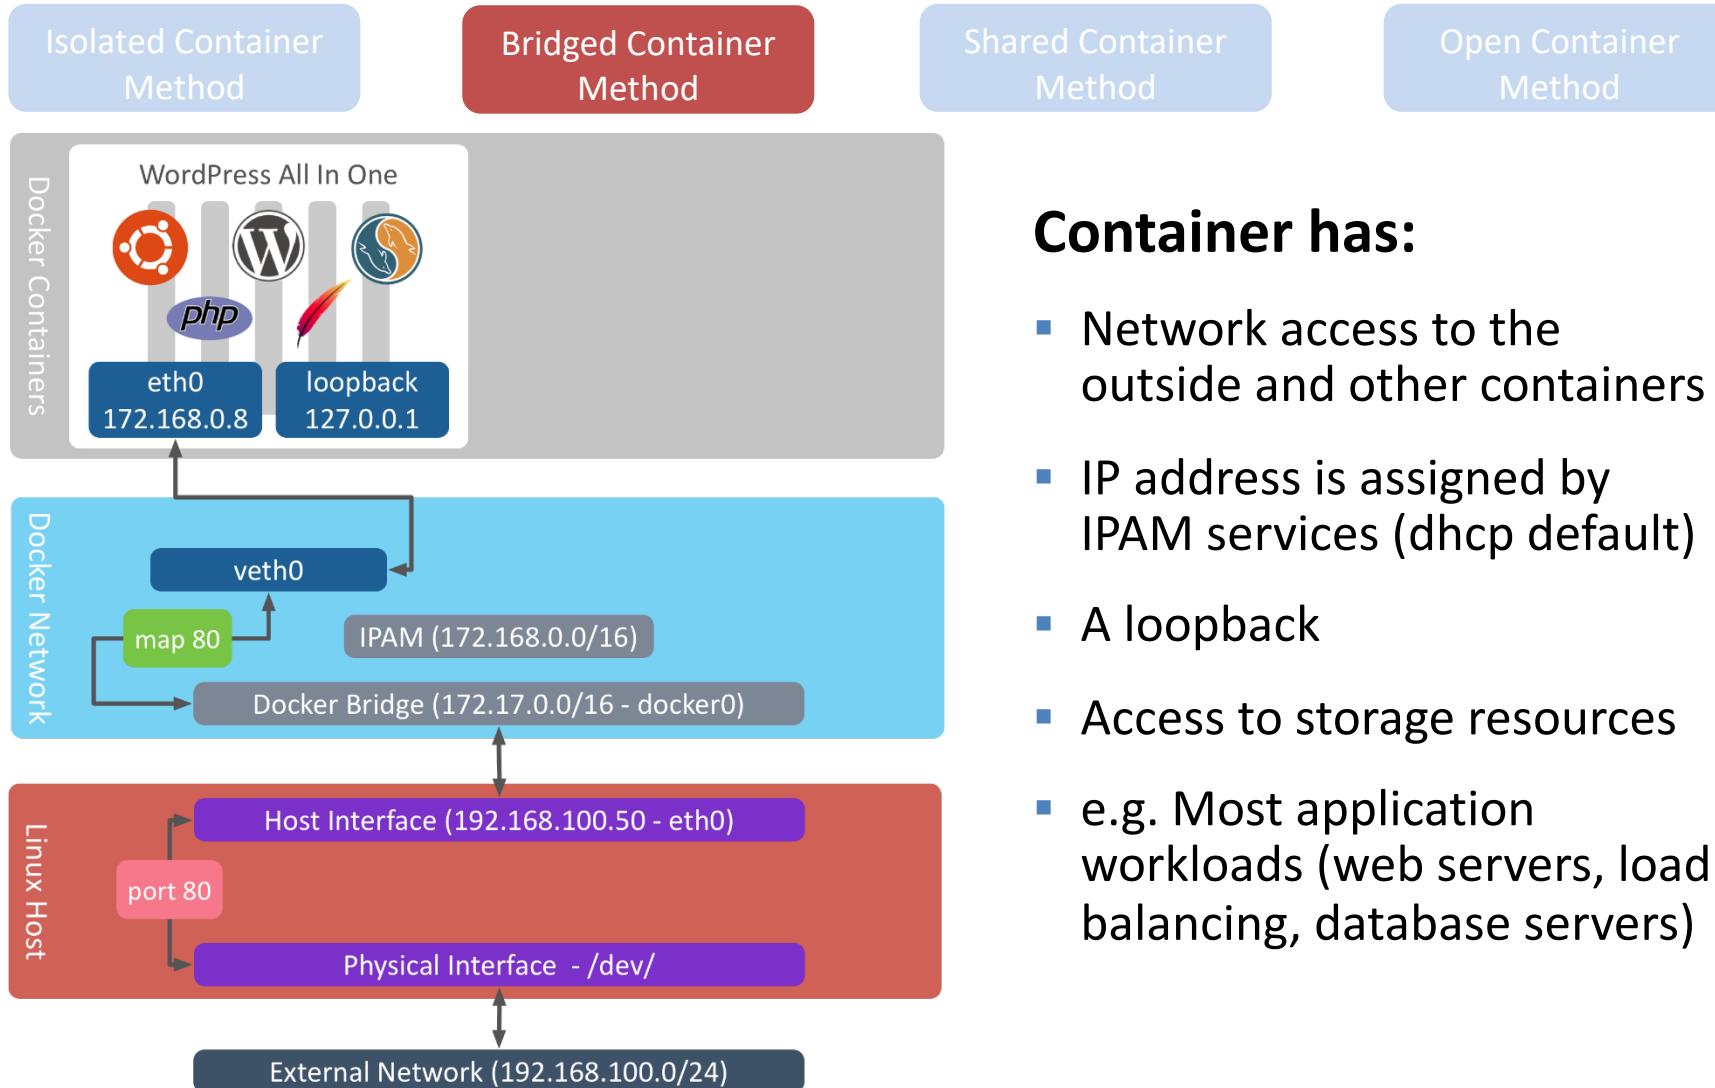


- Seldom used
 - Provides the most security
 - Has zero access to network resources
 - Has no eth0 but has a loopback
- Most commonly used
 - Has limited access to network resources
 - Has an eth0 and a loopback
- Least used
 - Has limited and no access to network resources
 - Each container has either a loopback or eth0
- Used as a last resort and with caution
 - Unrestricted access to host network
 - Uses host network and loopback

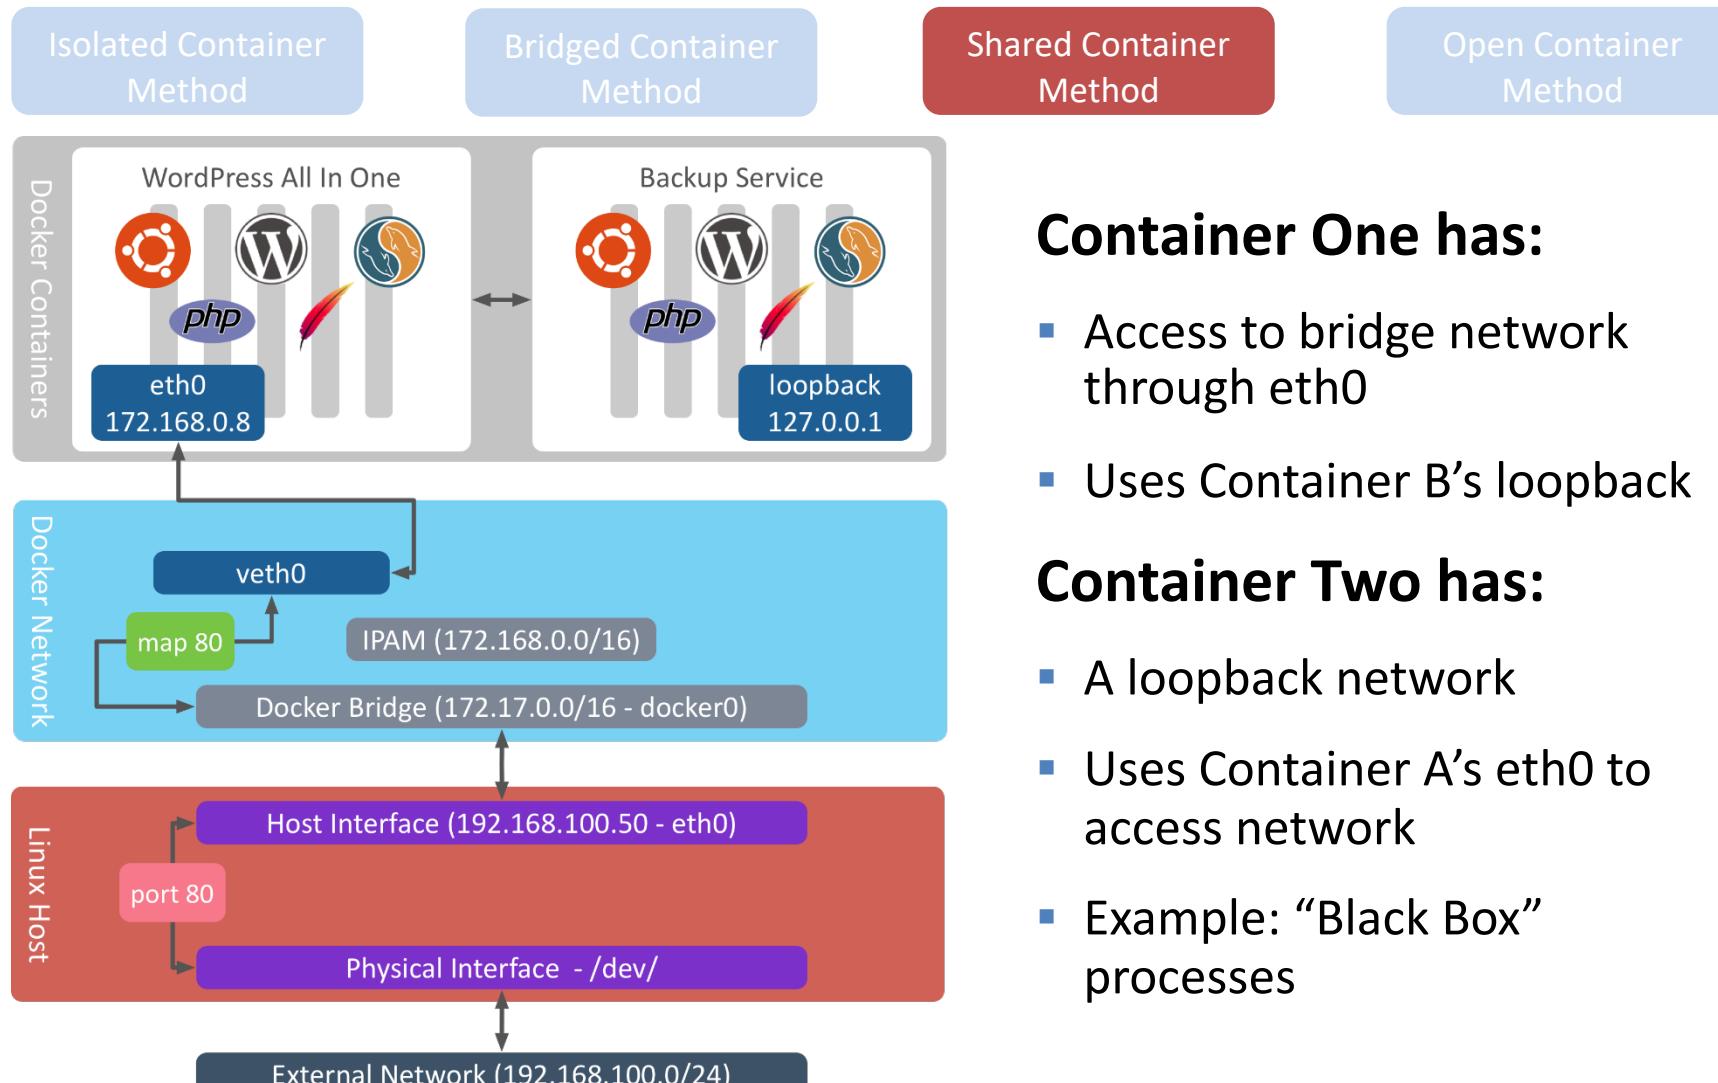
Isolated Container Method



Bridged Container Method (Default)



Shared Container Method



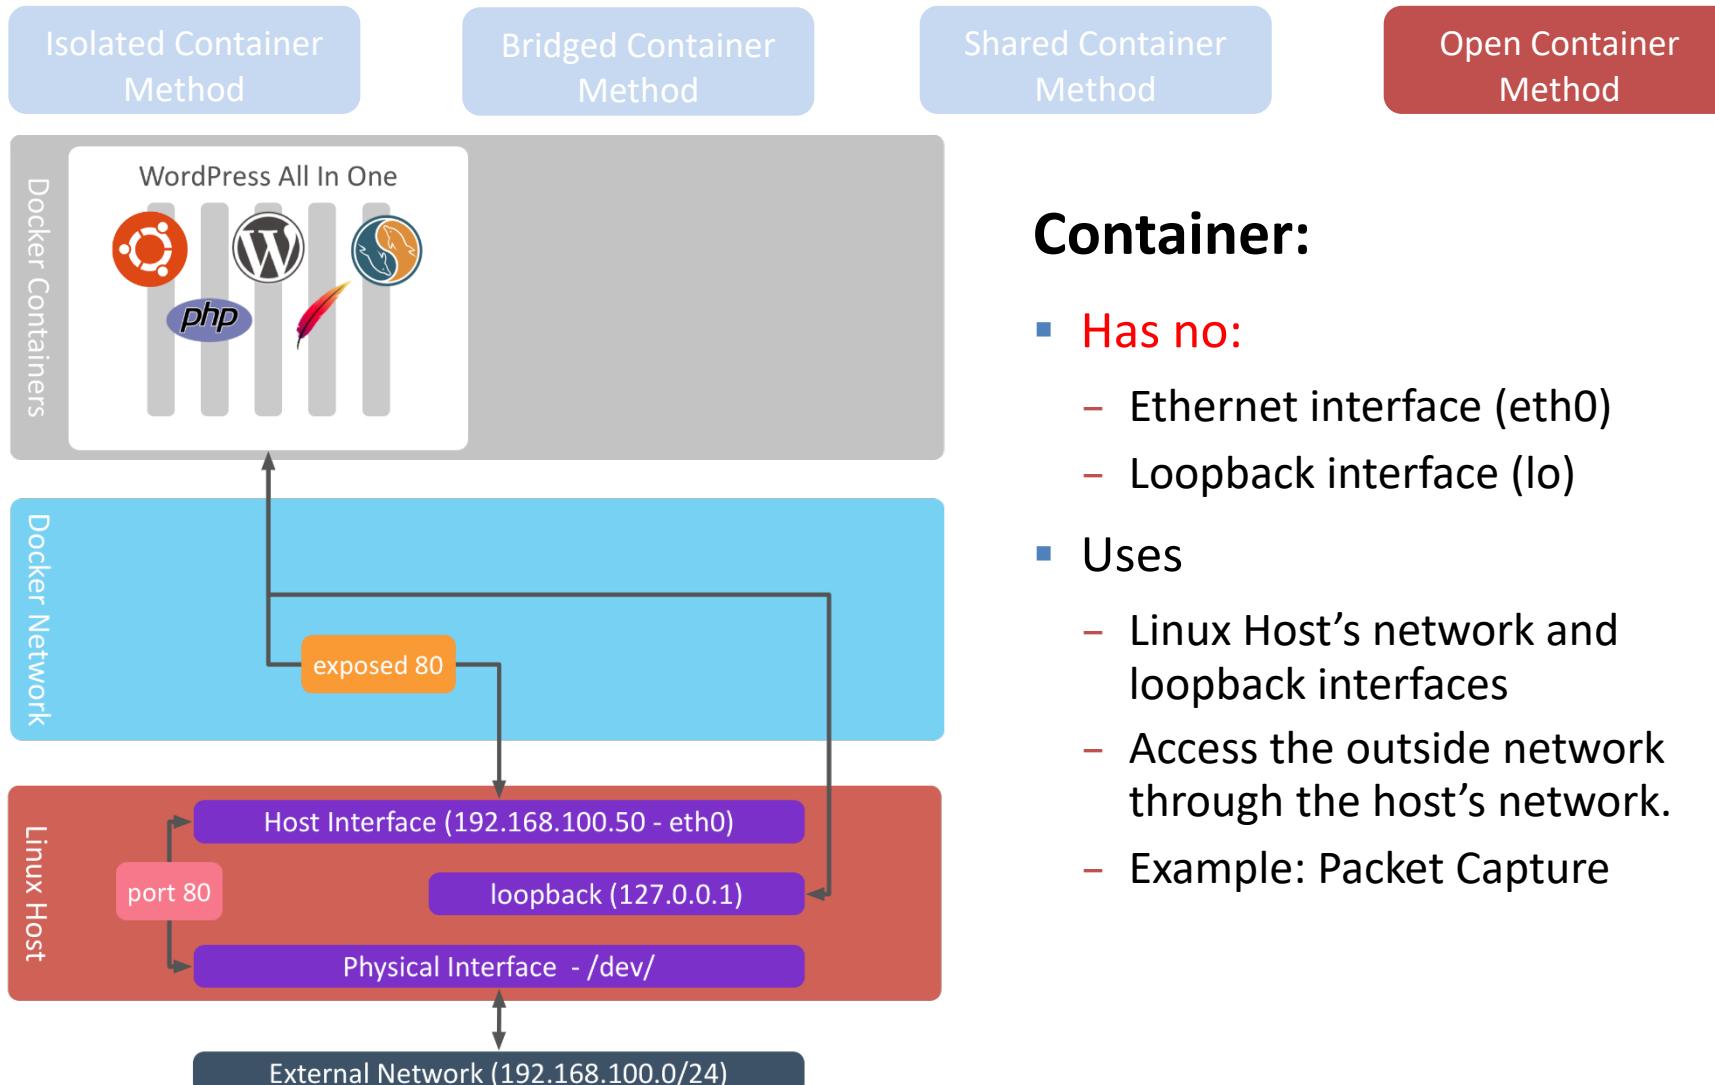
Container One has:

- Access to bridge network through eth0
- Uses Container B's loopback

Container Two has:

- A loopback network
- Uses Container A's eth0 to access network
- Example: “Black Box” processes

Open Container Method





Questions

Service Ports

© 2025 by Innovation In Software



Service Ports

Service names and ports are used to distinguish between different services. Many are **defined in RFC6335**

- Some Common Service Ports:
 - FTP TCP Port 21
 - SSH TCP Port 22
 - HTTP Web Server TCP Port 80
 - HTTP SSL Web Server TCP Port 443
 - LDAP TCP Port 389

RFC 6335

This RFC 6335 was published in 2011.

Abstract

This document defines the procedures that the Internet Assigned Numbers Authority (IANA) uses when handling assignment and other requests related to the Service Name and Transport Protocol Port Number registry. It also discusses the rationale and principles behind these procedures and how they facilitate the long-term sustainability of the registry.

RFC 6335 introduction

For many years, the assignment of new service names and port number values for use with the Transmission Control Protocol (TCP) [RFC0793] and the User Datagram Protocol (UDP) [RFC0768] have been less formal than the IANA registration of IP addresses and port numbers. New mechanisms have been added -- the Stream Control Transmission Protocol (SCTP) [RFC4960] and the Datagram Congestion Control Protocol (DCCP) [RFC4342] -- and new mechanisms like DNS SRV records [RFC2182] have been developed, each with separate registries and separate guidelines.

The community also recognized the need for additional procedures beyond just assignment; notably modification, revocation, and release.

Download links

Click here to download RFC 6335: [TXT format](#) [PDF format \(coming soon\)](#)

Related Request for Comments

- RFC 6766 -Traversal Using Relays around NAT (TURN): Relay Extensions to Session Traversal Utilities for NAT (STUN)
- RFC 6741 - RFC 5246 Headers and Integers
- RFC 5595 - The Datagram Congestion Control Protocol (DCCP) Service Codes
- RFC 5389 - Session Traversal Utilities for NAT (STUN)
- RFC 5237 - IANA Allocation Guidelines for the Protocol Field
- RFC 5226 - Guidelines for IANA Considerations: ABNF
- RFC 5226 - Guidelines for Writing an IANA Considerations Section in RFCs
- RFC 4960 - Stream Control Transmission Protocol
- RFC 4844 - The RFC Series and RFC Editor
- RFC 4722 - Experimental Values in IPv4, IPv6, ICMPv4, ICMPv6, UDP and TCP Headers
- RFC 4620 - Early IANA Allocation of Standards Track Code Points
- RFC 4540 - Datagram Congestion Control Protocol (DCCP)
- RFC 4020 - Early IANA Allocation of Standards Track Code Points
- RFC 3828 - The Lightweight User Datagram Protocol (UDP-Lite)
- RFC 3992 - Assigning Experimental and Testing Numbers Considered Useful

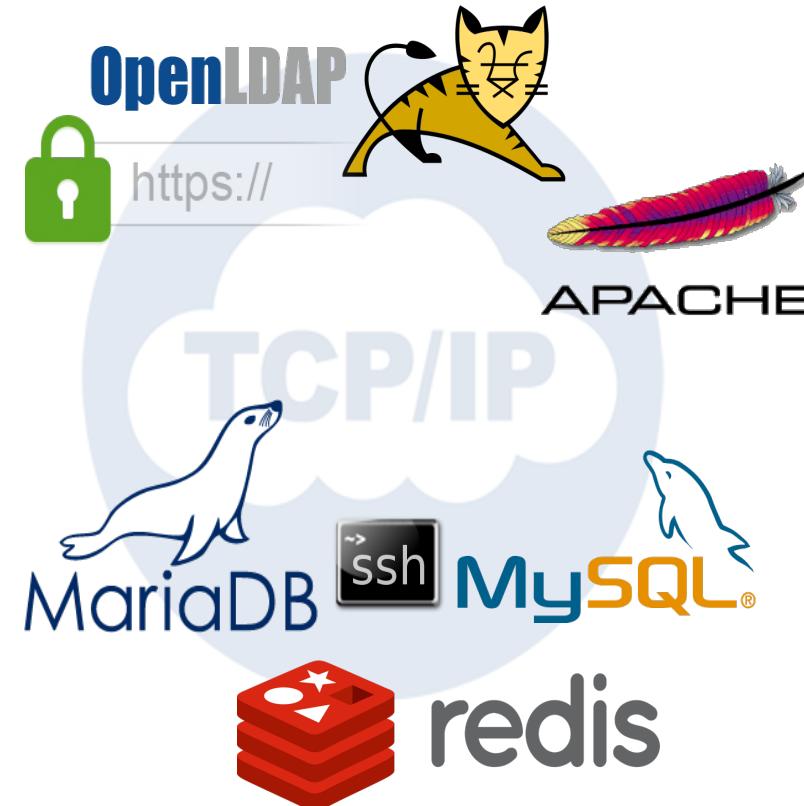
Popular RFCs

- [HTTP/1.1 Protocol - RFC 2616](#)
- [Uniform Resource Locator URL - RFC 1738](#)

Service Ports and Applications

Common TCP Service and Application Ports:

Service	Service Port
FTP	21
SSH	22
HTTP	80
Apache	80
LDAP	389
HTTPS	443
MariaDB	3306
MySQL	3307
redis	6379
tomcat	8080



Expose Service Ports

© 2025 by Innovation In Software



Expose Service Ports

Docker Engine provides a mechanism to expose required service ports

- Exposure is by port, not service name
- Service port exposure is required for communication with the container
- Service ports are exposed:
 - Docker Container creation
 - Docker Image creation



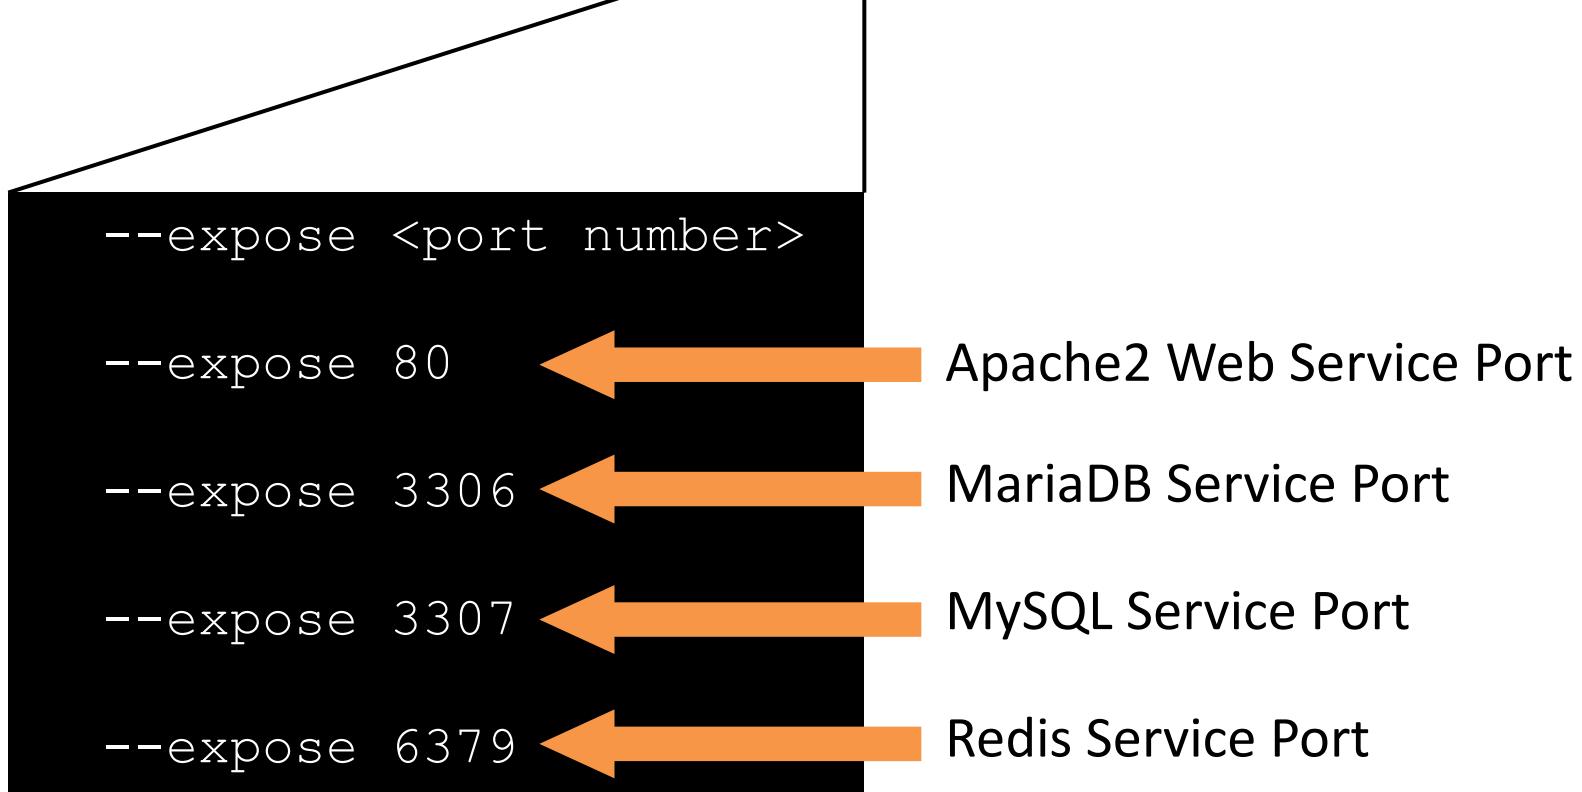
Expose Service Ports

Common TCP Service Ports:

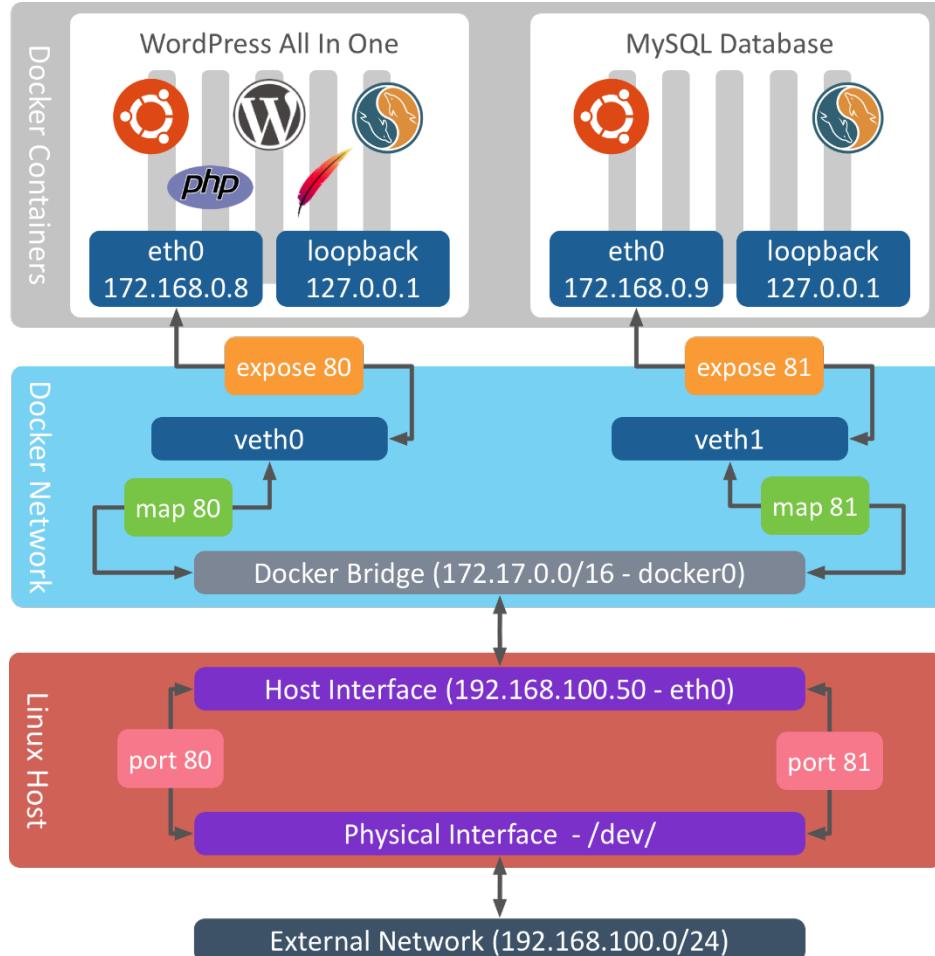
Service	Service Port	Exposed Port
FTP	21	21
SSH	22	22
HTTP	80	80
Apache	80	80
LDAP	389	389
HTTPS	443	443
MariaDB	3306	3306
MySQL	3307	3307
redis	6379	6379
tomcat	8080	8080

Map Ports on a Bridged Container

```
docker run -d --expose 80 -P busybox top
```



Expose Container Service Ports



Container:

- Service ports are defined
- Service ports are exposed from the container
- Ports are not accessible outside Docker network

Port Mapping Containers

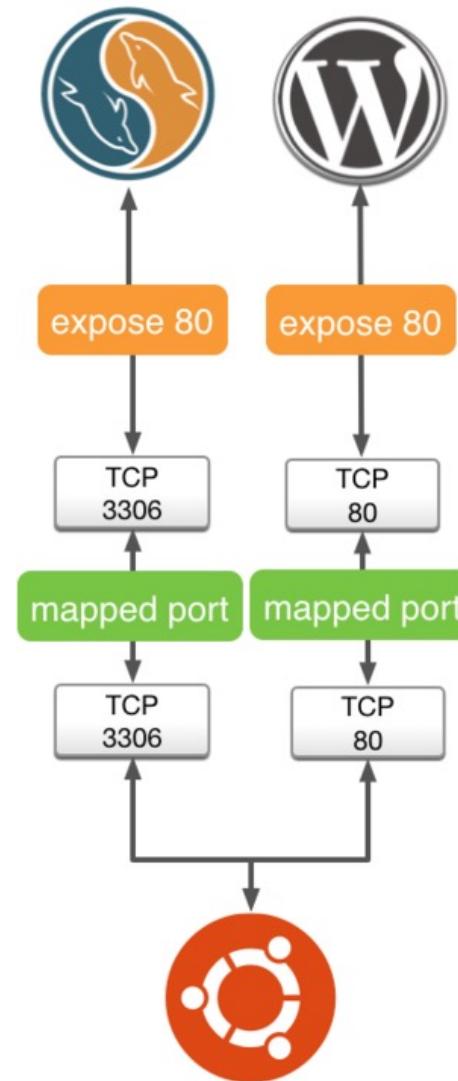
© 2025 by Innovation In Software



Map Exposed Service Ports

Docker provides a mechanism to map exposed service ports

- TCP ports can be mapped:
 - Dynamically
 - On all host interfaces
 - On user defined ports and interfaces
- Exposed service ports are only mapped during container creation:
 - Docker Container creation



Expose Service Ports

Common Service Ports:

Service	Service Port	Expose port	Mapped Port	Host Port
Apache2	80	80	80	33731
HTTPD	80	80	80	33732
LDAP	389	389	389	33733
MariaDB	3306	3306	3306	33734
MySQL	3307	3307	3307	33735
redis	6379	6379	6379	33736

Docker provides a mechanism to:

- map exposed service ports → mapped service ports → host TCP ports

- Note the two service ports using TCP Port 80

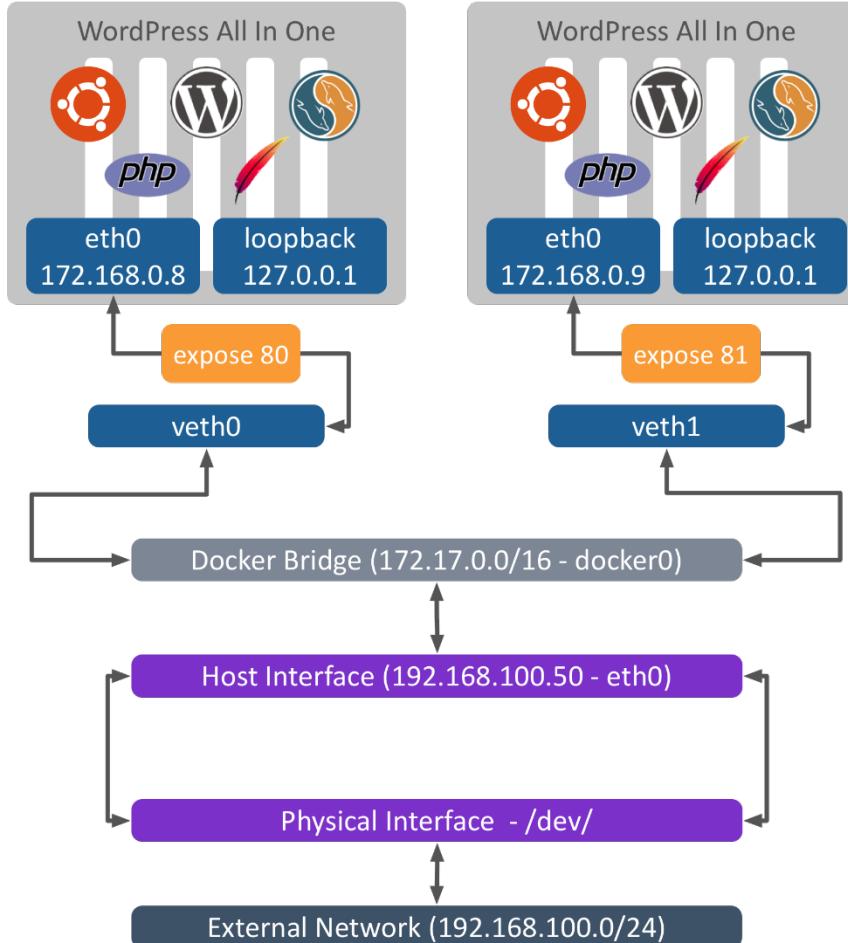
WordPress Service Ports

WordPress AIO Service Ports:

Service	Service Port	Exposed Port	Mapped Port	Host Port
WPAIO01	80	80	80	33731
WPAIO02	80	80	80	33732
WPAIO03	80	80	80	33733
WPAIO04	80	80	80	33734
WPAIO05	80	80	80	33735
WPAIO06	80	80	80	33736

The above table depicts how a single Linux Host can host up six WordPress AIO(s) on container port 80, using dynamically assigned ports on the Linux Host

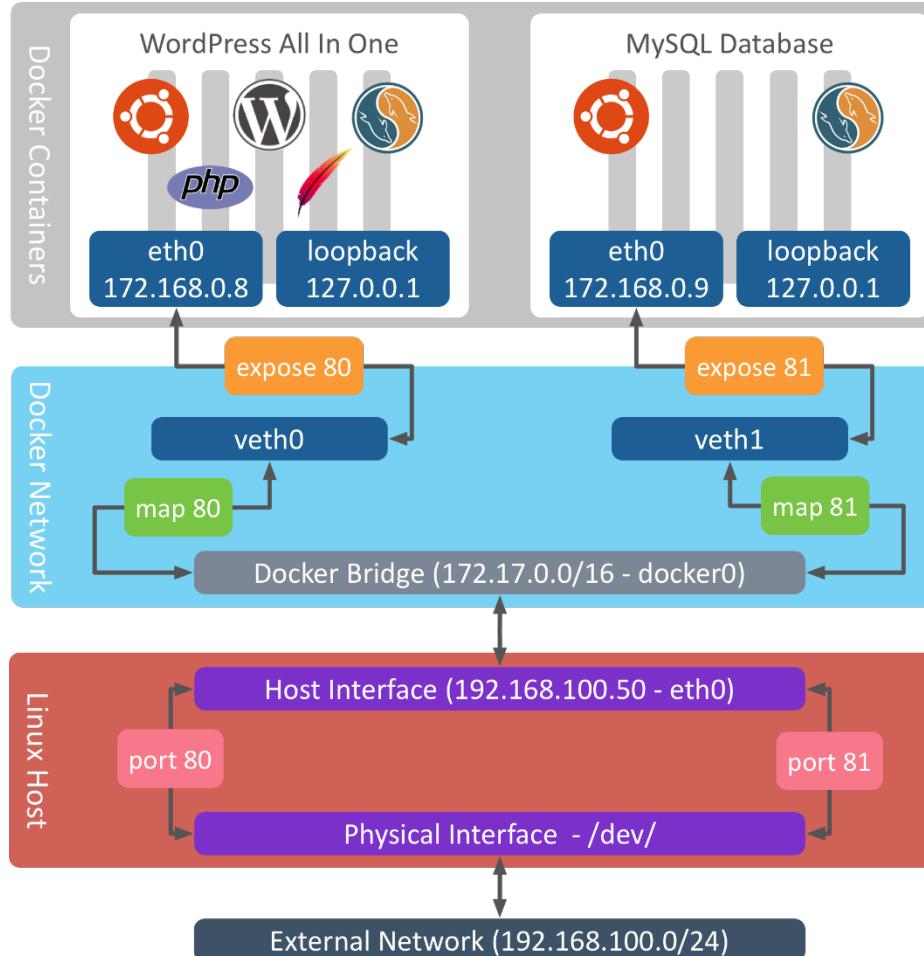
Expose Container Service Ports



Container:

- Service ports are defined
- Service ports are exposed from the container
- Ports are not accessible

Map Container Service Ports



Container:

- Service Ports are known
- Ports are exposed from the container
- Host ports are mapped to container ports for accessibility

Map Ports on a Bridged Container

```
docker run -d -p busybox:latest top
```

```
-P
```

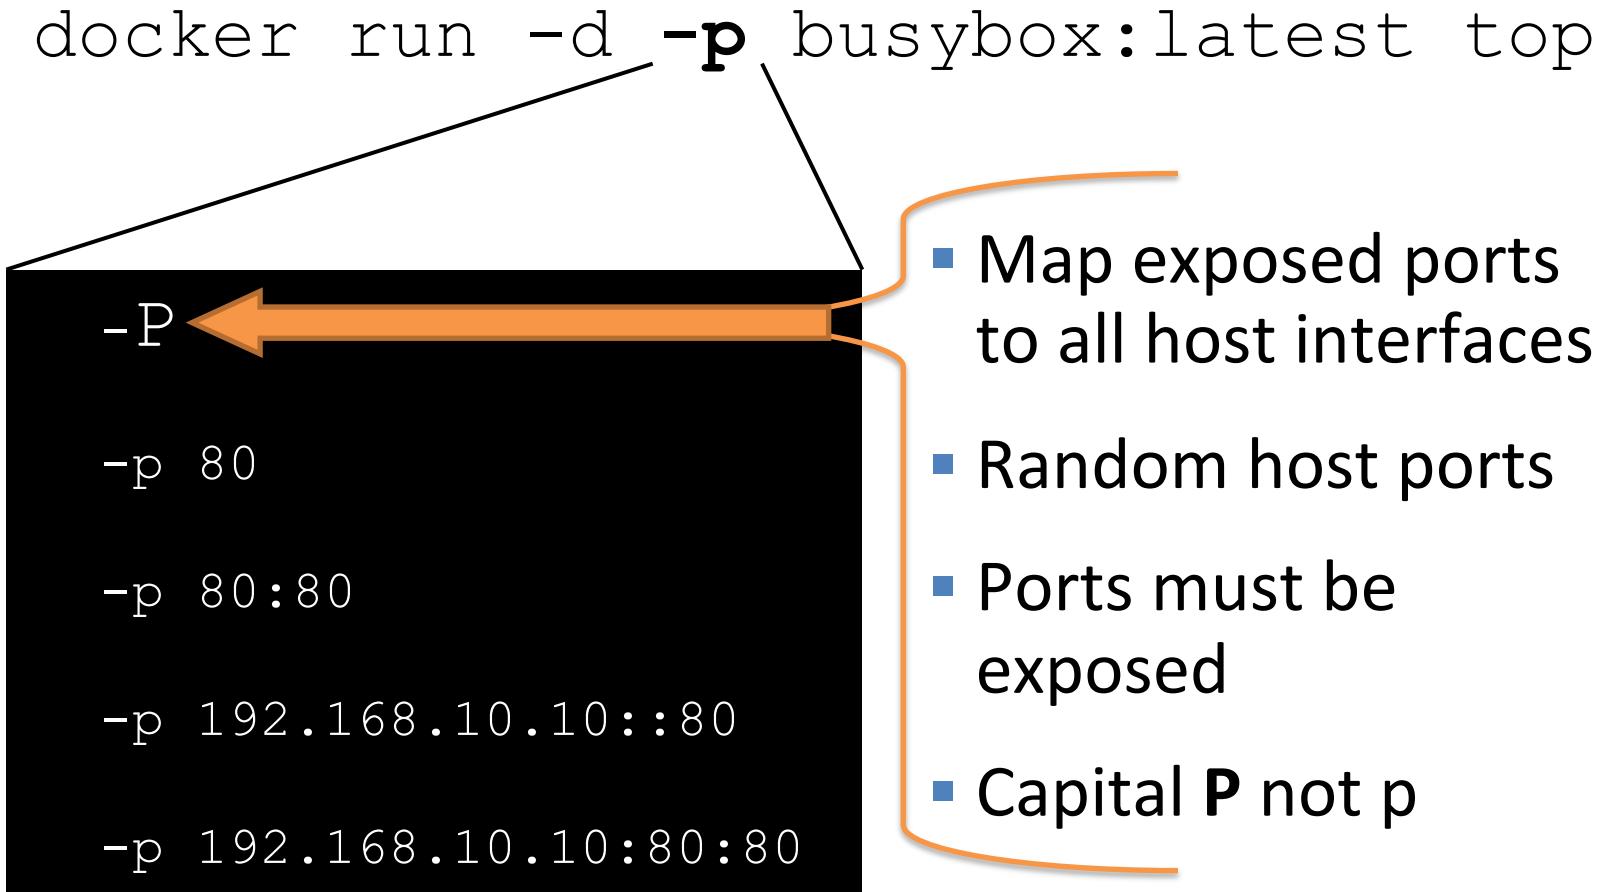
```
-p 80
```

```
-p 80:80
```

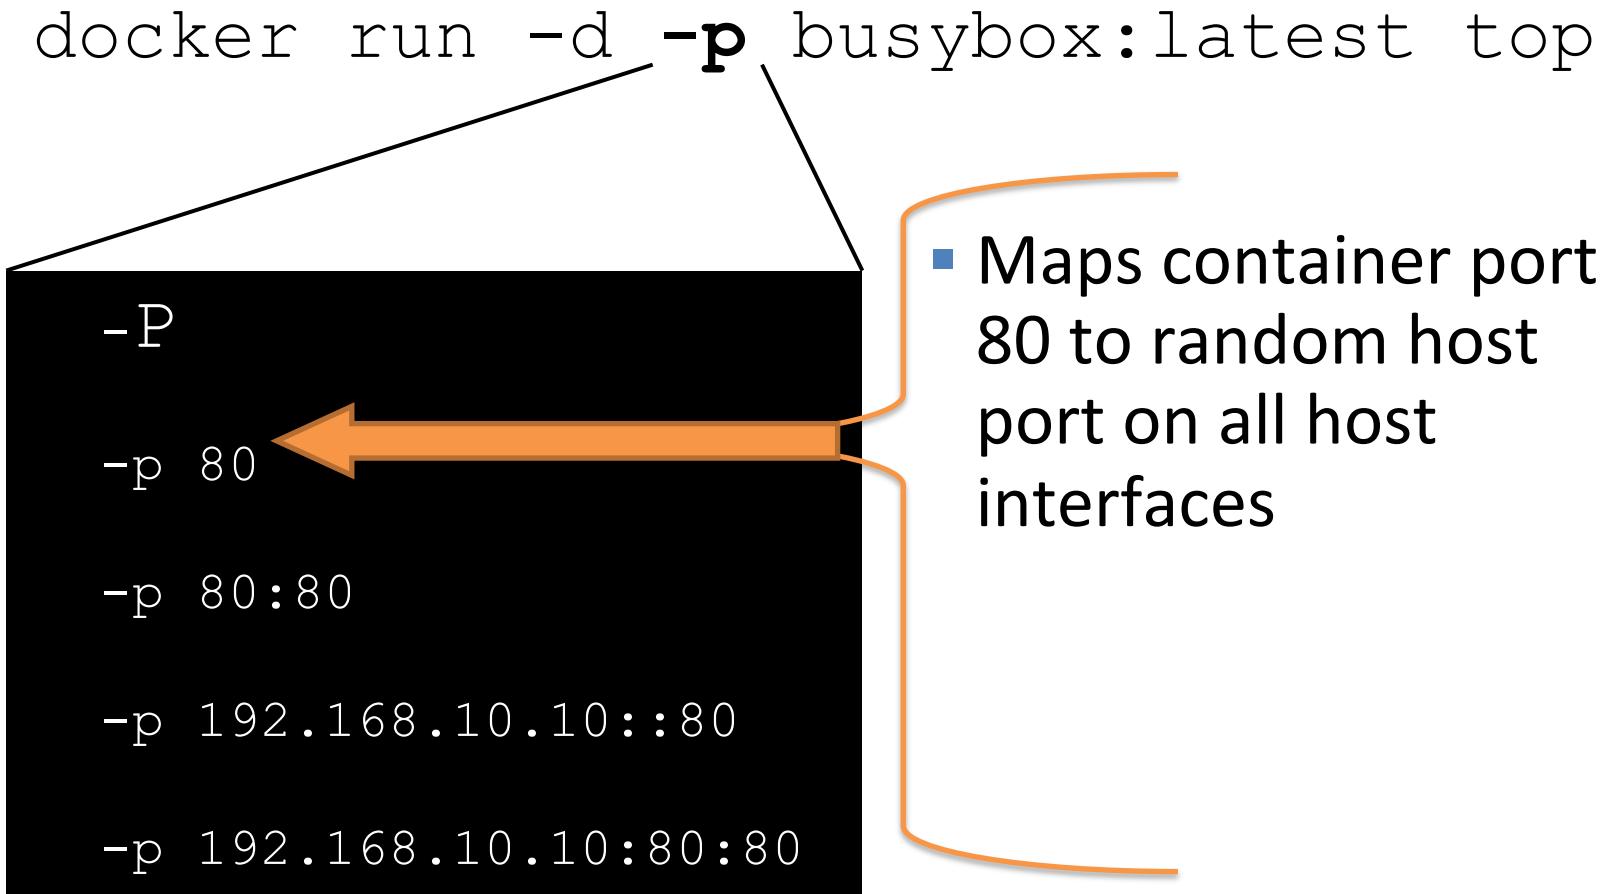
```
-p 192.168.10.10::80
```

```
-p 192.168.10.10:80:80
```

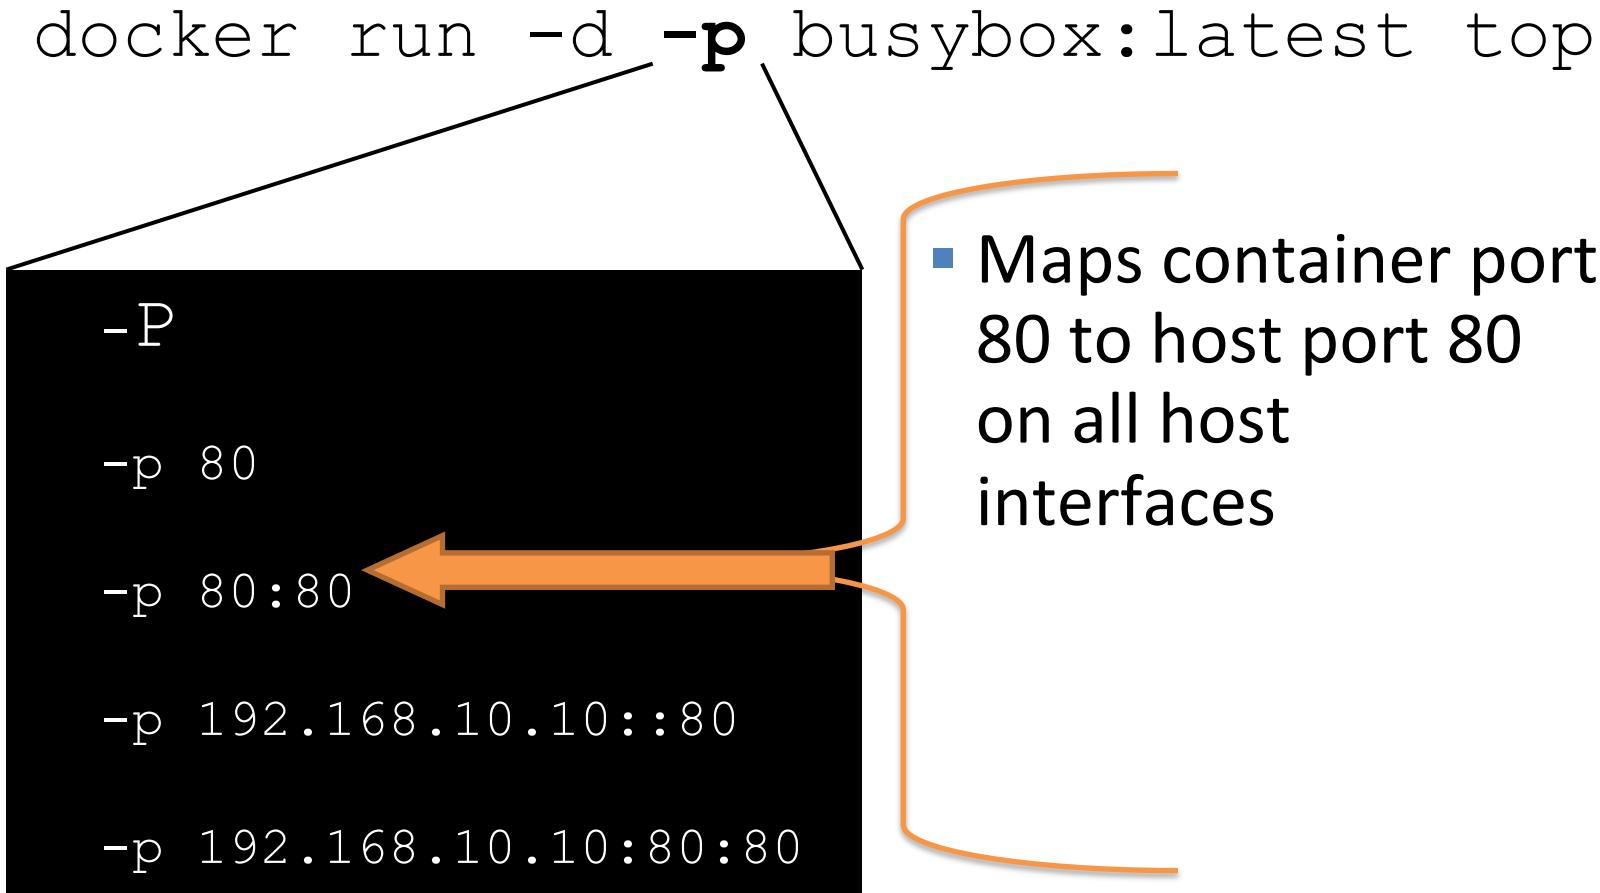
Map Ports on a Bridged Container



Map Ports on a Bridged Container



Map Ports on a Bridged Container



Map Ports on a Bridged Container

```
docker run -d -p busybox:latest top
```

```
-P  
-p 80  
-p 80:80  
-p 192.168.10.10::80  
-p 192.168.10.10:80:80
```

- Maps container port 80 to random host port on specified host interface

Map Ports on a Bridged Container

```
docker run -d -p busybox:latest top
```

```
-P  
-p 80  
-p 80:80  
-p 192.168.10.10::80  
-p 192.168.10.10:80:80
```

- Maps container port 80 to host port 80 on host specified host interface



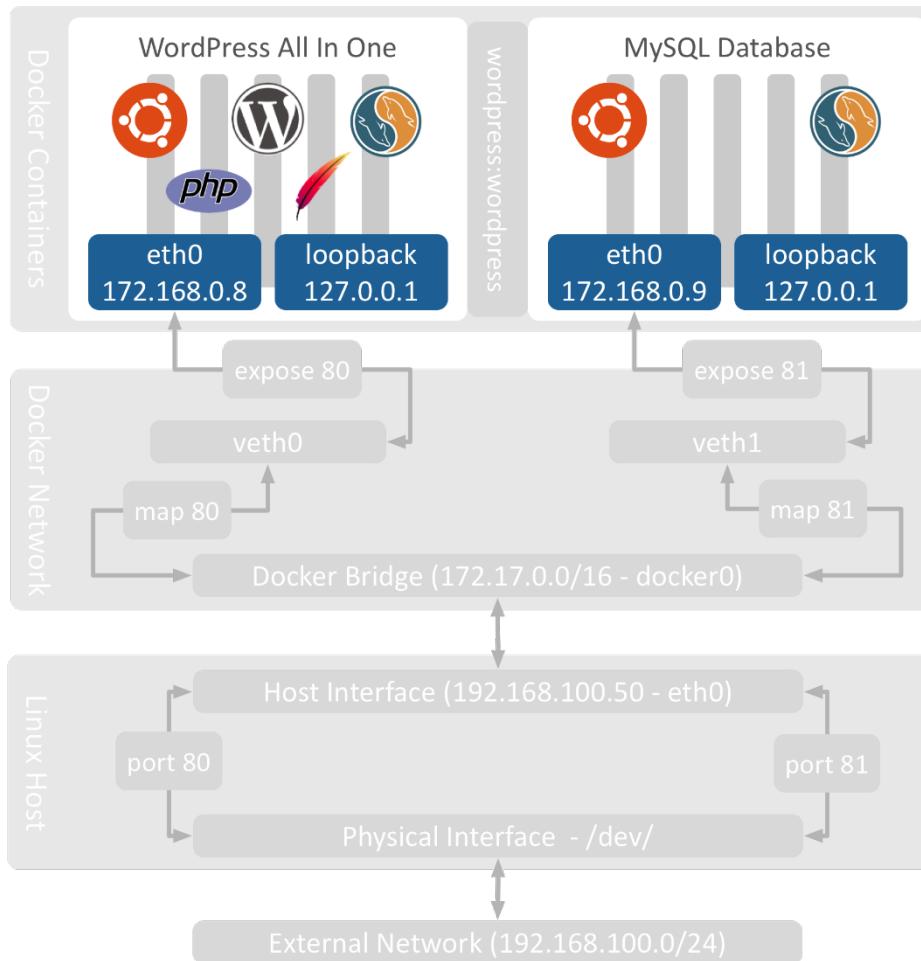
Questions

Link Containers

© 2025 by Innovation In Software



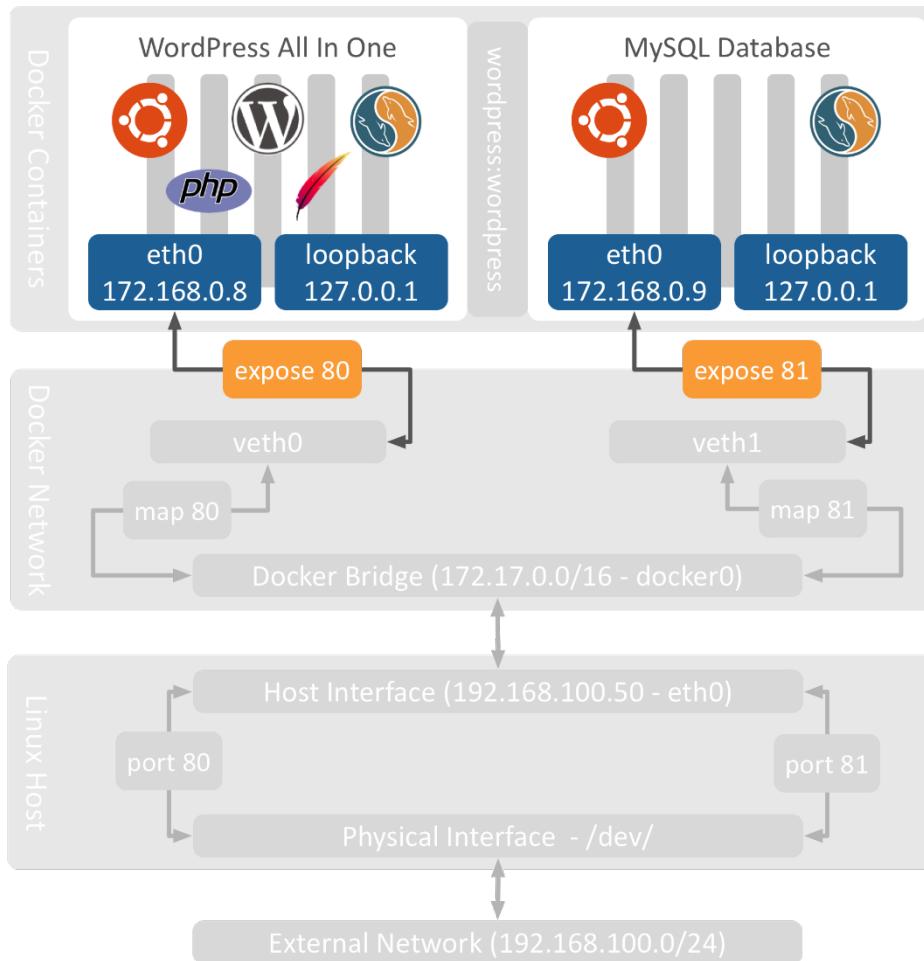
Link containers (LEGACY)



Linking:

- Service Ports are known to all containers
- Ports are exposed from the container
- Host ports are mapped to container ports for accessibility
- Linking containers create internal routing rules
- Linking containers uses aliases to connect containers

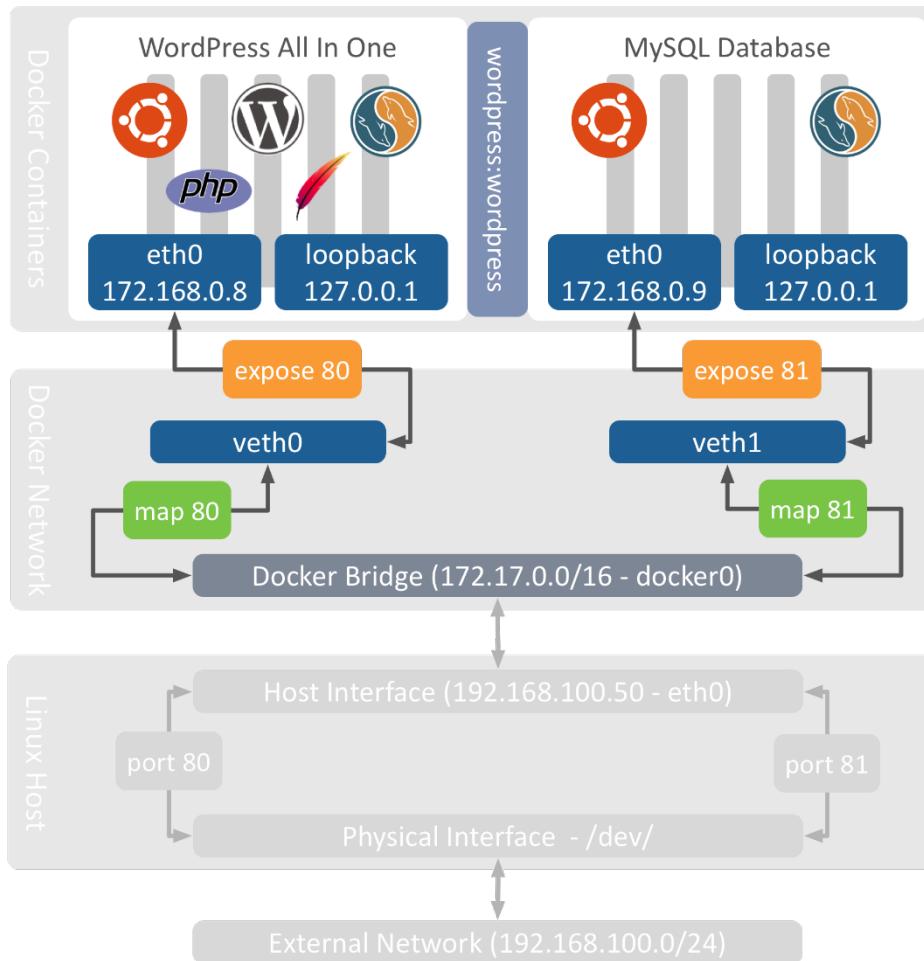
Link containers (LEGACY)



Linking:

- Service Ports are known to all containers
- Ports are exposed from the container
- Host ports are mapped to container ports for accessibility
- Linking containers create internal routing rules
- Linking containers uses aliases to connect containers

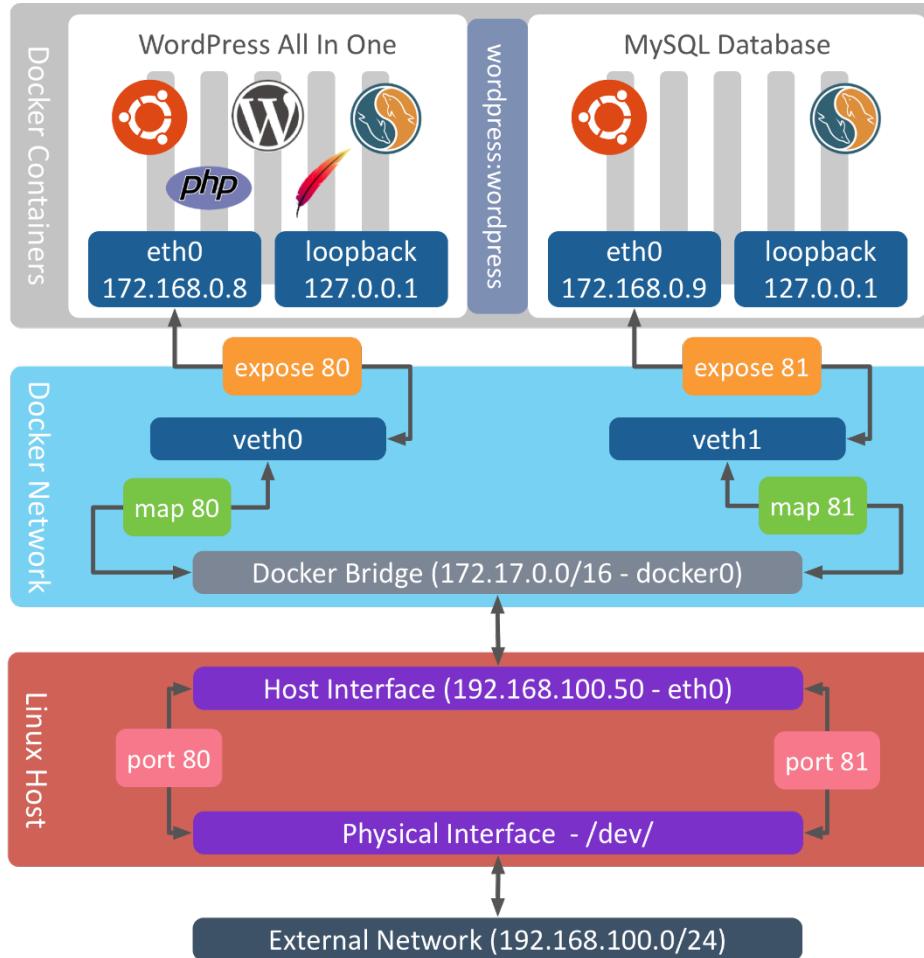
Link containers (LEGACY)



Linking:

- Service Ports are known to all containers
- Ports are exposed from the container
- Host ports are mapped to container ports for accessibility
- Linking containers create internal routing rules
- Linking containers uses aliases to connect containers

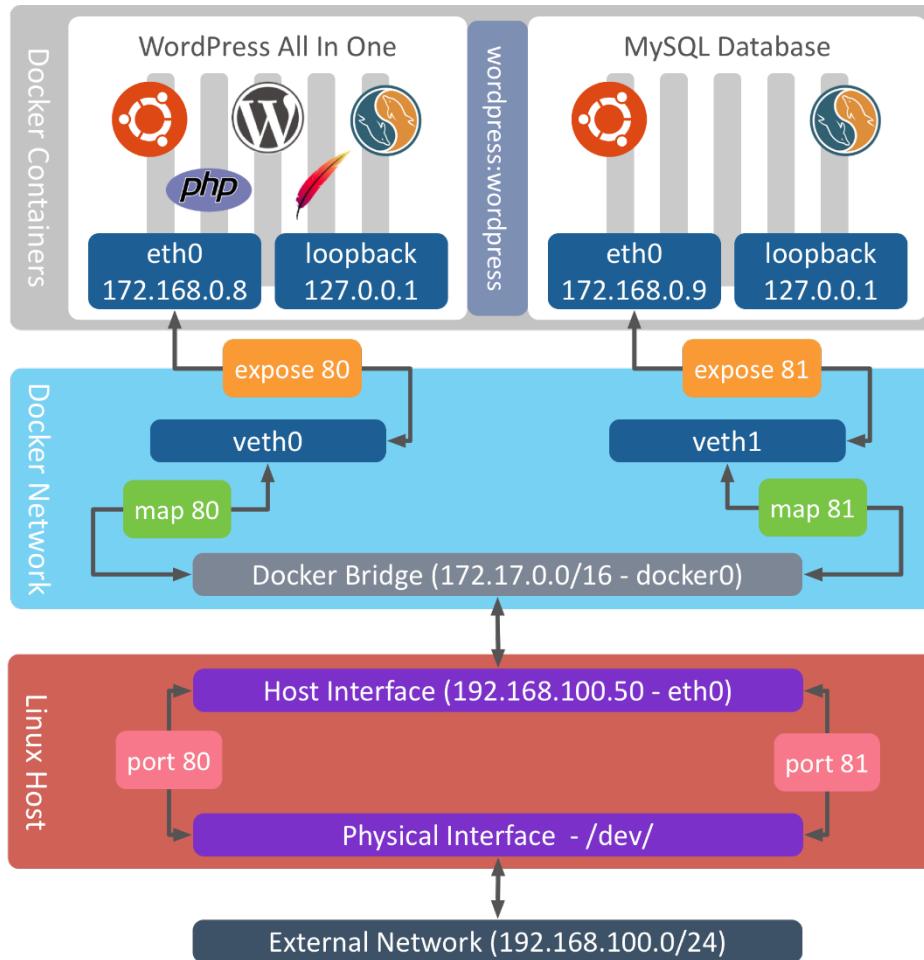
Link containers (LEGACY)



Linking:

- Service Ports are known to all containers
- Ports are exposed from the container
- Host ports are mapped to container ports for accessibility
- Linking containers create internal routing rules
- Linking containers uses aliases to connect containers

Link containers (LEGACY)



Linking:

- Service Ports are known to all containers
- Ports are exposed from the container
- Host ports are mapped to container ports for accessibility
- Linking containers create internal routing rules
- Linking containers uses aliases to connect containers

User Traffic Flow

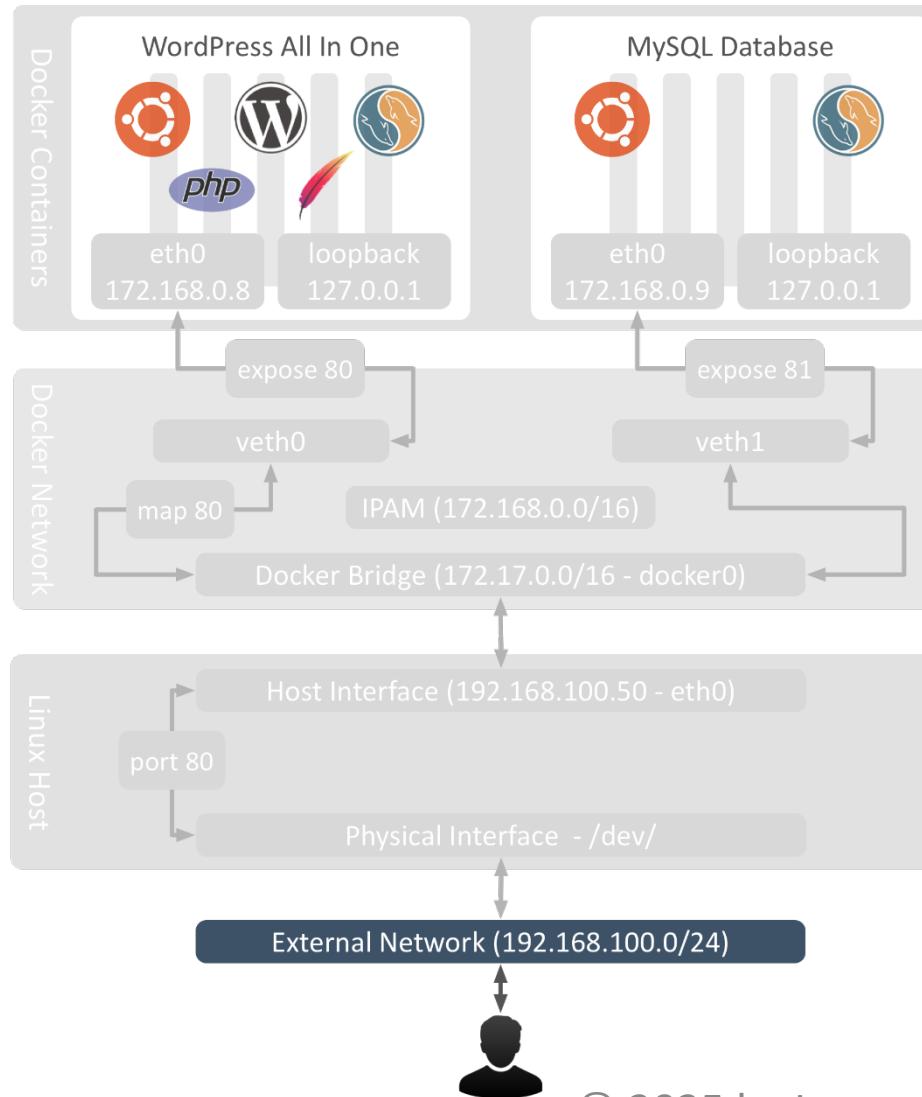
© 2025 by Innovation In Software



User Traffic Flow

- **Applications provide:**
 - Service Ports that need to be exposed
- **Docker Engine provides:**
 - Exposure of required Service Ports
 - Mapping of required Service Ports
- **Linux Host provides:**
 - Internal connectivity between containers
 - Internal connectivity between host and containers
 - External connectivity between containers and users

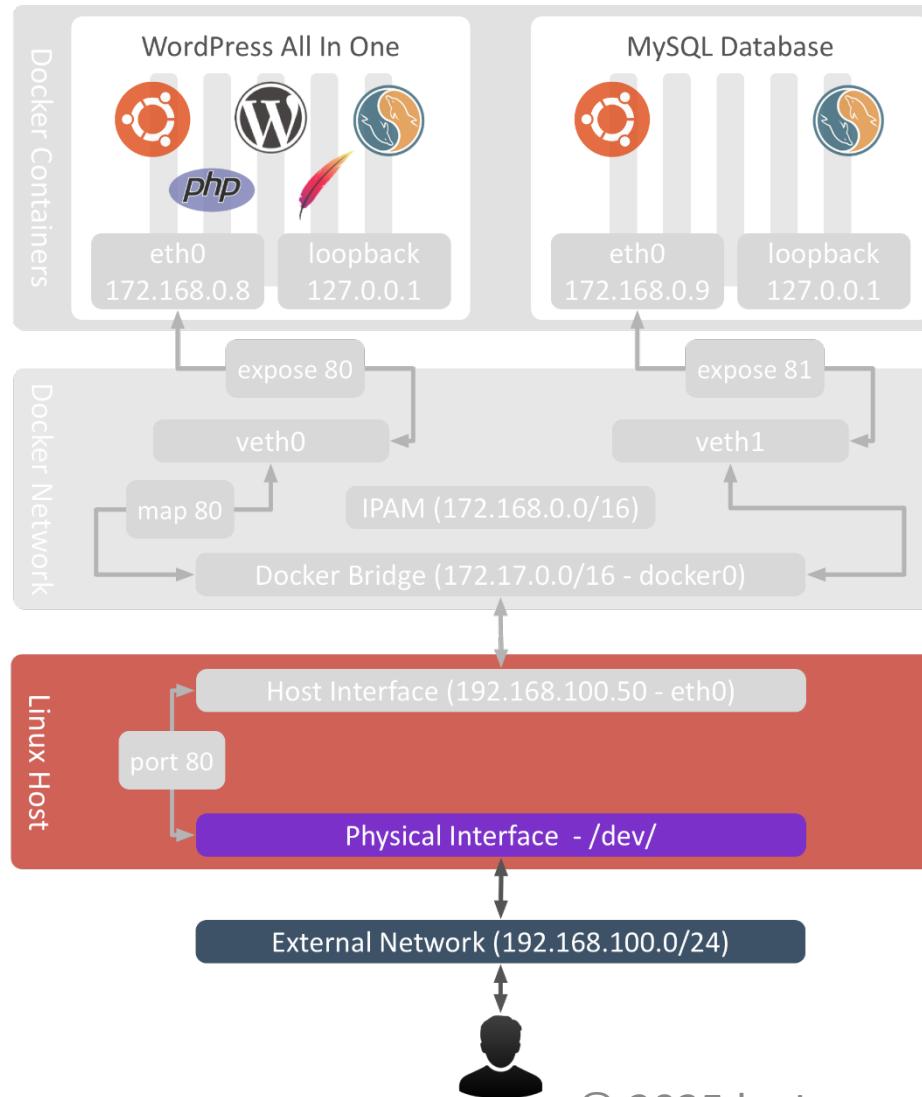
User Traffic Flow



User Traffic:

- Enters from the external network (outside the Docker host)
- Enters the physical interface
- Is passed to the host interface (eth0)
- Is passed to the Docker Bridge
- Is passed to the veth interface
- Is passed to eth0 in the container

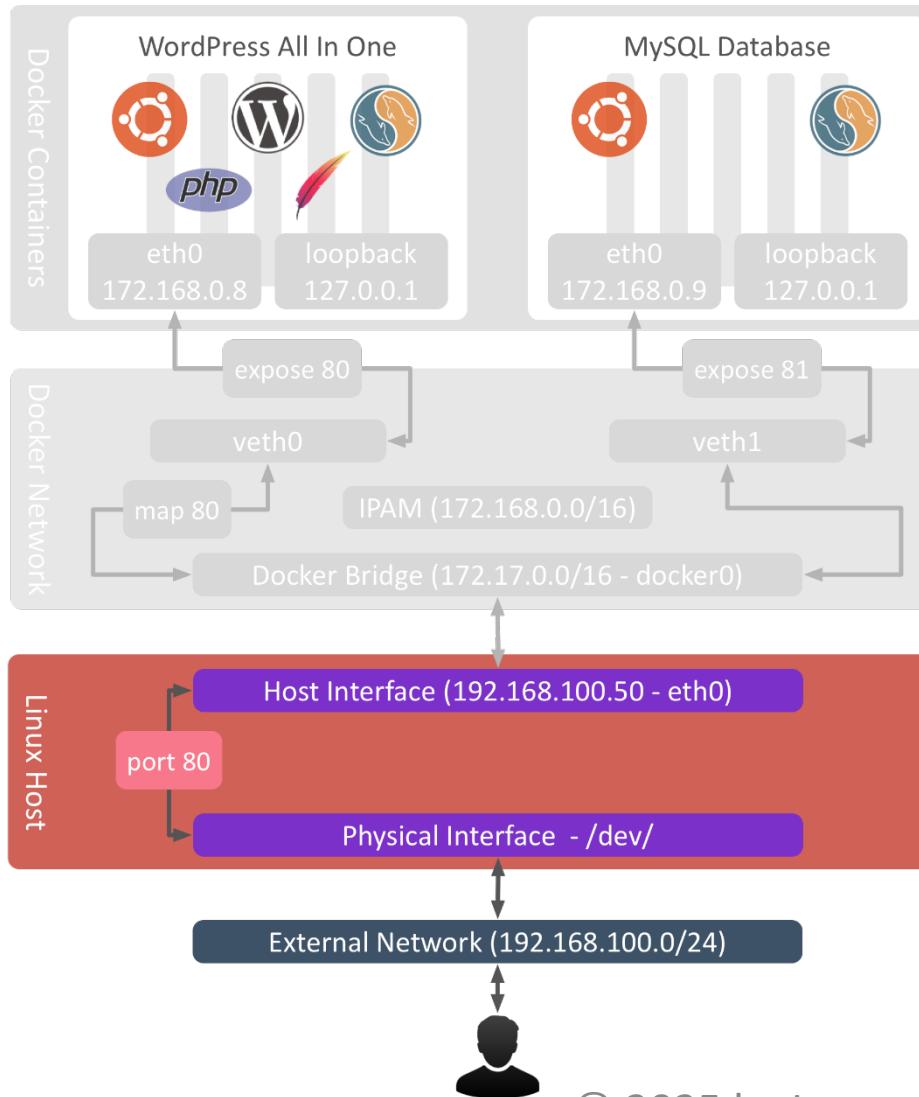
User Traffic Flow



User Traffic:

- Enters from the external network (outside the Docker host)
- Enters the physical interface
- Is passed to the host interface (eth0)
- Is passed to the Docker Bridge
- Is passed to the veth interface
- Is passed to eth0 in the container

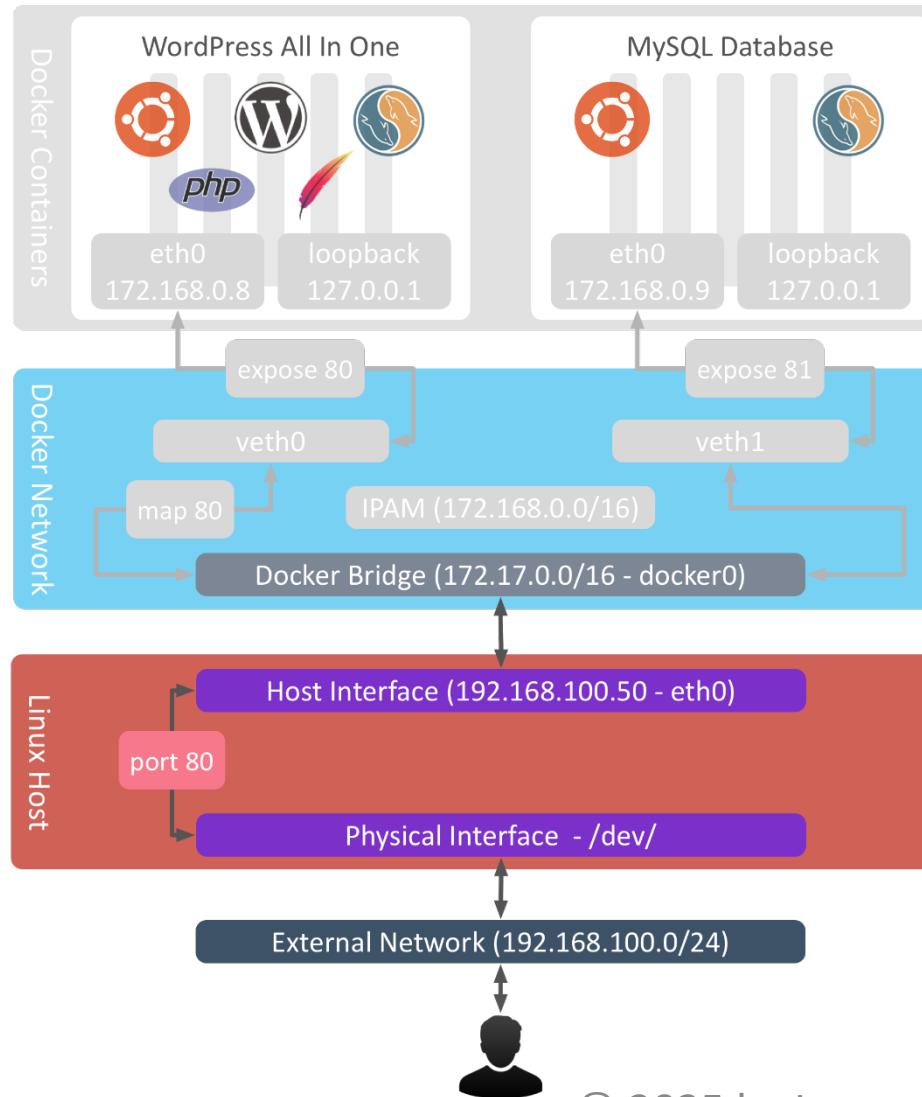
User Traffic Flow



User Traffic:

- Enters from the external network (outside the Docker host)
- Enters the physical interface
- Is passed to the host interface (eth0)
- Is passed to the Docker Bridge
- Is passed to the veth interface
- Is passed to eth0 in the container

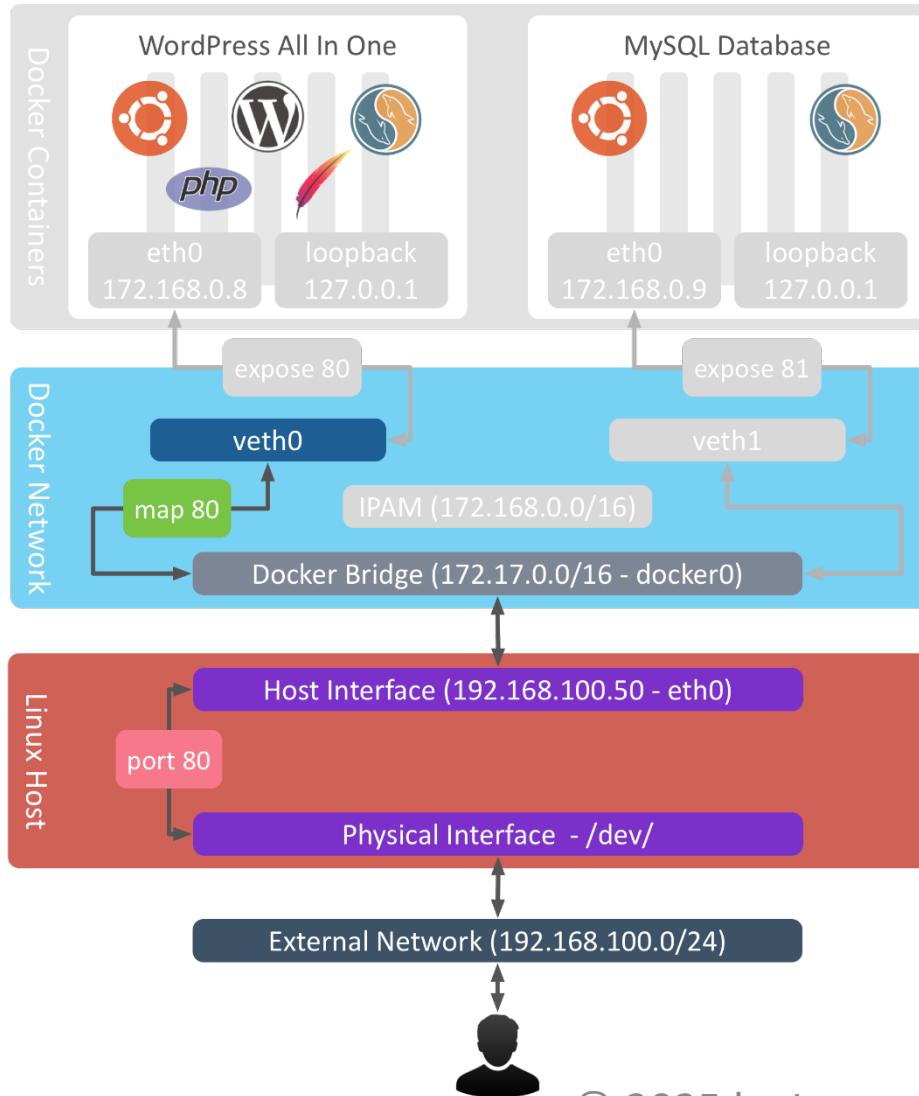
User Traffic Flow



User Traffic:

- Enters from the external network (outside the Docker host)
- Enters the physical interface
- Is passed to the host interface (eth0)
- Is passed to the Docker Bridge
- Is passed to the veth interface
- Is passed to eth0 in the container

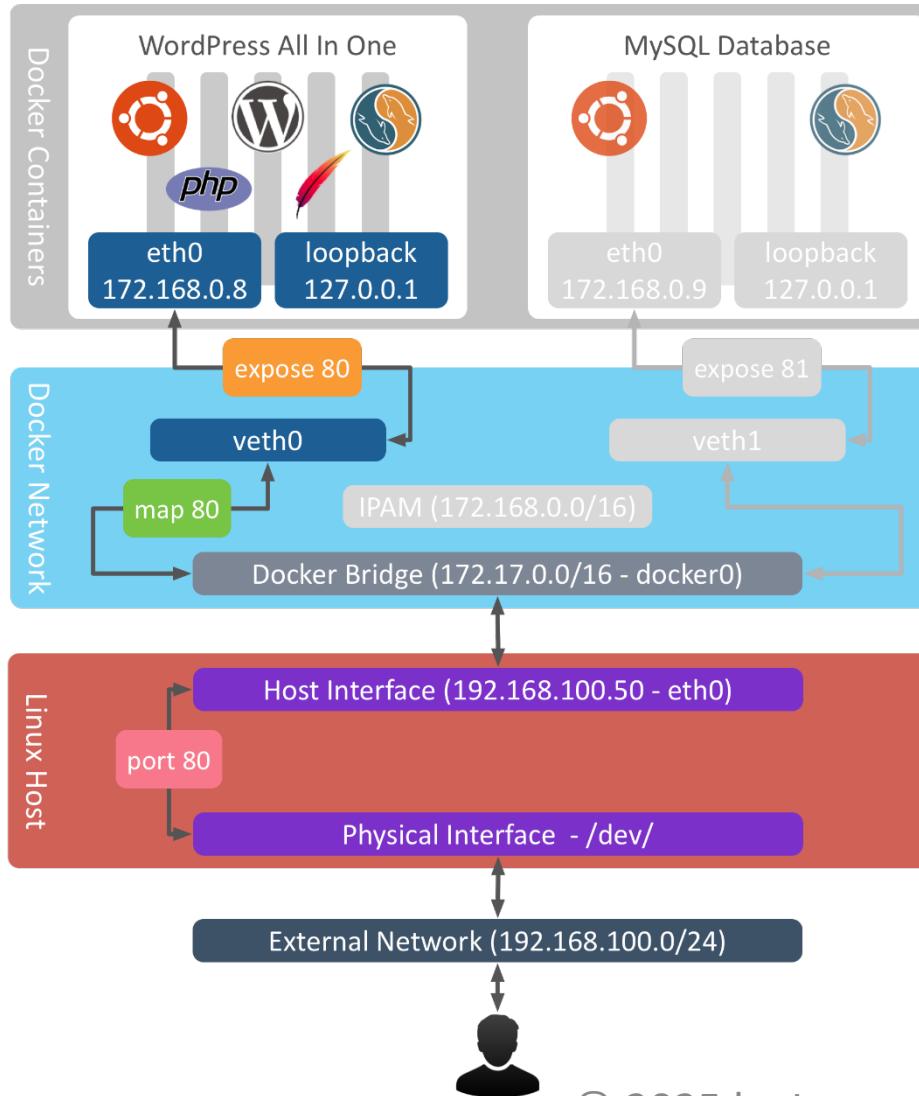
User Traffic Flow



User Traffic:

- Enters from the external network (outside the Docker host)
- Enters the physical interface
- Is passed to the host interface (`eth0`)
- Is passed to the Docker Bridge
- Is passed to the veth interface
- Is passed to `eth0` in the container

User Traffic Flow



User Traffic:

- Enters from the external network (outside the Docker host)
- Enters the physical interface
- Is passed to the host interface (`eth0`)
- Is passed to the Docker Bridge
- Is passed to the `veth` interface
- Is passed to `eth0` in the container**

Add Networks to Docker Bridge

Add Networks to Docker Bridge

Four types of Networks:

1. **none** (no network access, most secure)
2. **bridge** (limited access to host network, moderately secure)
3. **host** (no restrictions to host networks, least secure)
4. **overlay** (networks with other Docker hosts)

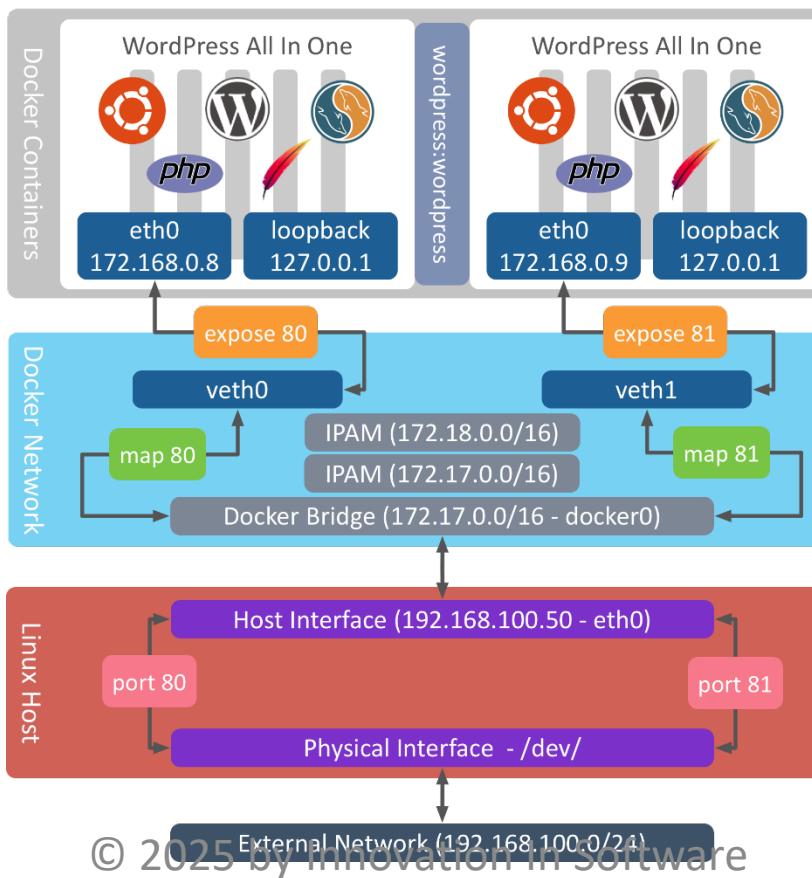
Why use Docker bridge Network segments?

- Use functions of docker bridge to isolate containers from each others by limited **inter-container communication (icc)**
- Provide levels of security in the same docker host

Add Networks to Docker Bridge

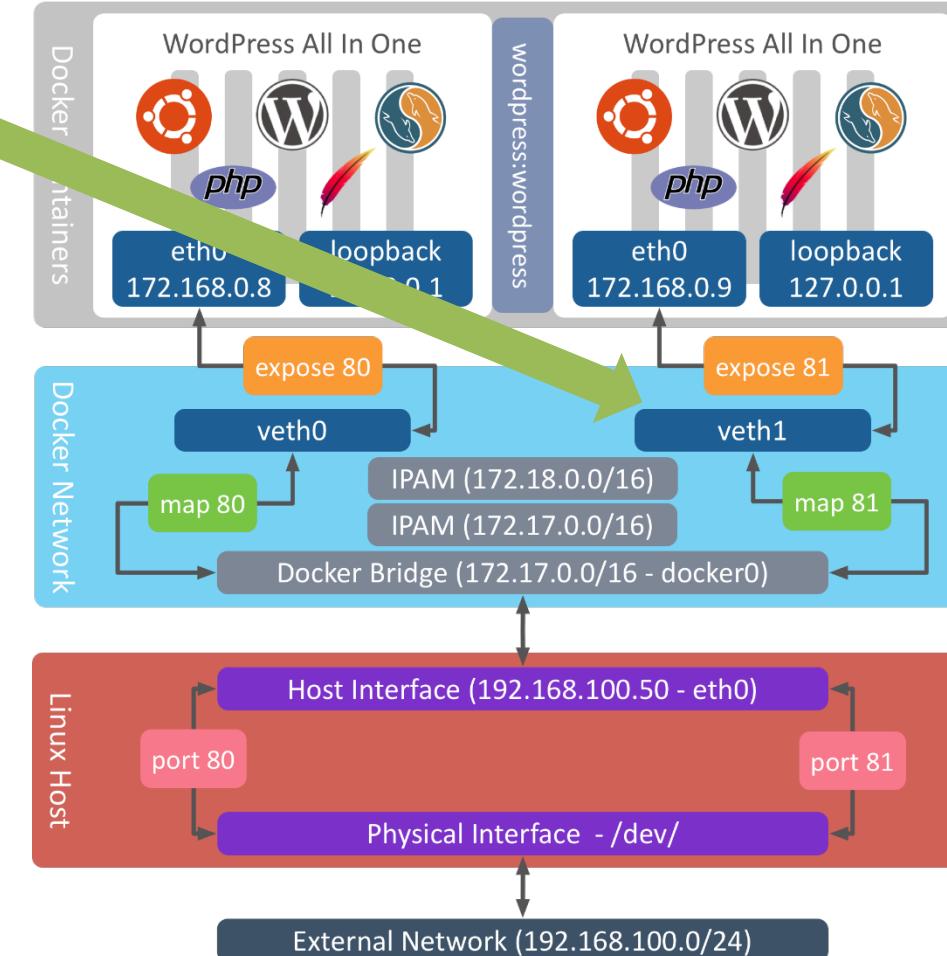
Create additional networks on bridge docker0

```
$ docker network create database --driver bridge
```



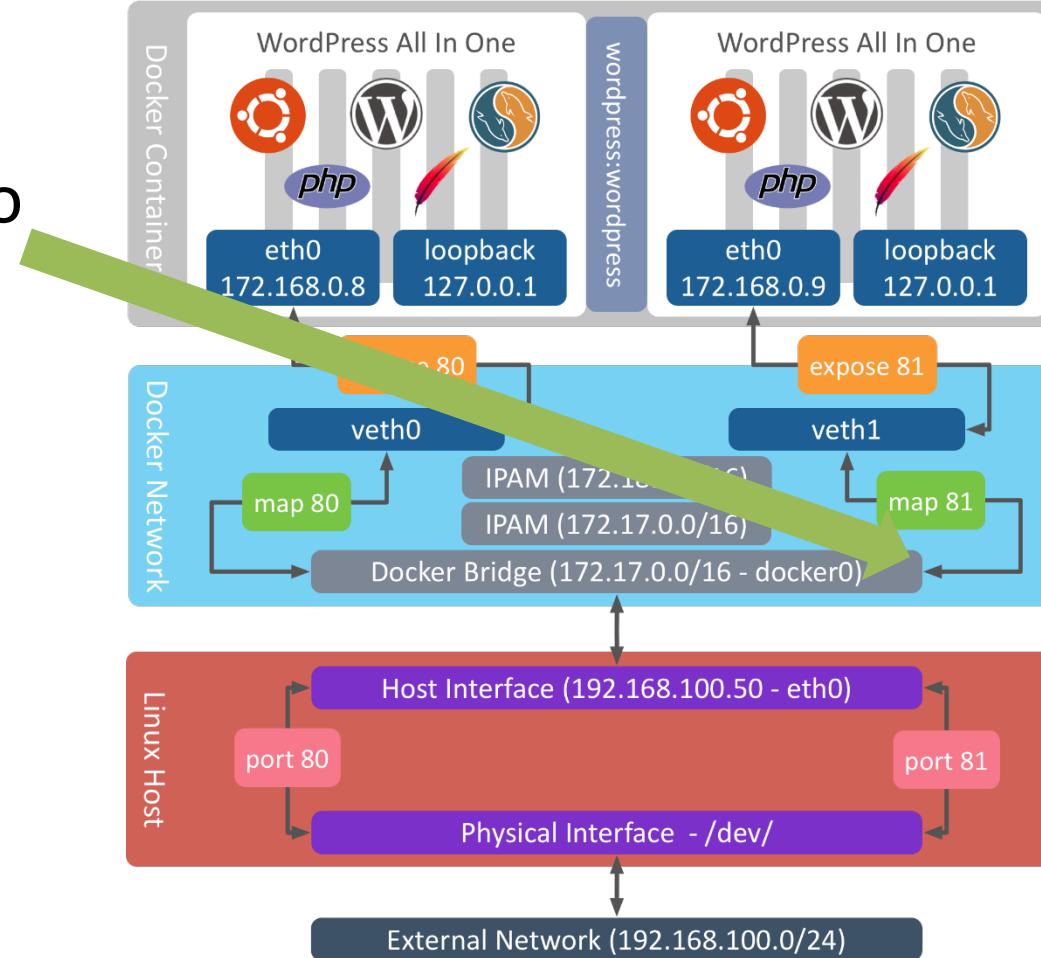
Add Networks to Docker Bridge

- Creates a veth interface pair
- Connects one end to docker0 bridge
- Connects other end to container eth0
- Assigns an IP address from docker0 IPAM for network



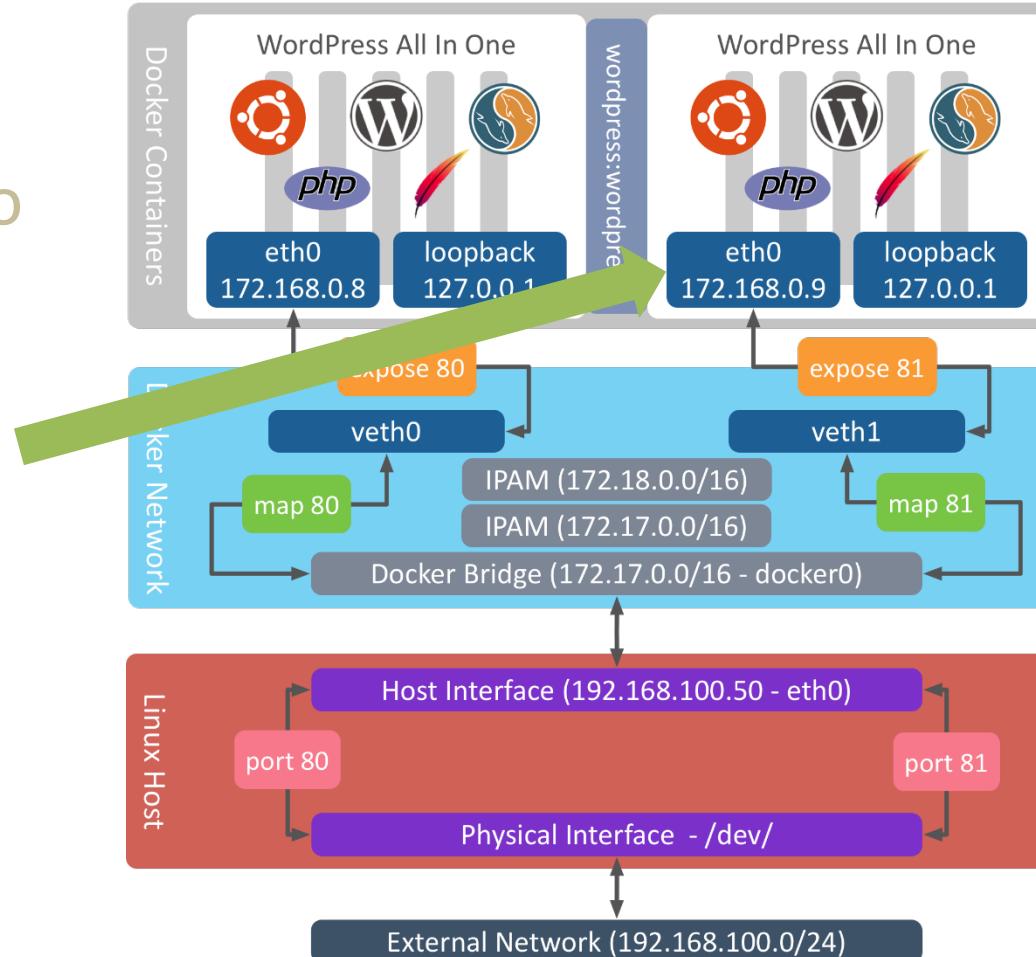
Add Networks to Docker Bridge

- Creates a veth interface pair
- Connects one end to docker0 bridge
- Connects other end to container eth0
- Assigns an IP address from docker0 IPAM for network



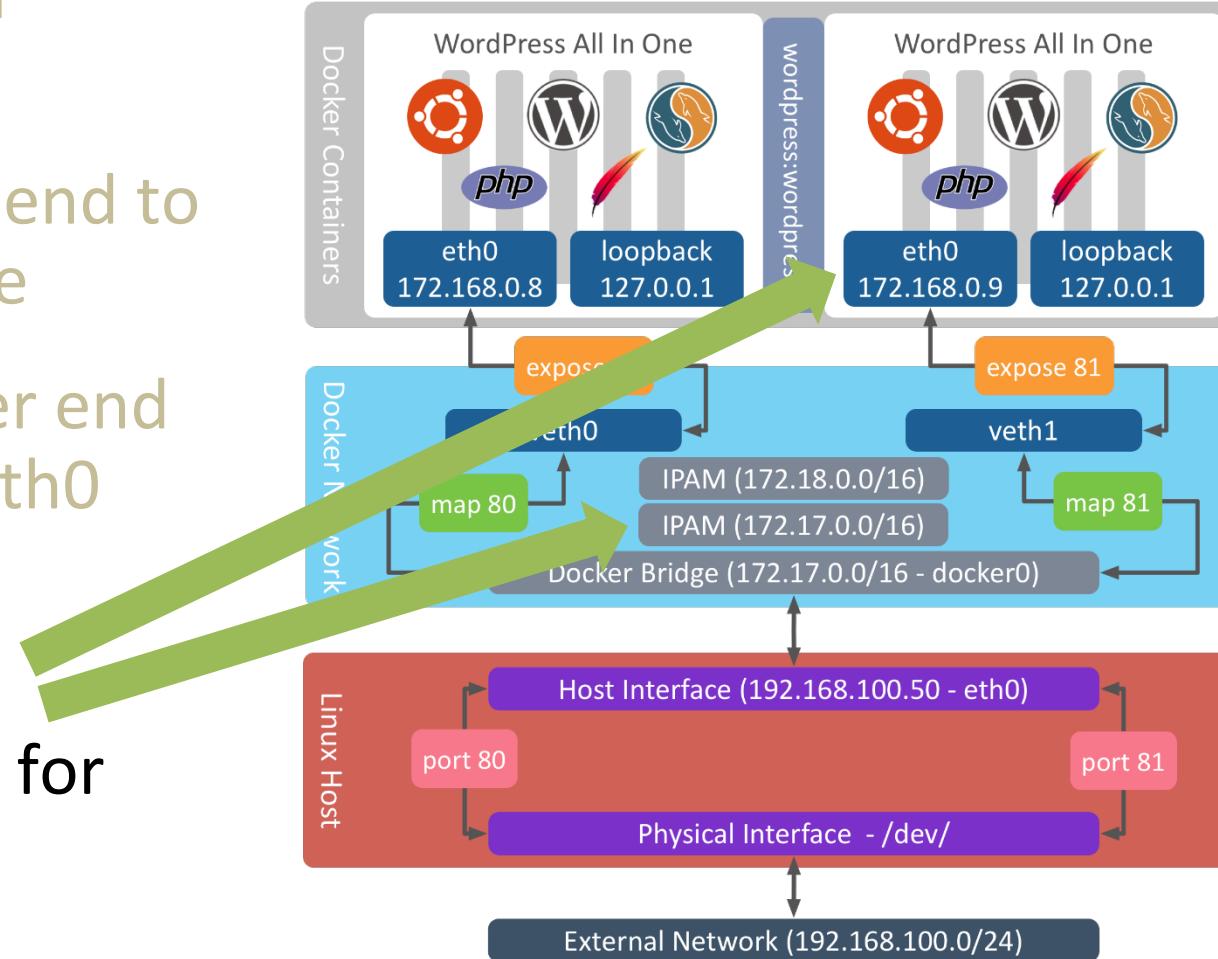
Add Networks to Docker Bridge

- Creates a veth interface pair
- Connects one end to docker0 bridge
- Connects other end to container eth0
- Assigns an IP address from docker0 IPAM for network



Add Networks to Docker Bridge

- Creates a veth interface pair
- Connects one end to docker0 bridge
- Connects other end to container eth0
- Assigns an IP address from docker0 IPAM for network



Add Networks to Docker Bridge

Create an additional network on docker0:

- defaults to docker bridge (docker0)

```
$ docker network create database--driver  
bridge
```

Network is in docker network ls output:

```
$ docker network ls  
NETWORK ID          NAME        DRIVER  
82cb7fdb657d        bridge      bridge  
2e46aa7d66a0        dbase      bridge  
957de60d5958        host       host  
6b64d7ed1c93        none       null  
2c257ae62bcd       wordpress  bridge
```

Add Networks to Docker Bridge

Docker network inspect <network>:

- Provides details on the network

```
$ docker network inspect database
```

- IPAM range, this is the range that the containers are assigned IP Addresses

```
"IPAM": {  
    "Driver": "default",  
    "Options": {},  
    "Config": [  
        {  
            "Subnet": "172.19.0.0/16",  
            "Gateway": "172.19.0.1/16"  
        }  
    ]  
}
```

Example Using default network

Deploy a demo container (default network):

- network option undefined using “bridge,” the default docker0 network

```
$ docker run -d busybox top
```

- docker inspect formatted for bridge and IPAddress of container.

```
$ docker inspect -f \
{.NetworkSettings.Networks.bridge.IPAddress} \
$(docker ps -lq)
```

```
172.17.0.5
```

Example using defined network

Deploy a demo container (**defined network**):

- network option undefined using “bridge,” the default docker0 network

```
$ docker run --network database -d busybox top
```

- docker inspect formatted for bridge and IPAddress of container.

```
$ docker inspect -f \
{ .NetworkSettings.Networks.database.IPAddress }
}
$ (docker ps -lq)
```

```
172.19.0.5
```

Docker Network Metadata

© 2025 by Innovation In Software



Network Metadata

What is it? – Network metadata describes what the Network is through the use of key value pairs

- **docker network inspect <id>**

- Queries the Docker Engine, a json formatted output is returned

```
$ docker network inspect <network>
```

```
  "Name"=
  "Id"=
  "Scope"=
  "Driver"=
  "EnableIPv6"=
  "IPAM"=
    "Driver"=
    "Options"=
    "Subnet"=
  "Containers"=
```



Questions

Lab: Docker Networking